

Daily Notes - 2026 April



Apr 30., Thu
Agents Team



Apr 29., Wed
Agents Team



Apr 28., Tue
Agents Team



Apr 27., Mon
Agents Team



Apr 24., Fri
Agents Team



Apr 23., Thu
Agents Team



Apr 22., Wed
Agent Team



Apr 21., Tue
Agents Team



Apr 20., Mon
Agents Team



Apr 17., Fri
Agents Team



Apr 16., Thu
Agents Team



Apr 15., Wed
Agents Team



Apr 14., Tue
Agents Team



Apr 13., Mon
Agents Team



Apr 10., Fri
Agent team



Apr 9., Thu
Agents Team



Apr 8., Wed
Agents



Apr 7., Tue
Agents Team



Apr 2., Thu
Agents Team



Apr 1., Wed
Agents team



Agents Team

Gyula Halmos · SDK goes native, Telegram gets topics, and three bugs get cornered · Claude Agent SDK Infrastruct...

↑ Apr 30., Thu

Agents Team



Gyula Halmos

SDK goes native, Telegram gets topics, and three bugs get cornered

Claude Agent SDK Infrastructure changes

- Uplifted Claude Agent SDK from v0.2.111 → v0.2.123: `cli.js` (11 MB) replaced by a per-platform native binary (216 MB). Bun `--preload` interceptor mechanism is dead in the new model — Phase 1 ships without it; Will evaluate what will happen without it.
- Ripgrep sourced from `@vscode/ripgrep` instead of the SDK-vendored copy; per-platform binary handling added to build scripts, `electron-builder.yml`, and CLI install path

Telegram Messaging

- Designed and planned Phase A: one supergroup per workspace, per-session topic binding via `/pair`; adds `threadId` to `ChannelBinding`, supergroup config to `MessagingConfig`, and a setup UI
- Phase B planned separately: automations auto-create Telegram topics and bind them at session spawn time
- Investigated multi-bot model; concluded per-chat binding on a single bot is the right first step
- Fixed supergroup pairing bugs

Bug Investigation & Fixes

- **Session list collapse on screen lock/sleep:** root-cause narrowed — `refreshSessionsMetadataAtom` was destructive on stale reconnect. Fix: `removeMissing: false` mode for recovery paths, plus detailed breadcrumb logging on both renderer and server side
- **Chat Display stuck on loading spinner:** `loadedSessionsAtom` and `sessionAtomFamily` could get out of sync. Fix: resilient `messagesReady` derivation in `ChatPage.tsx` that falls back to in-memory messages and `messageCount` metadata
- **Issue #636 — local models with API key - loopback endpoint with API key returns 401:** `isLoopbackBaseUrl` check in `llm-connections.ts` was too broad, leaving `piAuthProvider` unset for keyed local servers. Fix: `isKeylessLoopback` flag

UX Polish

- Group by Unread mode: always render both UNREAD/READ buckets with item counts; i18n updated for 7 locales (en, es, de, ja, zh-Hans, hu, pl)

Next steps

- Further test messaging and automation changes
- Release 0.8.14

↑ Daily Notes - 2026 April

Apr 29., Wed



Agents Team

Gyula Halmos · Shipped 0.8.13 · DeepSeek V4 tested - so far positive experiences, really cheap and looks competent...

↑ Apr 29., Wed

Agents Team



Gyula Halmos

Shipped 0.8.13

- DeepSeek V4 tested - so far positive experiences, really cheap and looks competent enough
- Various quality of life fixes and bugfixes
- Now thinking level can be also set for automations

DevBP Meetup

- Held a presentation about AI native development in front of 55+ developers
- Got positive feedback all around 😊

Next steps

- Evaluate how to uplift to latest Claude Agent SDK

↑ Daily Notes - 2026 April

Apr 28., Tue



Agents Team

Gyula Halmos · Reliability hardening and automation quality push · Platform & Features · Advanced automation capa...

↑ Apr 28., Tue

Agents Team

Reliability hardening and automation quality push

Platform & Features

- Advanced automation capabilities by adding per-action reasoning/model controls and aligning docs/UI around them.
- Improved visibility of automation configuration to reduce ambiguity in scheduled runs.

Stability & Bug Work

- Focused on critical reliability issues around tool-call handling and session flow.
- Drove root-cause analysis and fix plans for message queue behavior and OpenAI-compatible/DeepSeek integration failures.

Delivery Execution

- Progressed multiple bugfix and follow-up branches in parallel, with clear test/validation plans to reduce regressions.

Next steps

- Land and verify the automation control updates end-to-end.
- Ship the tool-call reliability fixes with regression coverage.
- Finalize queued-message UX fixes to eliminate perceived dropped messages.

↑ Daily Notes - 2026 April

Apr 27., Mon



Agents Team

Christoph · BYOK Interface · Went over the config UI once more now that everything is hooked up and cleaned up th...

↑ Apr 27., Mon

Agents Team

Christoph

BYOK Interface

Went over the config UI once more now that everything is hooked up and cleaned up the UI.

Assistant Table Support

Finalized, fixed review issues and merged.



Add table cell reorder support for AI agent edits

#8050 · Opened by christophlocher · Merged by christophlocher

Merged

Reviewer  WiesnerPeti

 Improvement

 Security Checked

Shipped 0.8.12

- I spent a bit more time than expected on Pi and Claude AgentSDK uplift than expected
- Pi SDK uplift: enabled GPT 5.5 and DeepSeek V4
 - They modified how tools are registered so had to adapt our implementation mainly for the session scoped tool injection
- Claude Agent SDK after Anthropic's Post-Mortem about Claude issues wanted to uplift BUT they changed to a standalone binary. This changes a lot of things in our integration like interceptor mechanic. So I had to keep the current version
- A lot of smaller bugfixes that were prepared on Thursday and Friday

DevBP rehearsal

- We had a 2 hour+ rehearsal
- It was fruitful, discussed the overlaps with Felhő and got some minor comments too

Next steps

- On  Apr 27., Mon I am out of office will be back on  Apr 28., Tue

↑ Daily Notes - 2026 April

Apr 24., Fri



Agents Team

Gyula Halmos · Preparing 0.8.12 — 5 bugs diagnosed, plans ready to ship · Preparing for 0.8.12 · 991-message deep...

↑ Apr 24., Fri

Agents Team

Preparing 0.8.12 — 5 bugs diagnosed, plans ready to ship

Preparing for 0.8.12

- 991-message deep dive on `bugfix/small-fixes` branch — 5 distinct issues surfaced and planned

- **source_test auto-activation** — today purely validative; will flip `enabled: true` + call existing `onSourceActivationRequest` callback to load tools mid-session (no restart). Backend-aware `availability` covers Pi's next-turn delay vs Claude's immediate
- **spawn_session gains thinkingLevel** — and fixes a latent bug where `createSession()` silently ignored `thinkingLevel` in options despite declaring it in the type
- **WebUI "Add New Label" bug** — typing `#NewLabel` was firing `craftagents://` deep link and launching the desktop app. Root cause: `FreeFormInput` cherry-picked 6 fields from `EditConfig` but dropped `inlineExecution`. One-line spread fix
- **Attachments leak across sessions** — `FreeFormInput` holds attachments in `useState`, no `key={sessionId}` remount. Attach in A → switch to B → file ships to B. Correctness bug, not cosmetic. Plan widens `SessionDraft` to `{text, attachments}` with path-only storage
- **Custom URL scheme passthrough (#590)** — `obsidian://`, `vscode://`, `zed://` links blocked by hard-coded allowlist. Switching to blocklist (`javascript:`, `data:`, `file:`, ...) + anchor-click fallback for WebUI

DevBP MeetUp slide polish (Zita's visual feedback)

- Dark theme now fully black background
- Fixed `git commit --no-verify` copy text in the slides
- SVG graphics redesign pass (long session, errored near the end — may need follow-up)

Next steps

- **Uplift Claude Agent SDK — high priority.** Boris confirmed the regression lives in the harness, not the model
- Rehearsal for the **DevBP MeetUp** talk
- Continue **MobileUI** polish

↑ Daily Notes - 2026 April

Apr 23., Thu



Agents Team

Gyula Halmos · Shipped 0.8.11 and unlocked the mobile dev loop · Release & Features · Shipped v0.8.11 — single-co...

↑ Apr 23., Thu

Agents Team



Gyula Halmos

Shipped 0.8.11 and unlocked the mobile dev loop

Release & Features

- Shipped **v0.8.11** — single-commit bugfix release for follow-up quote truncation (closes OSS #580). Extracted ChatDisplay follow-up helpers into a tested module, removed the dead 140-char noteLabel cap, and added a narrow subpath export so unit tests can import pure follow-up state helpers without the full UI barrel.

Bug Investigation & Plans

- **Issue #576** — Pi / Craft Agents Backend ignored `includeCoAuthoredBy: false`. Root-caused to the missing 7th arg on `getSystemPrompt`; plan adds a defensive fallback in the shared helper plus an explicit pass-through at the Pi call site.
- **Issue #577** — `.txt` attachment API error. Current `main` doesn't reproduce the reported shape; plan covers reproduction + a UX upgrade (inline small text content under 8KB to skip the Read roundtrip, also fixes a pasted-text side bug).
- **Issue #544** — `call_llm` max-turns hard fail. Plan supersedes PR #550 with a broader fix: raise `maxTurns` 1→10, recover partial text on both `SDKResultError` yields and thrown exceptions, `[Partial result - ...]` prefix, plus `main.log` observability.

Mobile UI Rework (mobile-ui-rework branch)

- Unblocked the actual fast dev loop — added a Vite proxy to `webui/vite.config.ts` so HMR on `:5175` can reach the server on `:9100` instead of falling back to full bundle rebuilds. Improved the single-instance lock error to mention `CRAFT_CONFIG_DIR`, and investigated a font `data:` CSP violation.
- Fluid onboarding on mobile — `StepFormLayout` was `w-[28rem]` fixed; swapped to `w-full max-w-[28rem]` (matching the workspace primitive), plus `p-4 sm:p-8` on `OnboardingWizard`. Two-line fix, all onboarding steps benefit.
- Compact `AppMenu` — hide keyboard-shortcut hints, Edit/View/Window submenus, and Quit on phone viewports; flatten Settings/Help subs to avoid sideways clipping.
- Floating '+' action button for New Chat in compact/mobile layout, scoped via the existing `isAutoCompact` predicate.

Research

- Explored 'Second Brain' (Tiago Forte's PARA / external knowledge-management approach) as inspiration for a Claude Code memory/knowledge layer.

Next steps

- Continue fine-tuning the Mobile UI and finish the outstanding graphic touch-ups.
- Evaluate a 'Second Brain'-style solution for Claude Code.
- Address Zita's remaining review comments.



Agent Team

Gyula Halmos · 0.8.10 out the door, image gen validated, memory feature designed · Release & Features · Cut Craft...

↑ Apr 22., Wed

Agent Team

 Gyula Halmos

0.8.10 out the door, image gen validated, memory feature designed

Release & Features

- Cut Craft Agent **0.8.10** from the bugfix branch
- Validated **image generation** end-to-end with both the OpenAI API and Gemini API as sources — worked great, ready to integrate when we want it

Planning & Architecture

- Designed a **memory implementation for Craft Agents** based on Hermes Agent's memory solution — full plan drafted
- Surveyed the **2026 agent landscape** — changelogs from Hermes Agent, Open Claw, Pi-mo etc. to inform the roadmap

Content & Outreach

- Prepped the "**AI-native development: what actually changed in my workflow**" Developer Meetup talk
- Reddit sweep for Craft Agents / agents.craft.do / craft-agents-oss mentions

Next steps

- Polish the **Mobile UI** to make it finally a nicely usable PWA
- Fix any urgent issues that surface with **0.8.10**

 **Zsolt Varnai**

 Apr 21., Tue

Assistant

- **BYOK**
 - Fixed an issue with help agent causing incorrect explore/execute mode switching after thread reset
 - Updated thread model with provider config
- **Other**
 - Strengthened cell reorder API based on review comments

Knowledge transfers

- Knowledge transfer sessions throughout the day

 Apr 22., Wed

- Continue knowledge sharing

Apr 21., Tue



Agents Team

Gyula Halmos · Messaging marathon — UX polish, headless parity, and a cross-model PR review · Messaging UX & S...

↑ Apr 21., Tue

Agents Team

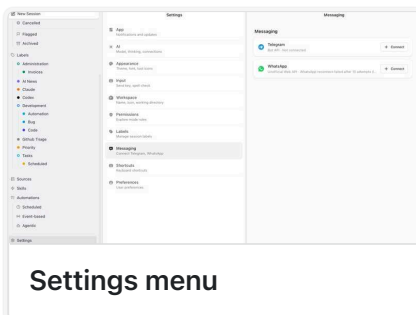
Gyula Halmos

Messaging marathon — UX polish, headless parity, and a cross-model PR review

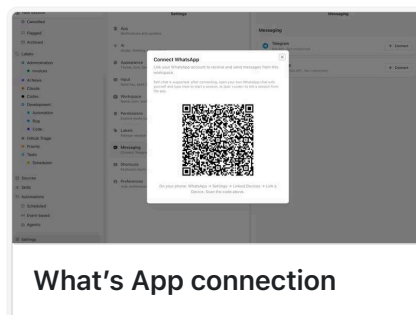
Messaging UX & Settings

- Ported the V2 `SettingsCard` redesign (Telegram + WhatsApp rows, inline bindings) from playground to production `MessagingSettingsPage`
- Session list now shows colored inline pills for bound platforms instead of the 📱 emoji
- Added playground preview wrappers so messaging UI can be iterated without running Electron main

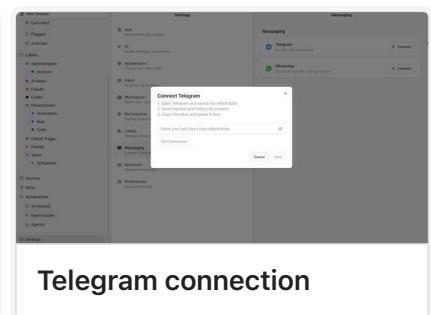
UI images



Settings menu



What's App connection



Telegram connection

Messaging — Depth & architecture

- Plan to forward `SubmitPlan` to Telegram as inline **Accept / Accept & compact** buttons
- Plan for headless server parity via a shared `createMessagingBootstrap` helper — single messaging stack across Electron and `packages/server`, no business-logic fork

Code Review

- Thorough cross-model review of PR #320 (messaging gateway) — separated real blockers from overstated review concerns and turned it into a concrete merge-gate plan

Bug Investigation & Planning

- **Follow-up island stuck-open bug** — root-caused the 220ms grace window silently swallowing close events; proposed deferred-close fix that still protects the enter animation
- **Issues #567 / #568 (0.8.10)** — `enable1MContext` defaults to `true` on Opus 4.7, triggering the 1M beta header without Tier 4+ access. Flip default, re-add Opus 4.6 as 200K fallback, drive auto-compaction from the app
- `electron-dev.ts`: re-enable `node --check` syntax verification (prior rationalization was wrong)

Next steps

- Final use-case testing of the Messaging feature end-to-end
- Create the DevBP presentation
- Release **0.8.10**

 **Zsolt Varnai**

 Apr 20., Mon

Assistant

- **BYOK**
 - Restructured BYOK provider / model structure
 - Introduced new providers for local and debug options
 - Updated core model handling
- **Model switching**
 - Improved model switch and fallback behavior
 - Force craft credits for help agent
 - Auto switch non agent models when agentic support is required
 - Handle missing model fallback
 - Update when default model is saved

Knowledge transfers

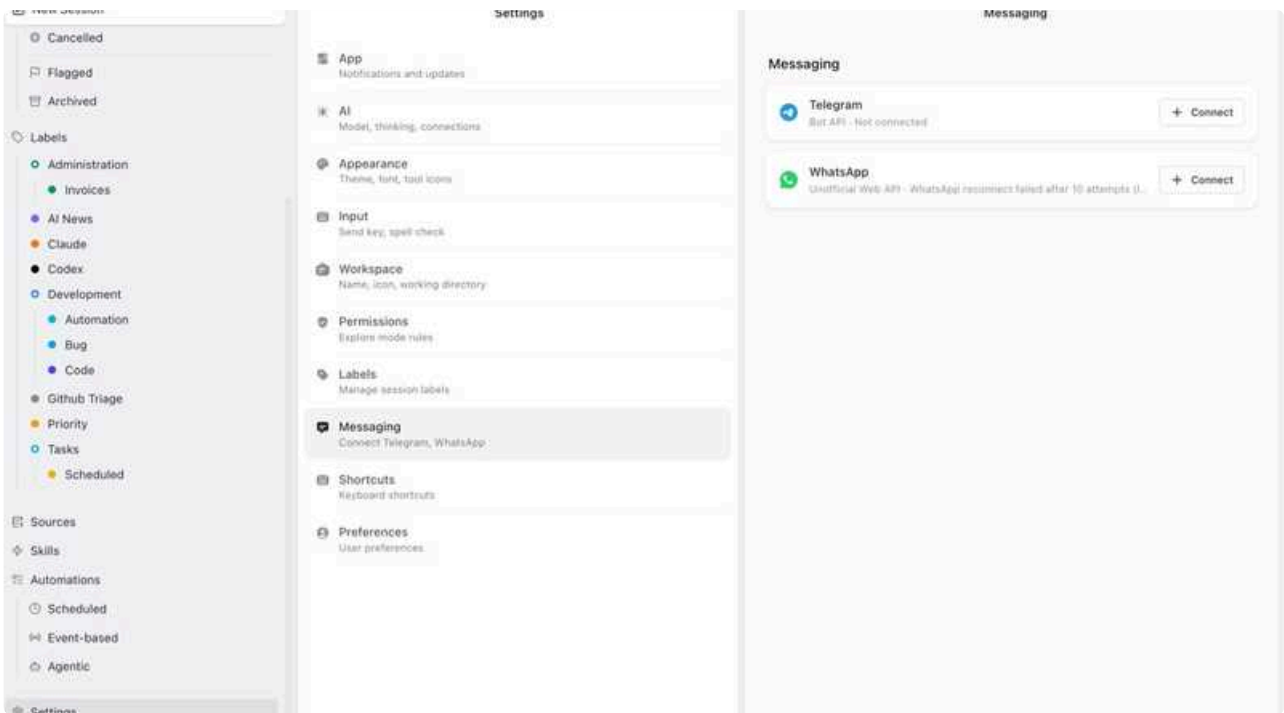
- Started services knowledge transfer sessions
 - Observability platform knowledge sharing
 - Generic service knowledge sharing session

 Apr 21., Tue

- Continue knowledge sharing
- Continue background mode improvements


↑ Apr 21., Tue · Agents Team

Settings menu




↑ Apr 21., Tue · Agents Team

What's App connection



↑ Apr 21., Tue · Agents Team

Telegram connection



↑ Daily Notes - 2026 April

Apr 20., Mon



Agents Team

Christoph · Fri, Apr 17 · Continued with BYOK Settings implementation · Settings look good now · For the configurati...

↑ Apr 20., Mon

Agents Team

Christoph



Apr 17., Fri

- Continued with BYOK Settings implementation
 - Settings look good now
 - For the configuration panel to be opened from the AI Chat I build a prototype version, for the cleanup of the UI, the structure needs to be updated
 - For Zsolt, these are the next changes that are necessary:
 - **Structure changes**
 - Add GPT Nano as "Core" to the Craft Credit Models

- Introduce "Local" as a new provider
- Maybe we should display "Debug" as a separate provider?
- **new structure after:**
 - Craft Credits
 - Core (GPT Nano)
 - Fast (Haiku)
 - Max (Sonnet)
 - Local
 - AFM
 - Ollama
 - OpenAI
 - GPT Nano
 - GPT 5.4 Mini
 - GPT 5.4
 - Anthropic
 - Haiku
 - Sonnet
 - Opus
 - Debug (craft employees only)
 - ...
- Discuss how we handle **automatic model switching**, e.g. when using smart search or the help agent mode
 - Help Mode → Always switch to Craft Credits/Fast (but new thread then uses User Preference)
 - Otherwise try to keep user preference, don't switch
 - Exception:
 - Preference is local model, mode needs agentic → Switch Craft Credits /Fast just for this thread, new thread honors User Preference
 - Preference is non-agentic → Keep current provider, switch to first agentic model

 Apr 20., Mon

- Create Plan for the Task Management improvements

Goal of the Week

- Create presentation for DevBP
- Introduce Messaging integration (Telegram and WhatsApp, iMessage best effort)
- Compose a roadmap for Craft Agents for the next 2 weeks

Messaging gateway deep-dive: Telegram wired, WhatsApp policy locked

Planning & Architecture

- Produced full messaging-gateway plan covering i18n fixes, Telegram pairing end-to-end, Baileys-based WhatsApp worker, and hardening — sequenced as 6 PRs (~2370 LOC). Decisions locked: **Baileys** for WhatsApp (precedent: Hermes Agent, OpenClaw), **per-workspace** bot scope, disclaimer surfaced in UI.
- WhatsApp platform policy: **transport-only**. No plan acceptance, permission approval, credential submission, or destructive confirmation happens over WhatsApp — those stay desktop-only. `/pair`, `/bind`, `/new`, `/unbind` remain allowed as routing operations.
- Bundled systemic fixes into Phase 1: relocate `bindings.json` out of the user's workspace folder (risk: accidental git-commit / iCloud sync), fix `TelegramAdapter.webhookSupport` metadata lie, rate-limit `generatePairingCode` RPC, add an RPC integration test that would have caught the `messagingRegistry` -not-in-deps bug.

Bug Investigation

- Root-caused API error `message_end` for Pi/codex-mini on `bugfixes-0.8.9` — three plans: add `xhigh` thinking level (Anthropic's new Opus 4.7 tier between `high` and `max`), DRY up thinking-level lists, fix codex-mini ChatGPT account mapping.
- Identified runtime-state bug in messaging: app treats persisted `config.enabled` as live connectivity → WhatsApp appears Connected after restart even with no adapter running. Fix requires a separate runtime-state model (`configured`, `connected`, `connecting`, `reconnectRequired`, `lastError`).
- Binding data-model honesty: backend supports N bindings per session but renderer collapses to `sessionId → one binding`, hiding rows and making disconnect misleading.

Infrastructure & DevOps

- Created `craft-agents-bugfixes-0.8.9` worktree from latest main for 0.8.9 bug triage.
- Summarized Ruff linting configuration across the monorepo.
- Proposed a dedicated structured log at `~/.craft-agent/logs/messaging-gateway.log` (JSONL, per-component `workspaceId/sessionId/platform/channelId` fields, rotation) to replace ad-hoc `console.*` usage in the messaging stack.

Next steps

- 0.8.10 release prep
- Messaging integration polish and architecture review

 **Zsolt Varnai**

 Apr 16., Thu

Assistant

- **Thread/session handling**
 - More testing with session state
 - Fixed thread unread state on panel open
 - Fixed socket disconnect race condition with unread state handling
- **Background sessions**
 - End to end tested various background session scenarios
 - Investigated and fixed issue with background session not being able to reconnect properly in certain cases
- **Support**
 - Checked app and backend changes about collections and tables
 - Created a fix for the table reorder API

This week

- Continue with thread/session handling improvements
 - Other BYOK improvements
 - Finish background mode + thread sync improvements
 - BYOK key remote storage + sync
- Start services handovers

↑ Daily Notes - 2026 April

Apr 17., Fri



Agents Team

Gyula Halmos · Apr 16., Thu · Double-release day: 0.8.8 + 0.8.9 out the door (Opus 4.7 default), messaging gateway...

↑ Apr 17., Fri

Agents Team

 **Gyula Halmos**

 Apr 16., Thu

Double-release day: 0.8.8 + 0.8.9 out the door (Opus 4.7 default), messaging gateway gets a shape

Release & Features

- Shipped **v0.8.8** — inter-session `send_agent_message` tool, zh-Hans translations, local model docs clarified; bug fixes for retry button, local model setup (no-key flow), ConfigWatcher duplication on headless Bun/Linux (#538), thinking-level IPC handler
- Shipped **v0.8.9** — **Opus 4.7 now the default Claude model**; also uplifted Claude Agent SDK from **0.2.78 → 0.2.111** (watch for surprises)

Bug Investigation & Fixes

- **Floating button race condition** — deep-dive identified a mount/unmount race in the follow-up button that caused intermittent sticking
- **Retry button** — now resends the last user message via position-aware lookup (handles the case where user typed another message after the error) instead of just focusing the textarea
- **Local model auth** — root-caused two interacting bugs: UI not sending `customEndpoint`, backend downgrading `pi_compat` → `pi` on loopback endpoints. Loopback base URLs now get `authType: 'none'` instead of prompting for an API key
- **Local model selector** — fixed display in selector

Planning & Architecture

- **Messaging gateway design** — 866-msg design session for first-pass gateway supporting **Telegram + WhatsApp** (iMessage to follow). Covers: `MessagingGatewayRegistry`, post-then-edit Telegram streaming at 3500ms cadence, inline permission buttons with `perm:` button IDs, 6-digit pairing-code flow (5-min expiry, rate-limited), Settings → Messaging page, session binding badges via Jotai atom, `confirm_messaging_pairing` / `unbind` / `list` tools on `SessionToolCallbacks`
- **Repo migration exploration** — mapped potential migration paths if the Craft Agents app name changes (opensource + private repos to new GitHub org)
- **Hermes Agent & OpenClaw analysis** — reviewed how each connects to Telegram/iMessage/WhatsApp; compared against Vercel Messaging SDK as alternative for gateway abstractions

Experiments

- **Local LLM testing** — installed LM Studio via Homebrew, spawned local Gemma session to verify answer accuracy on the Hungarian PM question

Next steps

- Test messaging gateway end-to-end — exercise Telegram + WhatsApp paths, fix bugs, fine-tune UI flows
- Monitor Claude Agent SDK 0.2.111 behaviours after the bump and fix regressions if they surface

 **Zsolt Varnai**

 Apr 16., Thu

Assistant

- **Thread/session handling**
 - Refactored thread state handling and read/unread handling
 - Removed reliance on local server state to manage unread threads as it was unreliable and use socket connected state
 - Added new REST API to fetch and manage thread state without having to connect to the socket
 - Improved socket lifecycle, so the socket is not connected before the assistant panel is open
 - Threads are fetched earlier, even before panel open, so the unread badge can be correctly shown on the collapsed panel
 - Now this can be done via the REST API without having to connect to a socket
 - Fixed thread unread state handling on stashed sessions
 - Fixed thread unread state handling while the panel is collapsed
 - Removed redundant API calls during app start
 - Fixed issues when the unread state goes out of sync between the client and SQL

 Apr 17., Fri

- Continue with thread/session handling improvements

↑ Daily Notes - 2026 April

Apr 16., Thu



Agents Team

Gyula Halmos · Squashing bugs for 0.8.8 while mapping the agentic battlefield · Bug Fixes (bugfixes-0.8.7) · Investig...

↑ Apr 16., Thu

Agents Team

 **Gyula Halmos**

Squashing bugs for 0.8.8 while mapping the agentic battlefield

Bug Fixes (bugfixes-0.8.7)

- Investigated and fixed headless server hang (craft-agents-oss #538) — root cause was duplicate ConfigWatcher creating recursive fs.watch on the same dir tree; added skipConfigWatcher flag + runtime duplicate guard
- Fixed SET_DEFAULT_THINKING_LEVEL RPC handler silently returning undefined instead of success true — config saved fine but UI threw a TypeError
- Fixed Retry button not focusing ChatInput — stale data-chat-input DOM query replaced with session-scoped dispatchFocusInputEvent

Features & Design

- Designed send_agent_message tool for inter-session messaging — sender envelopes, self-send guard, fire-and-forget delivery with reply support - for multi agentic framework
- Designed Dispatch Session QR code overlay for WebUI access from Electron - SPIKE ONLY

Research & Content

- Deep research on agentic AI landscape — surveyed X.com influencers, Reddit communities, recent developments
- I am still working on the research doc, but here is the current version: (cc  [Peter Adam Wiesner](#))



Current Agentic Landscape

Executive Summary · The agentic AI landscape in Q2 2026 is defined by five converging d...

Last Edited 4/16/2026

Next steps

- Investigate reported UI issues for 0.8.8 release
- Test agent orchestration harness — send_agent_message end-to-end
- Test LM Studio with Craft Agents together with Ollama
- Clean up Agentic Landscape doc, put together a Craft Agents wishlist

Christoph

 Apr 16., Thu

- Syncs & Planning
- BYOK Settings Design & Alignment with Zsolt and Christian

 Apr 17., Fri

- BYOK Settings Implementation

 Apr 15., Wed

Assistant

- **BYOK**
 - **OAuth**
 - Store refresh token in keychain and track expiration time and add automatic refresh mechanism when close to expiration or expired
 - Improved error handling for auth issues
 - **Client connections integration**
 - Discussion with Christoph about the connections data model and implementation approach
 - Allow switching between existing connections as default, including craft credits
 - Updated backend model list and tagging schema to allow the client to fully filter on the relevant criteria
 - Calculate available models based on connection and mode (craft token vs BYOK)
 - Fix default model selection and fallback based on available models
 - Add reauthentication option for oauth connections
- **Thread/session handling**
 - Started testing and making the session/thread handling more robust
 - The S3 based approach seems unreliable and performance scaling badly with old threads accumulating over time, so decided to have a cleaner approach, introducing an SQL table in DB2 for storing thread metadata
 - Created new table and wired thread related APIs to use hybrid approach and on the fly migration of existing threads
 - Fixed issue with unread thread state going out of sync between in memory and persisted state

Other

- Catchup and knowledge sharing with Peter about app related topics
- Review and cleanup API keys with Christian and Christoph

 Apr 16., Thu

- Continue with thread/session handling improvements

↑ Daily Notes - 2026 April

Apr 15., Wed



Agents Team

Christoph · Tue, Apr 14 · Change of plans, instead BYOK I focused on planning and discussing future projects with Vi...

↑ Apr 15., Wed

Agents Team

Christoph



Apr 14., Tue

- Change of plans, instead BYOK I focused on planning and discussing future projects with Viktor
- Another round of testing & bugfixing the AI sidebar with Peti – we are getting there!
- Still not feeling to good, so I left earlier to lay down



Apr 15., Wed

- Design team planning & alignment
- BYOK UI (hopefully)

Shipping 0.8.7, hunting bugs, and continuing exploring the current agentic landscape

Release & Bug Fixes

- Released Agent SDK v0.8.7, coordinated via Slack
- Investigated stale .server.lock causing windowless launch (#522) — planned fix with JSON lock format, boot-time stale detection, and process.on('exit') safety net
- Fixed Bedrock model validation error messages — one-line pattern match addition

CX Feedback & Fixes

- Session with CX Team to collect feedback on agent usage pain points
- Implemented and delivered optional token renew endpoint for API sources (MVP) — non-OAuth bearer token auto-renewal
- Added _rawBody support for API sources — cherry-picked into bugfixes-0.8.6

Research & Exploration

- Deep research on agentic memory systems — surveyed GitHub repos and approaches for agent memory (146-message session)
- Continued evaluation of agentic tools: Claude Code, Codex App, Cursor 3, Open Code — trends and gaps analysis

Next steps

- Conclude the agentic tools exploration and draft a summary document of findings
- Investigate local model usage in Craft Agents (Ollama, LM Studio, etc.)

 Apr 14., Tue

- **BYOK**
 - **OAuth integration**
 - Continue with the solution based on the HTTP server hosted on localhost
 - Created end to end oauth flow from the mac app allowing sign in with OpenAI credentials
 - Setup the server and temporary host the callback endpoint during oauth
 - Added PI integration for the codex open AI endpoint which has a different contract (supported by gateway already)
 - Forward the oauth token to the codex gateway endpoint
 - After fixing some issues managed to make the flow work end to end
 - Investigated some issues with Christian about the gateway record not appearing in the DB
 - **Client connections integration**
 - Added client support to manage different connections
 - Added new connections list in the settings page
 - Anthropic/OpenAI API keys and OpenAI OAuth credentials can be added as connection, which are stored in the data store (currently local only user defaults, can be synced later)
 - The keys are stored in keychain
 - Multiple connections on the same type can be added, connections can be renamed or deleted
 - Default connection can be selected for new threads
 - Thread state will be extended with connection info, so a thread will store both the connection and the selected model

 Apr 15., Wed

- Continue with BYOK implementation

↑ Daily Notes - 2026 April

Apr 14., Tue



Agents Team

Christoph · Mon, Apr 13 · BYOK planing with Zsolt · Started working on BYOK settings interface · Testing session wit...

↑ Apr 14., Tue

Agents Team

Christoph

 Apr 13., Mon

- BYOK planing with Zsolt
- Started working on BYOK settings interface
- Testing session with Peti to finalize his AI sidebar PR
- Long planning session with Viktor
- One on one with Zita

 Apr 14., Tue

- Continue with BYOK
- Planning with Viktor

Unfortunately I picked up a cold on the airplane and am feeling a bit feverish now. Will check tomorrow morning how I'm feeling.

Teaching the agent Magyar, Deutsch, and Polski while patching the pipes

Localization

- Added Hungarian, German, and Polish translations to Craft Agents (~3,180 UI keys total), PR #307 submitted
- Quality-reviewed translation output across all three languages

Bug Investigation & Planning

- Planned fix for session list not re-rendering after macOS sleep/wake — root cause: scattered `store.set()` calls need batching into a single Jotai write atom transaction
- Planned fix for file path validation (#526) blocking files outside home directory — refactoring 3 duplicate validators into one shared function with Windows path support

Research & Tooling

- Researched latest Claude Code, Codex, and agentic coding tool updates from X.com and Reddit
- Installed and looked at UI features of Codex app and Cursor 3, evaluated projects feature of ChatGPT
- Multi-model session day: Claude Opus 4.6, GPT-5.3-codex, GPT-5.4 all used across sessions

Next steps

- Continue evaluation of agentic tools (Claude Code, Claude Cowork, Codex App, Cursor 3, Open Code) — trends and gaps analysis
- Release 0.8.7 with the localization changes
- Implement the two planned bug fixes (session list reconnect, file path validation)

- **BYOK**
 - Started planning and implementing the various BYOK flows
 - **Anthropic API key**
 - Added custom key configuration to the generate API to be able to pass custom API keys
 - Christian deployed a testable proto of the new gateway endpoint to sandbox, so the goal was to test end to end with that
 - The custom API key is sent in an extra header to the gateway, while keeping the internal token in the Auth header. Then the gateway forwards the users key to the provider
 - The generation worked with a custom API key and the gateway stored it properly as a **non billable** request 🎉
 - **OpenAI API key**
 - Added openai key config option as well. The PI SDK picks it up and uses an OpenAI compatible contract and redirects to the new OpenAI compatible proxy endpoints of the gateway.
 - **This flow also worked end to end**, however we still need to spend time on validating the assistant behavior with these models. Back then with 5.0 they were much less capable than the claude models, but the assumption is that with 5.4 these gaps closed
 - Did some price comparison between the Anthropic and OpenAI models and based on that the counterpart of **haiku** is **GPT 5.4** mini and the counterpart of **sonnet** is the **regular GPT 5.4**
 - **Model filtering**
 - Updated the available model list and added new tags so the client can apply smart filtering
 - When using craft credits or anthropic API key, then the Max/Fast models continue to be the claude ones
 - When using OpenAI API key, then the GPT models become the default options for Max/Fast
 - The non-agentic Core GPT nano option is only available with craft credits or OpenAI API key
 - **OpenAI oauth**
 - Did some research about the options to implement oauth, including how Craft Agents does it
 - The core mechanism to get the token is an HTTP server on localhost, that receives the token in the client
 - Created a prototype oauth flow to test the options and override the redirect URI to non localhost or a custom URI scheme for iOS, but it did not work due to OpenAI rejecting the request
 - Just like OpenClaw and Craft Agents we are mimicing Codex CLI and the only allowed redirect URI is localhost
- **Did some research about the possibility of spinning up an HTTP server on a custom port on localhost from an iOS app. Technically it is possible and some people even said, that their app got approved in the AppStore. I also know an application in the appstore that does it, so it does not seem to be impossible, but it involves some risk until we get approved.**

- I was also looking for an alternative option, which would be just copy pasting the redirect URL from the browser to the Craft app, but this flow won't have that good UX due to the localhost redirect failing, so it might be worth to continue with the localhost server approach

📅 Apr 14., Tue

- Continue with BYOK implementation

↑ Daily Notes - 2026 April

Apr 13., Mon



Agents Team

Christoph · Fri, Apr 10 · Continued tweaking the sidebar trying to fix some of the animations, then pushed all of it an...


↑ Apr 13., Mon

Agents Team

Christoph

📅 Apr 10., Fri

- Continued tweaking the sidebar trying to fix some of the animations, then pushed all of it and handed it over to the app team. This doc documents the current change and existing issues:



Current State Fri, Apr 10

✗ Changing widow size → Main panel content becomes invisible!? · Wide Window Width – S...



Last Edited 4/10/2026

- Extend the assistant collection editing capabilities to include all the new collection formats and properties. Luckily this requires only server-side changes, and all the applying of the changes on the client worked right away.

📅 Apr 13., Mon

- Prepare UI for BYOK

Goal of the week

- Create an overview of what are the current trends and where are the major players in the market (CC, Cowork, Codex, OpenCode...erc)
- Create a Roadmap for next major features for Craft Agents, finalize and summarize collab findings
- Continue the maintenance and localization

Shipped 0.8.6

- Concluded a long bug fixing patch of patches
- Lot of bug hunting, debugging with users
- Most of the issues and the more ceiticals were found and fixed

Next steps

- Continue the research of current directions and trends in the market. Pilot new features in CC and Cowork and Codex app
- Add German and Hungarian localization
- Critical bug fixes only

This week

- Progress with roadmap for the next release
 - Bring Your Own Key
 - Agentic GPT
 - Full background mode

↑ Daily Notes - 2026 April

Apr 10., Fri



Agent team

Christoph · Thu, Apr 9 · Full focus on the assistant sidebar mode PR. I had a sync with Peter in the afternoon. Overall...

↑ Apr 10., Fri

Agent team

Christoph

 Apr 9., Thu

- Full focus on the assistant sidebar mode PR. I had a sync with Peter in the afternoon. Overall all of the behaviors worked with the refactored architecture but some animations didn't anymore (other ones than before! the complex ones now worked, the simple ones didn't). Peter gave me some suggestion for debugging them, which I started but didn't finish yet.
- Did a quick fix I've been wanting to do for a long time: On Github branch names can contain a slash, but from the URL you cannot know what is part of the branch name and what is part of the file path. Because of this all Smart Links to a branch with a slash in the name (as we usually prefix our names these were a lot) showed "Directory not found". I asked the agent if there is quick fix, and it turned out Ferenc already implemented a fix for the /blob route but not the /tree route – if the first request fails try again, inlacing the first part behind the slash in the branch name. Levy deployed it and now it works, eg. here:



lukilabs/Luki

Directory

 lukilabs / Luki Branch christoph/assistant-sidebar-2

Which before looked like this:



Github

Either the directory doesn't exist or you don't have permission to view it

Directory not found

Data is Outdated

 Apr 10., Fri

- Finish Sidebar Mode PR & Assistant Collections Schemas

Pi SDK Ships, Localization Gets the Full Treatment

Release & SDK

- Pi SDK update: major refactoring and bug fixes across a 418-message session, culminating in PR #297
- Investigated Pi SDK ChatGPT Plus authentication flow
- 0.8.5 release prep — updated release notes with new model compatibility (GLM 5, GLM 5.1, Minimax 2.7) thanks to Pi SDK upgrade

Localization

- Comprehensive localization architecture review (955-message marathon session) — analyzed implementation, identified unlocalised strings, moved work to release-0.8.5 branch
- Simplified Chinese localization status check and plan review
- Performance analysis of localization system — reviewed potential runtime impacts
- Quick investigation into Chinese translation differences (Simplified vs Traditional)

Next steps

- Monitor localization release
- Do some performance checks
- Investigate bugs reported by users related to headless setup

↑ Daily Notes - 2026 April

Apr 9., Thu



Agents Team

Christoph · Wed, Apr 8 · Misc · Created graphics for What's New and aligned on copy with Lisa · Assistant Capabilitie...

↑ Apr 9., Thu

Agents Team

Christoph



Apr 8., Wed

- **Misc**
 - Created graphics for What's New and aligned on copy with Lisa

- **Assistant Capabilities**

- Started extending collections schemas to support the new modes

- **Assistant Sidebar Mode**

- Prepared a planning doc
- Had a very long sync with Peter, we looked at the code together. Peter recommended not starting a new implementation for production, but rather refactoring the current branch. We discussed how to clean it up, switching from centralized control of the derived properties by the SplitView, to a more local approach with computed properties for deriving component specific state, and the updates flowing through the UI through the automatic calling of the layout functions. *"UI is just a river."* (said either Confucius or Peti)
- Got started on these refactoring and made good progress. It's quite complex for me to refactor the data flows while keeping the behavior we agreed on the same time, but I think I now see how it will all work out.
- We agreed to have another call tomorrow to check my progress and plan the next changes

 Apr 9., Thu

- Finish AI sidebar refactoring and prepare PR
- Continue assistant editing fixes

 Gyula Halmos

Shipped 0.8.4

- Sharing resources is now out
- Lot of small fixes and refactors

Next steps

- Finally turn to localisation so next week with the help of  Xuefang Li we can release simplified chinese.
- Prepare and test Pi SDK uplift
- Monitor latest release feedbacks

↑ Daily Notes - 2026 April

Apr 8., Wed



Agents

Christoph Locher · Tue, Apr 7 · AI Sidebar · Tweaked the drop behavior to solve a few edge cases. I think now the lo...

↑ Apr 8., Wed

Agents

 Christoph Locher

 Apr 7., Tue

- **AI Sidebar**
 - Tweaked the drop behavior to solve a few edge cases. I think now the logic is solid. Because this is a prototype, there are a few cases where the animation is skipped, but overall the logical behavior is as I designed it – I'm now looking for some feedback on cases where it doesn't match user's expectation.



Github

- **Assistant Editing Fixes**



Github



- **Cards: Included in Styling, Improved setting Colors**

- Chris Messina mentioned some issues with cards in the community. I think they were due to an ambiguous prompt (the agent assumed by "also style this" he meant style the subpage, not style the card) so there is no super obvious fix, but I extended the Style skill so the agent also updates Cards when applying a style to a page, which was Chris' initial expectation
- Unrelated to this, I also noticed some issues with setting custom gradient and solid colors because dark mode variants were missing and isLight wasn't being set, so I adjusted it. Now when the agent sets cards' styles it follows all the same properties as when the user does it via UI.
- Unrelated to this, there was a bug in the Card configuration UI, where the preview would switch to the first default gradient whenever the background panel was opened, no matter which gradient was selected (one of the defaults or custom, set manually or by agent – didn't matter). This was a UI bug, which I also fixed.

- **Tables: Fixed Moving Existing Cells**

- The agent was able to edit tables, but it had difficulties reordering existing tables. We observed this ourselves and got some user reports, so I looked into it. It turned out to be an oversight, where the server diff between the old and new html only detected content changes, not position changes, so I added a new function for that which sends an updated cell position call to the client.
- However this was only half of it – it gets really grarly when the move at the same time also changes the structure of the table, e.g. when moving cells to a newly added column. This took a lot of back and forth and tweaking, because cell creation is asynchronous, so we need to record the old cells, then adjust the table's structure, then move the old cells to their new position even if they were deleted in the meantime (eg. because their row was removed and instead a new column inserted). This took many iterations, but now it seems to work relatively reliably. I think Zsolt should review the code once he's back, because it added a lot of new code.

 Apr 8., Wed

- Continue with Tables & Collection Editing Fixes
- Call with Peter to prepare production implementation of the AI Sidebar mode
- Prepare "What's New" graphics for the Growth team

0.8.4 preparation marathon, new source integrations, and prepping for the stage

Release Prep & Bugfixes (0.8.4)

- Massive dependency analysis and bugfix session
 - Session-scoped-tools refactor and simplification
 - Introduction of OAuth for API sources
 - Added end-to-end test for session-scoped tools (Claude and Amazon Bedrock for Pi in CI)
- Want to include resource sharing in the next release hence not releasing it yet

Infrastructure & DevOps

- Cloudflare Tunnel — investigated adding 2FA / additional authentication layer to the tunnel
 - Enabled additional OTP authentication on the tunnel to make sure that access to Craft Server is protected by 2FA
 - It was important because the CX Team is setting up Sources that can manipulate subscriptions

Planning & Architecture

- Designed approach for sharing automations and sources between workspaces (309-message deep dive in share-resources)
- Started testing it as well, needs some more manual tests to be sure that all scenarios work. This is the light version, as credentials are not transferred yet

Presentation Prep

- Prepared topic ideas for a 30-minute developer meetup talk on craft-agents-oss

Next Steps

- Finalize testing of shared resources and release it in 0.8.4
- Do a preliminary test of the localisation PR

IMPORTANT: I will be out of office after lunch. So i will be in the office only half day.

↑ Daily Notes - 2026 April

Apr 7., Tue



Agents Team

Christoph · Thu, Apr 2 · I kept working on the behavior for the assistant sidebar mode. · I updated my description to...

Christoph

📅 Apr 2., Thu

I kept working on the behavior for the assistant sidebar mode.

I updated my description to a state machine / state chart style description (it's not quite a state machine in the strict sense due to multiple inputs and outputs).

CRAFT

AI Assistant Sidebar Mode Documentation

There are

- Constants – like the thresholds for 1 or 2 sidebars
- Context – things that can be changed by the user by interacting with UI: window width, desired mode for the panel, open/close state of the sidebar etc.
- Actions – things the user does that can change the context (e.g. clicking the toolbar button)
- Derived State – the state properties as they are calculated from the constants plus the current context, like the actual mode for the panel

The logic of the behavior is encoded in the action transitions (how actions change context) and of course the `deriveState(from: context)` function.

During iteration I maintained a pseudocode description in TypeScript to give to the agent as a reference for the implementation.

I did lots of manual testing and also created a list of tests to check the behavior, and added a debug view to be able to observe if the states are updated correctly.

```
SIDEBAR DEBUG
contentWidth: 1511
threadActive: false
```

```
USER INTENT:
  right:  hidden
  aiMode: pinned
  aiState: closed
```

```
DERIVED:
  right:  hidden
  aiMode: pinned
  aiState: closed
```

```
BUTTONS:
  rightDisabled: false...
```

Overall most behavior now works as I want it in the prototype, there are few state transitions where the animations are currently not working as I want (e.g. when the sidebar changes to float sometimes the content width is not animated), but I think this is mostly because of the hacky prototype implementation. I highlighted these areas in my doc to keep an eye on them during implementation.

I also started working on the dragging, so you can drag the sidebar out to turn it floating or drag the floating panel to the side to snap it into sidebar mode.



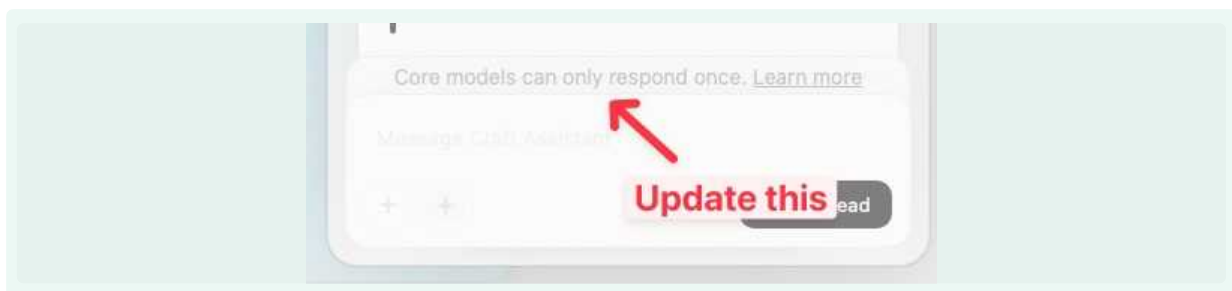
Github
Directory not found

<https://github.com/lukilabs/Luki/tree/christoph/assistant-sidebar-2>

📅 Apr 7, Tue

For the release from Viktor - cc: 👤 Christoph Locher

- **Craft Assistant bug / limitation** - <https://craft-community.slack.com/archives/C07FDJ6F8BG/p1775239134076449> - can you please check it?
- **Update wording**
We show this when ChatGPT is selected (prob we will need to update the Help site article as well)



Plan for the Week:

- Pinned Sidebar + Threading (1.5d)
I will tweak the dragging implementation a bit more, then ask for feedback and schedule an implementation planning session with Peter around mid-week.
- Collections & Tables Editing Fixes (2d)
We have some known issues with table editing via the assistant and I want to double check everything existing still works with collections, especially with Gallery and Kanban mode changes, and also support the new properties
- Free: 0.5d

Weekly goal

- Finish localization that  Peter Bobula prepared and have it ready for release with testing
- Evaluate collab features based on CX Teams's feedback and polish it - have a review of the UI
- Bugfixes
- Finish up small features that are prepared like sharing of skills,sources,automations (needs test)

Released 0.8.3

•

 **lukilabs / craft-agents-oss**
github.com/lukilabs/craft-agents-oss

 lukilabs  TypeScript Stars 5688

Customer care

- Followed up some threads and emails from users, also investigated some bugs, the ones i could reproduce fixed as well, some needs a bit more investigating

Assistant

• Guides

- Created a guide for **AI Agent Service V2** with basic info about which branch to deploy from, dashboards and recent service history

-  AI Agent V2 - Service Guide

- I still kept the original AI Agent Service guide for the V1 version as web is still using that service

-  AI Agent V1 (legacy) - Service Guide

- Shared it with the OOH rota channel too

• Stability

- Added PI subprocess inactivity timer + auto shutdown mechanism after 1 hour
 - This is to fully prevent even occasional instances of the subprocess being stuck in memory permanently

• Logging

- I was looking at log quality and log noise and making improvements
- Some errors didn't have an actual reason still. It turned out those are the user cancelled generations, so added termination reason to both the subprocess and the main process side to make a user cancellation caused it.
- Added subprocess stacktrace to error logs
- Gated verbose debug logs behind isLocal flag to reduce datadog noise
- Downgraded noisy errors to warnings as they were tied to disconnected clients

• AI tools

- I continued with end to end testing in sandbox the new workflows, including the interactions between the **agent-service-v2** and **ai-tools** services
- Fixed some configuration issues to enable proper space bundle caching behavior
- **Cache observability**
 - Tested uncached and cached space loading behavior
 - Added logs + stopwatch metrics to understand when full space bundle cache rebuilding or incremental build or full SQL fetch is needed
 - Added similar logs to the snapshot fetch + incremental build for individual document content

This week

On holiday this week 🌴



Agents Team

Christoph Locher · Wed, Apr 1 · Completed testing of AFM Turn-Mode and merged the PR, recorded an improved vi...

↑ Apr 2., Thu

Agents Team

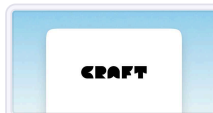


Christoph Locher



Apr 1., Wed

- Completed testing of AFM Turn-Mode and merged the PR, recorded an improved video
- Rebuild the pinned sidebar mode prototype on the new master and started fine-tuning all the different pinning and unpinning behaviors and combinations with the right sidebar and window size changes. I'm document everything as code in the prototype and as a written description, so we have all behaviors explicitly defined when creating the production implementation:



AI Assistant Sidebar Mode Documentation

Adding a pinned mode for the Assistant panel · Rules: · [DCW = Document Content Width] · ...



Last Edited 4/10/2026

This is tricky because to feel natural we can't just override the current state, but then also need to restore it when there is enough space again, following certain rules to avoid conflicts between right sidebar and AI sidebar



Apr 2., Thu

- Sync with Customer Success about AFM
- Continue Sidebar Prototype

22 sessions, 3 branches, zero chill

Mobile / Compact UI Overhaul

- Full compact mode plumbing + 6-workstream UX plan: session list redesign, composer cleanup, header simplification, accessibility pass
- Compact permission mode selector with bottom drawer, restored attach & sources in compact toolbar
- Fixed session info popover clipping → bottom sheet/modal pattern for narrow screens
- Browser-based UX testing on mobile emulation to collect actionable feedback

Stability & Infrastructure

- Provider type cleanup: LlmProviderType from 6 → 3 (removed dead bedrock/vertex/anthropic_compat paths), with startup migration
- User-friendly error handling: provider-aware messages, rendered action buttons (retry, status page, settings), MCP reconnect action
- Fixed Bedrock models not appearing until after first failed session — await model refresh during setup
- Fixed file:// links misrouted as URLs (blocked Office file opening from agent responses)
- Fixed headless server automations not firing — eager workspace init at boot + scheduler tick concurrency guard
- New session self-management tools: set_session_labels, set_session_status, get_session_info, list_sessions
- Axios supply chain vulnerability check (not impacted)

Collaborative Features

- ParticipantIdentity design and implementation: transport-scoped identity, remote workspace + WebUI display name capture
- Collaborative chat rendering: multi-sender detection, initials avatars, grouped sender runs, "you" marker
- Interruption attribution: "Response interrupted by {name}"

Code Review & PRs

- PR #291 i18n review — 14 issues found (stale translations from module-load-time t(), race condition in init, HTML rendered as literal text, unpersisted language pref)
- PR #293 fixes: credential sanitization (isTemplatedValue), slug consistency validation, removed unrelated settings changes
- Granular automations in resource bundles: per-automation export/import with conflict resolution, webhook auth sanitization
- Co-authored-by preference (#500): includeCoAuthoredBy flag to conditionally include Craft Agent trailer

Next steps

- Extended manual testing across all platforms and interfaces (web, mobile, desktop — Windows, Mac, Linux, and Docker) given the large volume of changes
- If all good, prepare the release
- Collaborative features won't ship before the holiday, but will be released to the CX Team tomorrow
- Discuss localization and see how much effort it is to get it to the following countries

( [Peter Bobula](#)):

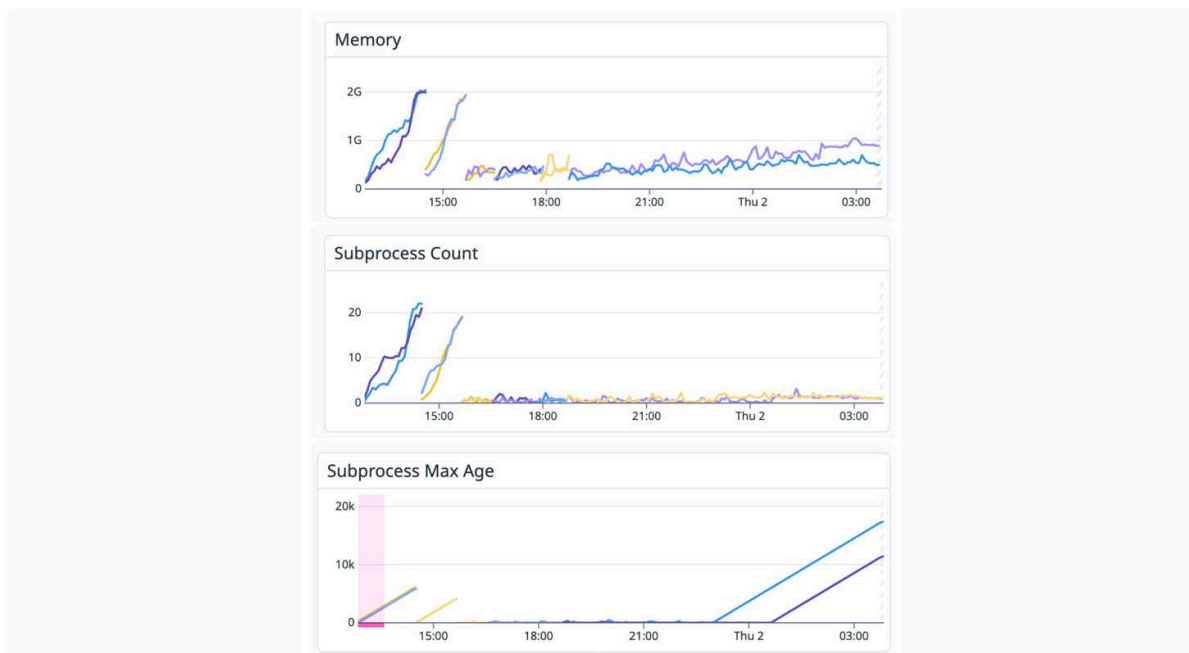
- **Chinese is the single clearest localization win** — CN+HK together represent **20.8%** of all downloads, nearly matching the US. Add Singapore's Mandarin users and it's the #1 non-English audience by a wide margin.
- **Japan is the #2 non-English priority** — strong market, and Japanese users tend to strongly prefer localized software.

2	 China	728	+15.4%
3	 Hong Kong	255	+5.4%
4	 Singapore	253	+5.4%
5	 Japan	217	+4.6%
6	 Germany	173	+3.7%
7	 Hungary	131	+2.8%
8	 United Kingdom	122	+2.6%

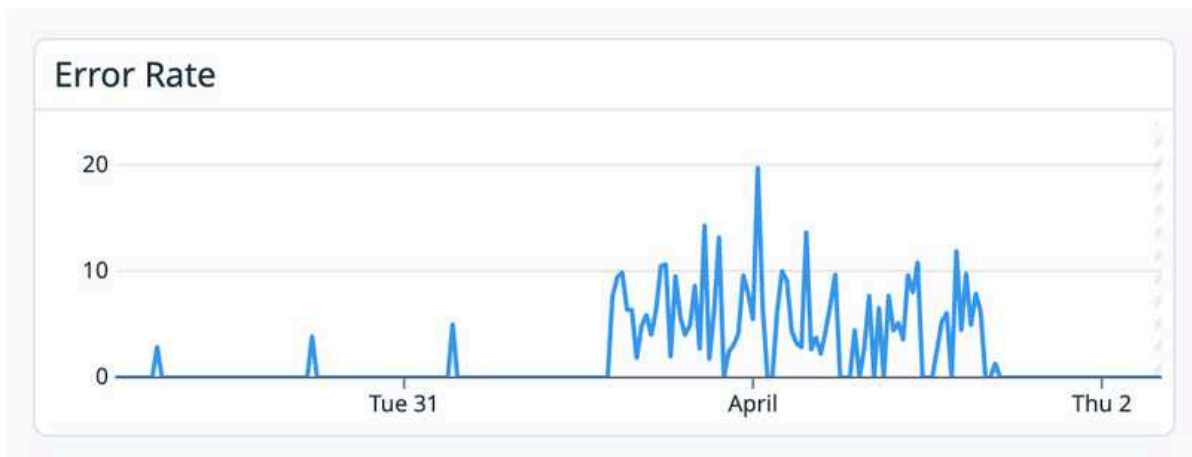
Assistant

• Out of memory issue

- Added various new metrics and extra instrumentation to find the root cause
- Managed to find the root cause, which was in the coordinator of the PI subprocesses. The issue caused some of the subprocesses to not shutdown correctly, which lead to the used memory not getting deallocated.
- The improvement is very significant, the subprocess count is almost fully flat over time, memory is close to being flat, although still slightly increasing over time, just like it did with Claude Agent SDK
 - This is not a problem, eventually the instances will recycle, but this won't happen too often
- Subprocess max age showed during the night, that very occasionally 1 PI subprocess might get stuck still. This is less urgent, but extra measures could be added to make sure they are shut down.



- Also talked to the services team, we are fine to increase memory to the maximum allowed level for the type of fargate instance. I increased it to 8GB per instance just in case, but we can reduce this back later on if it seems unnecessary.
- **Other stability issues**
 - Improved propagation of subprocess errors to datadog, although after the subprocess fixes no failed generation occurred, so I will keep checking this if it happens today.
 - Updated dashboard memory usage with new container level metrics and the other metrics also went back to normal after the subprocesses are correctly shutting down



- Reduced error noise by changing validation and other generation related events to warnings
- **Testing and improving new capabilities**
 - As some of the tools still had feature parity issues, I refactored 3 space tools that return text result, moving the text generation to the AI agent service and requiring JSON output schema from the mac and AI-tools implementations
 - Improved filtering behavior to hide template documents from AI-tools output to reduce behavior differences

Current State

- **New service stability in prod**
 - Managed to find and fix the issue in the component that was coordinating the PI subprocesses ✓
 - Memory usage went down drastically to normal levels (minimal increase over time as before) ✓
 - Very rarely a single subprocess might still leak as I saw it happen during the night, which can still be improved, but this is not in the urgent territory anymore ⓘ
 - Increased available memory to 8GB just in case ✓
 - With this setup service restarts should be very rare ✓
 - Error rate also went back to normal levels ✓
 - Concurrent generation metrics also went back to normal levels ✓
 - Updated the dashboards to show total instance memory including subprocesses ✓
 - Improved logging of detailed PI subprocess errors propagating up all the way to datadog ✓
 - Haven't seen any production errors happening since then, so I will keep checking on this ⓘ
- **New service testing with new capabilities**
 - Overall we spent less time on the new feature stability as planned due to prod issues this week, but in the remaining time I will keep testing/improving the stability and get back to it when I get back (we planned with enough buffer in the remaining weeks of the roadmap to do it)

📅 Apr 2., Thu

- I will update the service guide with basic info about the service and deployment in case changes has to be made next week

- Continue testing/improving the stability of the new feature set
 - AI tools service and server side tools, background sessions

↑ Daily Notes - 2026 April

Apr 1., Wed



Agents team

Christoph Locher · Tue, Mar 31 · Lots of syncing and planning with Zsolt · Sync with Viktor · Sync with Peter, discuss...

↑ Apr 1., Wed

Agents team




Christoph Locher



Mar 31., Tue

- Lots of syncing and planning with Zsolt
- Sync with Viktor
- Sync with Peter, discussed planning, App team support and Apple Foundation Models
- Adapted Apple Foundation Model to allow Turn-Based conversations




Fix AFM multi-turn flow and context-window disable UX

#7983 · Opened by christophlocher

Merged
Reviewer  WiesnerPeti
 Improvement
 Security Checked

Apr 1., Wed

- Finalize PR for AFM Turn-Based Conversation
After using it for a while, it reaches the token limit right away, even in a new thread, with the same doc and prompt that worked before. I need to investigate why this happens – an app restart solves it.
- Improve Pinned Sidebar POC to finalize all behavior (automatic sidebar closing, dragging between states) [next step once all behavior is decided is to create plan for production implementation with Peter's support]

Released 0.8.2

- Stability fixes
- Aesthetic improvements on webUI
- OAuth flow rework to accomodate webUI use cases

Remote agent for CX Team

- Spinned up a docker based remote agent on an old MacMini for CX Team
- Details:
 - URL: <https://cxteamagent.craftdocs-cf-t1.com>
 - Token for login: 31a99407b87deb2cd2f75efea31c7c9a53b776f0b59562a45c351ee69cdcf78a
 - cc  Peter Sajevics and  Tamas Fazekas and  Abner Cruz
- Also discussed some use cases and the feature parity with Claude Cowork with the CX Team

Next steps

- Follow what we discussed with Balint about the short term roadmap
 - Stability improvements
 - Mobile view polishing
 - Multi user shared workspaces with lightweight login

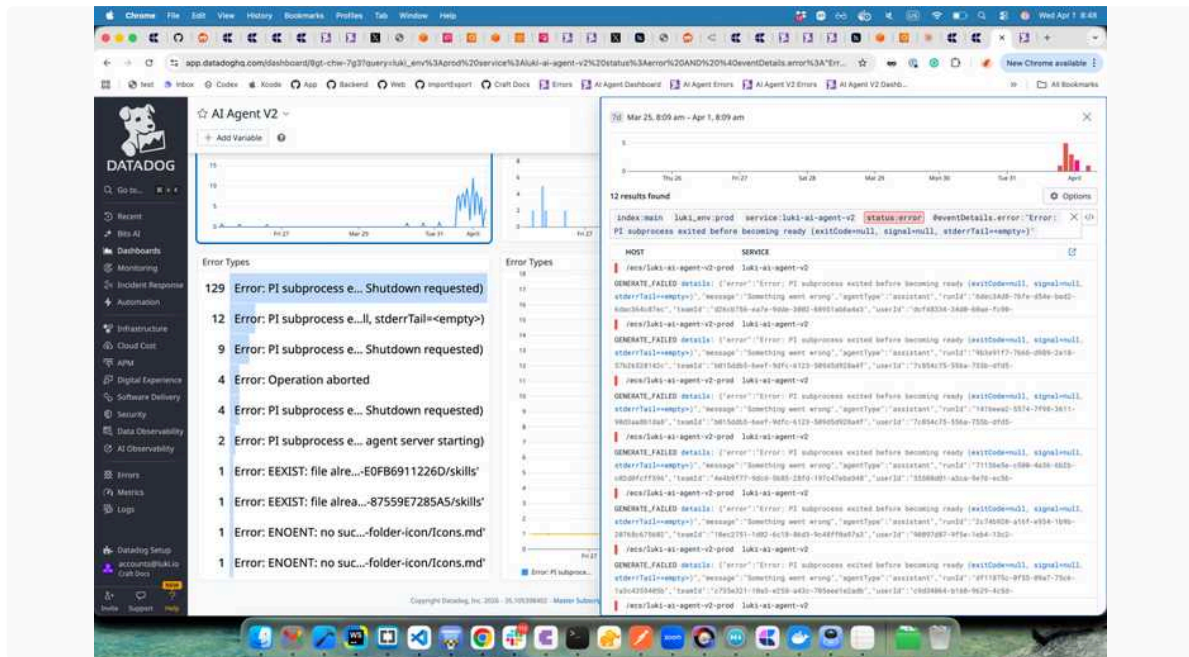
Assistant

- **Planning**

- Several discussions with Christoph to agree on the [next feature plans and weekly schedule](#)
- We are overall aligned on the main goals
- The decision has been made to not support oauth for claude models due to risk of being banned
- Due to this we pivoted to introduce **full agentic GPT support in "bring your own key" mode**
- **AI gateway dependencies**
 - Had planning discussion with Christian about AI gateway dependencies
- **Bring your own key**
 - It was requested that we keep the AI gateway proxy route for bring your own key, so we don't lose tracking capabilities
 - For that the gateway needs to support custom API key passthrough mode and differentiate postgres logs to not mix our costs with users custom key usage
 - This was discussed to be 0.5 day, lower risk work and suggested that Agents Team can do it on our own, so adjusted our schedule to fit into Week 16
- **OpenAI support**
 - This is a completely new integration, that we don't yet have with new proxy routes and new usage tracking
 - **It's estimated to be 3 days of work and services team can work on it next week**

- **Production Deployment**

- I deployed and smoke tested the V2 service in prod
- The service is mostly stable and fulfilling requests, but seeing multiple restarts every hour, even though the memory usage seems very low
- By checking AWS events I found, that the cluster kills the instances due to insufficient heap memory.
- It turned out, that due to our subprocess based PI setup we lost visibility on the subprocess heap and either subprocess might be leaking or the heap of the concurrently running subprocesses might be too large
- I started to add subprocess level fine grained instrumentation to have more visibility on what's happening on the subprocesses and a periodic memory sampler within the subprocess
- Also created a fix for the PI runner, that controls the subprocesses and potentially leaking
- I'm a bit worried about the complexity of the PI subprocess architecture and the issues it causing. I ran an investigation about potentially removing it and running PI in-process, but according to codex it's a complex and high risk change, so I'm holding off with that for now.
- ▼ There are also other subprocess errors that will require investigation



• Testing New Service Capabilities

◦ Tool Parity

- Basic generate sessions seem to work fine during end to end testing (server side prompt generation, server side ID mapping works fine)
- I created a tool result logging mechanism, that stores all tool call inputs + output pairs and replayed a long session with fully real space data with only client side tools and then only server side tools
- There were no failures, but diffing the results there are still quite a few differences to fix

◦ Remote sessions / session resume

- Found an issue with user prompts getting duplicated when a remote session is restored
- Added some fixes to the session storage mechanism to deduplicate and not sure unnecessary messages
- Also fixed storing system prompts and context messages properly

Current State

- **Plans and service dependency discussions were quite productive** ✅
- **New service stability in prod** ⚠️
 - Agent service with PI was deployed to prod ✅
 - It's stable enough, so I did not need to revert yet ✅
 - I'm not aware of user reports so far about significant issues ✅
 - The service restarts several times every hour by AWS ⚠️
 - Error rate is higher than usual ⚠️
 - Concurrent generation count is reported to be 5-6 higher than the previous version, which should be investigated, probably a tracking/session counting bug ⚠️

- Our custom memory reporting shows false data due to the PI subprocesses having their own heap ⚠️
- There are also PI subprocess errors and the subprocess exiting, but we don't have full visibility about the reasons ⚠️
- Im personally not sure at the moment if having PI subprocesses was the right approach in AWS, now we have a significantly more complex setup and less visibility on whats happening. I'm thinking about trying a setup without subprocesses, but the exploration I started about the switch considers it a complex and high risk refactor, so I'm exploring other options first. ⚠️
- There are many things to investigate that worry me and compete for my focus, so Im much less confident than yesterday ❌
- **New service testing with new capabilities**
 - Core things like server side prompt generation and server side ID mapping seems to work fine so far ✅
 - Still found many issues during full manual end to end testing with the new capabilities (tool result differences between mac and AI-tools service, remote session issues) ⚠️
 - These are lower urgency issues as they are needed in May release and we have buffers in the weekly plans, so for now my focus is more on the service, than the new capabilities ✅

📅 Apr 1., Wed

- Continue investigating prod issues
 - Might create a version without PI subprocesses at least to see how stable that is or even just to give us more visibility on whats happening, but it might break hooks and other things, so its a high risk change.