

Daily Notes - 2026 March

-  **Mar 31., Tue**
Agents Team
-  **Mar 30., Mon**
Agents Team
-  **Mar 27., Fri**
Agents Team
-  **Mar 26., Thu**
Agents
-  **Mar 25., Wed**
Agents Team
-  **Mar 24., Tue**
Agents team
-  **Mar 23., Mon**
Agents Team
-  **Mar 20., Fri**
Agents Team
-  **Mar 19., Thu**
Agents team
-  **Mar 18., Wed**
Agents team
-  **Mar 17., Tue**
Agents Team
-  **Mar 16., Mon**
Agents Team
-  **Mar 13., Fri**
Agents
-  **Mar 12., Thu**
Agents Team
-  **Mar 11., Wed**
Agents
-  **Mar 10., Tue**
Agents
-  **Mar 9., Mon**
Agents
-  **Mar 6., Fri**
Agents
-  **Mar 5., Thu**
Agents
-  **Mar 4., Wed**
Agents Team



Mar 3., Tue

Agents Team



Mar 2., Mon

Agents Team

↑ Daily Notes - 2026 March

Mar 31., Tue



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Calendar/Timeline coordination w...

↑ Mar 31., Tue

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent


- **Calendar/Timeline coordination with App Team** — Benedek is prototyping the timeline UI (events, attachments, Apple Reminders integration) and plans to write a summary for cross-team sync. Meanwhile, Peter Adam (App Team) is pausing other work to focus on Calendar planning after Balint clarified key requirements: must be fast, independent from current data, replace Google/Reminders, no sync. Christoph also plans to align on who contributes what — this sync is timely.
- **C4 release scope shift** — Christoph confirmed the team is focusing on quality-of-life assistant improvements for C4, with more emphasis on Calendar and less on assistant features. The minimal set includes: background mode, synced thread state, title generation, sidebar mode, per-tab state, BYOK, and collections hardening. Worth confirming priorities within this list.
- **Apple Foundation Models exploration** — Viktor asked the team to estimate complexity for multi-turn conversations and tool calling with Foundation Models. CX Team (Peter) also discussed this with Viktor, highlighting on-device/secure/private advantages. Zsombi's earlier [multi-turn POC](#) showed the 4000 token window works surprisingly well for many turns.

 Mar 30., Mon

Timeline

Continued with adding new features to the prototype. My focus today was to:

- Simplify calendar events to use mock data so I can generate a whole week with multiple calendars and to make them more easily prototype-able for editing, adding, deleting them.
- Add URLs as attachments. My approach is to use the design of card type smart links as the enlarged versions of link previews. The collapse/expand works the same way as with images (click the group to expand it into a focused interface, then you can add, open, and remove them)
- Added Apple Reminders as integrable tasks to the timeline view.
- Event configurator UI started, but I ended up only drafting its static version to see how it fits. I like it, so my plan is to integrate my earlier [Figma design](#) into the prototype.



 Mar 31., Tue

I added quite a few things, but they are not looking / acting perfectly. I'll focus on improving all the stuff I added:

- ☐ Event overlap issues
- ☐ Event interaction, and config panel UI improvements
- ☒ ~~Link attachment open and close UI iteration~~
- ☒ ~~Attachment exposé improvements (expand and collapse animation, and UI improvements)~~

I also plan to write a summary of what I'm planning and prototyping right now. Then we can sync on who will contribute to what in the coming weeks.

Christoph

📅 Mar 30., Mon

Based on this morning's discussion, we are focusing on quality-of-life improvements for the assistant for the C4 release, as all the additional features would need more explorations. So we put more focus in the calendar feature and less on the assistant.

This is the set of improvements I believe we should keep pushing for:

Assistant Improvements – Minimal Set

Chats Infrastructure

- **Full Background Mode** – Smooth background execution of threads, even when the app is closed
- **Synced Thread state** – when a thread has new message since the user looked at it this should be stored in the thread and synced across devices. Currently the badge count calculation is very unreliable and should be hardened
- **Thread title generation** – when the users sends the first message, we should generate a short title for the thread, instead of using the message directly

Assistant UI

- **Pinned Mode/Sidebar** – The argument for adding a permanent sidebar mode hasn't changed – with assistant becoming an integral part of working inside Craft, there should be a way to the assistant in your window more permanently and without blocking other panels or UI.

Threading Mode

- **Per-Tab Assistant State** **Experiment** – Right now, we have one global assistant state that is kept across threads and only changes when you switch between spaces. However from the way users have started using the assistant and the longer execution times in edit mode, I believe it is more common to have on thread going on one tab, then switch to another document in another tab and continue working with the assistant there. Therefore I believe a model of one assistant state/thread per tab is better fitting to how Craft and the assistant are used. This also matches the behavior of AI chats in other tab-focused tools (Comet, Dia).

Bring Your Own Key

- **User-Provided Keys** – Users want to use their own Anthropic subscription instead of using Craft credits.

Capabilities Improvements

- **Known Issues** – We have a few smaller [known issues](#) to be investigated
- **Collections Hardening** – All collections operations should be retested and improved, especially the new view modes and their new settings

Apple Foundation Models

Viktor asked us to look into improving the Apple Foundation Model integration. I believe we should do estimates for how complex it would be to test these ideas:

- **Multi-Turn Conversations** – Right now, we only allow one question and response per thread and force users to create a new thread for a second question. In November Zsombi created [a POC](#) for multi-turn conversations with Foundation Models, which seemed to have worked well ("The 4000 token window is actually not that small, you can have many turns even on longer documents.")



AFM multi-turn conversations — Enabling Multi-Turn Conversations — Session persiste...

- **Tool Calling** – Apple Foundation Models also support tool calling, it would be an interesting experiment to see if we can wrap our tools to support the `Tool` protocol (it seems pretty basic) to support some agentic behavior with local models, like search.



Expanding generation with tool calling | Apple Developer Documentation

- Idea: If the foundation models turn out to be good at adhering to simple instructions but not the complex system prompt, we could test creating special model+tool+instructions combinations for specific tasks, e.g. a simple prompt with a `list folder structure`, `create folder` and `move doc` for a "cleanup unsorted docs" action.

📅 Mar 31., Tue

- Align on tasks for the calendar/timeline feature, to see how I can best support Bene

Activities Yesterday

- **Release Preparation**

- Deploying and end to end testing the latest service version in sandbox
- Production readiness review and fixing findings and issues
- Merged release versions fixes into the new service
 - The edit folder validation required some adjustments with PI, so the skill loading is properly enforced
- Found and fixed some configuration issues with PI's grep and find tools
- Disabled automation REST APIs
- Fixed test and lint issues
- Security review and adding some vulnerability improvements
- Testing the AI gateway rate limit error with PI
 - The rate limit error was successfully propagated without multiple retries
 - Added propagation of the rate limit headers via the fetch interceptor throughout the PI event handlers, so we can send the remaining wait time to the client
 - Adjusted the client's error message to show the remaining wait time when its available
- AWS deployment caused various crashes and subprocess hanging
 - After some investigation managed to find the root cause, which was a bug in PI. One of PI's dependencies was imported from its code, but not properly declared as a dependency, which caused issues in the Docker build
- As the first step I switched my mac client to the release build and was focusing on end to end testing with that
 - This means new feature flags were off, the prompt generation is still client side, the ID mapping is still client side and the original client tools are integrated into the new tool server side tool wrappers

PLAN FOR TODAY

I will deploy the service to prod today, so we have time to follow up on anything that comes up, while I will continue end to end testing and strengthening the code pathes that are triggered when we use the new client with the service.

I will also prepare in the next 2 days before I go to holidays a proper runbook, and talk with the services team in order so they can fix any issues that might arise when I'm on holidays.

- Deploy new service to prod
- Testing and preparaing the new client's code pathes with the deployed service



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Gyula is out of office Monday an...

↑ Mar 30., Mon

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

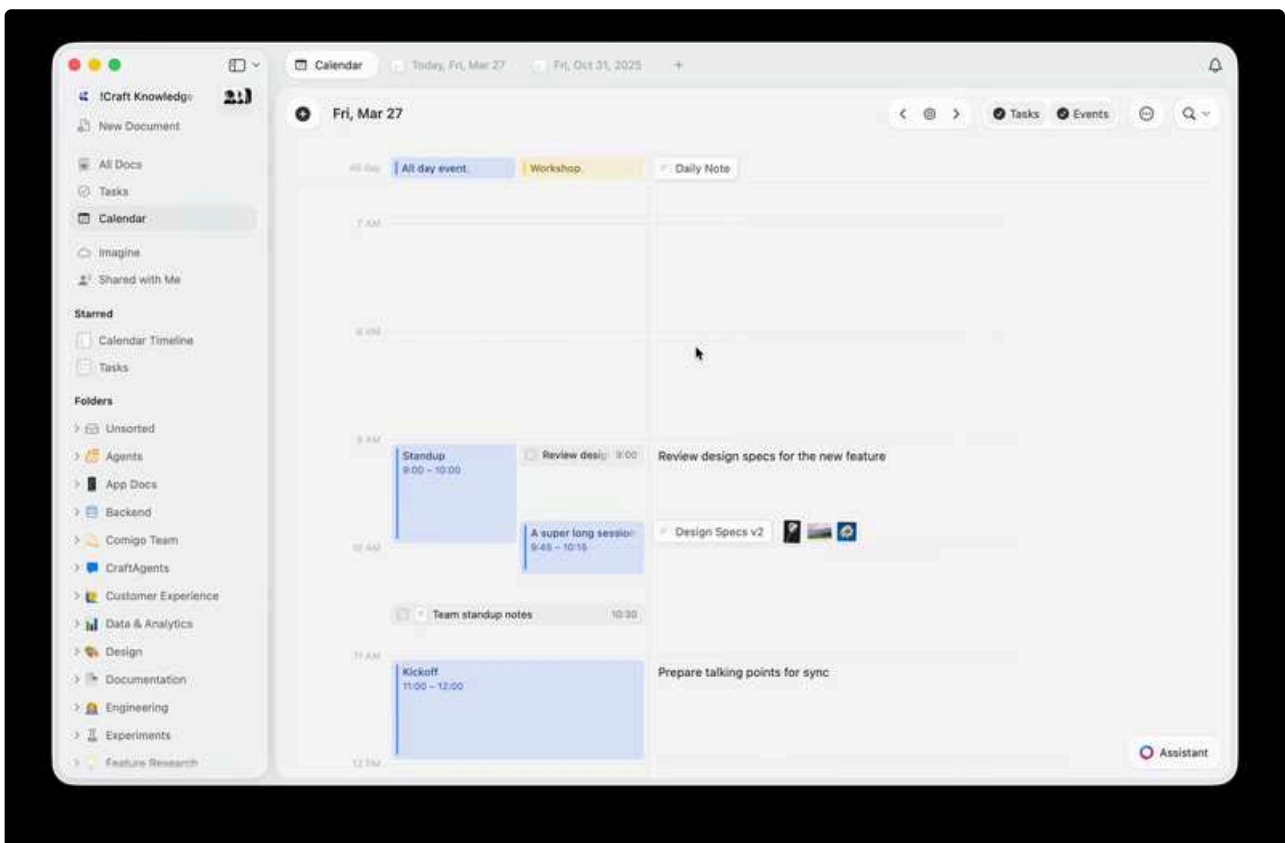
- **Gyula is out of office Monday and Tuesday** (vacation, Tuesday depends on return flight). Plan around his availability for the Cloud relay and search rework items.
- **Zsolt is planning to create ai-agent-v3 service** — the service stack has gone through "massive changes" since the current release (agent harness swap, service-side ID mapping, server-side tools, threading, automations, postgres dependency). He feels this is the right time to clone to v3 for safe deployment and testing. Discuss scope and timeline for the v3 switchover.
- **Zsolt's automation work spans 3 standalone services** (ai-agent-v2, ai-agent-automation, ai-tools) plus a new client with feature flags. Increasing complexity worth discussing from an operational perspective — deployment ordering, testing strategy.
- **Christoph's timeline check** — targeting early May release with 4.5 weeks left. This week (W14) is focused on deciding information architecture, mobile, and full Automations coverage. Good moment to validate priorities and confirm the plan is on track.
- **Cross-team: Marton (Services) wants to sync with Levy about MCP v2 timeline** and create a beta scope+checklist, plus give Anthropic a heads-up. Relevant context for the Agents team's API/MCP work.

 Mar 27., Fri

My whole day was spent working on the prototype. It now supports Event preview, Task config + completion, and image and linked doc attachments.

Interactivity with Doc links and inside note taking is limited, the only things working are the ones shown in the video below.

- One new feature is image attachment handling. Changes are that other attachments could have this same scale up centre stage interaction when the group is clicked.
- Another change I made when doing the prototype is move the attachments row to the bottom of the hour cell. This way we can have uninterrupted space to display an "Add" button users can use to quick add images, links, liked docs, etc.
- The current day and daily navigation was sent up to the navigation for now. I kinda like it and it mirrors how we use the glass buttons for nav purposes in other pages, and this also saves space for the pages.



 Mar 30., Mon

Continue adding features to the proto, focusing on the quick notes feature:

- Formatting text.
- Editing the quick note should elevate the cell into a platform that can stretch to accept more text.
- Doc link add, edit, delete flow
- Improve image add, edit, delete flow both in terms of UI and cover more cases (how would 20 images open for example)

Cloud relay groundwork, cleaner branch boundaries, and Docker deployment cleanup

Planning & Architecture

- Defined the Craft Agents Cloud direction as a self-contained relay service: Cloudflare Worker + Durable Object tunnel, signed visitor identity tokens, auth, billing, and an app-facing contract.
- Reworked the rollout to keep cloud infra isolated from the bugfix branch: created a dedicated `subscription-tunnel-infra` worktree/branch and removed the relay infra commit from `bugfix/smallfixes`.

Release & Features

- Implemented a pluggable settings extension API in Electron so optional packages can register settings pages, components, icons, and channels at runtime without leaking cloud-specific logic into core files.
- Added the Cloud settings integration/package and pushed the cloud work to `origin/subscription-tunnel-infra`, alongside the relay infra work.

Code Review & Bug Investigation

- Re-evaluated the in-chat search fix around CSS Custom Highlight API + cross-node matching, with remote `rg` fallback and a proper "search unavailable" state instead of silent zero-result failures.
- Investigated Docker/WebUI startup failure modes from the user environment side, not just app logic, and tightened the explanation of where the failures can come from.

Docs & Developer Experience

- Added Docker troubleshooting guidance for writable `HOME` requirements and WebUI asset/startup issues in the headless server docs.
- Committed and pushed the Docker docs update on `bugfix/smallfixes`.

. . .

- Validate the Cloud relay + settings flow end-to-end and lock down the OSS/private boundary.
- Land the search rework after testing CSS Highlight API behavior on large, code-heavy sessions.
- Tighten Docker/WebUI deployment docs based on real failure modes.
- Out of office on Monday and Tuesday due to vacation; Tuesday office presence depends on the return flight.

Christoph Locher

At the beginning of the week I stated the goal to have a full UI overview including the mobile view. I moved away from this goal in order to take more time to explore various variants of the information architecture/navigation as this required more consideration. After the end of last week felt very fast, getting to the prototype with lots of details I enjoyed exploring, this week felt more full of friction, just because everything took a bit longer than expected. Nothing major, just more friction. So today, I took a step back to reflect and review the current state:



Assistant Project Overview Fri, Mar 27

Information Architecture · The biggest challenge right now is figuring out the information architect...

 Last Edited 3/27/2026

Targeting a release early in May, there are 4.5 weeks for me left (I am OoO from 30.04). For the next weeks, I plan these focus areas:

- W14 (30.03 – 4.04) 4 days : Decide on information architecture, mobile, full-coverage for Automations
- W15 (7.04 – 10.04) 4 days : full-coverage for Connections, Panel behavior (sidebar), Thread behavior (thread per tab?)
- W16 (13.04 – 17.04) Week Together : full-coverage for Memory, Thread View (Message format & turn group)
- W17 (20.04 – 24.04): Testing & Buffer

 Mar 27., Fri

Assistant

- **Automations**

- Testing and improving automation end to end across **ai-agent**, **ai-agent-automation** and **ai-tools** services
- Iteration on DB schema
- Fixed scope related issues with spaceld and teamld
- Setup unified API access for ai-agent
- Added proper authorization based on selected space/team
- Updated thread handling, so each automation run creates its own remote thread
- Fixed message schema issues causing user messages to not appear in the history

- **Thread handling**

- Introduced server side thread read/unread state
- Automation runs automatically become unread
- Runs from the client where the thread is moved to background is made unread by the client via a new socket action
- The client manages read state based on UI thread changes
- The server side thread state is synced to the client when the panel is opened
 - Still have client side unread tracking and hybrid operation, because we don't have real time sync of the thread state from the server

- **Service stack in production**

- The service stack is getting more complex now having 3 standalone services (ai-agent-v2, ai-agent-automation and ai-tools) and a new client again with some new feature flags (service side ID mapping)
- The service went through some massive changes since the currently released version
 - full agent harness swap to PI with abstraction layer
 - service side id mapping
 - introduction of server side tools with hybrid tool operation and headless mode
 - significant refactor of all tooling behavior
 - threading changes
 - background mode
 - automations with scheduling
 - core dependency on postgres and unified API
- I feel this is the right time where we have to make the switch to a v3 service in order to safely deploy and test this internally
- I started to make the plans for cloning the service and rewire the stack to work with v3

 Mar 30., Mon

Weekly Goal

- Continue work on automation support and make thread handling more robust
- Create ai-agent-v3 service and rewire the service stack + client to work with the v3 service
- Deploy all 3 services to production for end to end testing and fixing issues

↑ Daily Notes - 2026 March

Mar 27., Fri



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Benedek's 1.0 roadmap has most...

↑ Mar 27., Fri

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Benedek's 1.0 roadmap has most items ~80% done but unchecked** — he's shifting to a sequential end-to-end pass to surface real progress. Worth reviewing the [roadmap](#) together to align on what's truly blocking 1.0 vs nice-to-have.
- **Benedek raised a Reminders vs Tasks distinction topic** — he collected his thoughts in a [writeup](#). This is an open design question that could affect the product model. Worth discussing alignment.
- **Christoph is comparing three assistant UI/IA variants today** — pinned sidebar with 2-level nav, top-level assistant page with swappable sidebar, and the previous approach. He plans a writeup with implementation requirements. Good moment to align on direction before more prototype work.
- **Tamas Flamich (Services) found that Craft Agents doesn't support OAuth flow for API sources** — only MCP sources are supported. He's pivoting to bearer tokens for now, but this may be worth discussing as a future capability.

 Benedek Gobolos

 Mar 26., Thu

Reminders vs Tasks

Collected my thoughts on the Reminders vs Tasks topic I raised yesterday

 CRAFT

Reminders vs tasks distinction

Roadmap

I got a bit confused about what I'd need to finish for a good 1.0, so I collected them super thoroughly. This list involves all the smaller / larger stuff I need to get done. As you can see most checkboxes are unchecked, but most things are about 80% done. What I'll do is start focusing on covering everything from one end to the other, this way we'll see the actual progress from now.


 CRAFT

Project Roadmap

Prototype

I got inspiration from Christoph's time working on his prototype right insid the app. For all the integration and interaction topics to be achivable I wanted a prototype that is using the codebase with all the bells and whistles. It's not in a good enough state, but I'm progressing fine.

I'm mostly having challenges setting up the layout, but after that is done it will be easier adding isolated features.

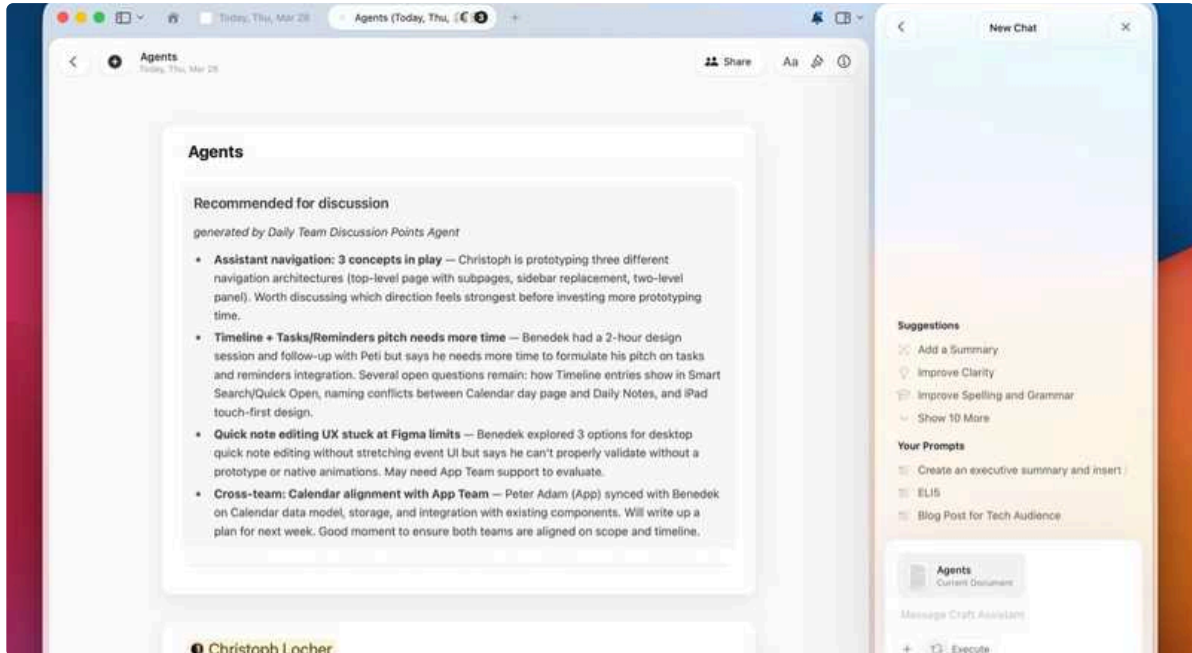


- Continue with setting up the prototype. I'll inherently arrive at points where I'll figure out interaction on macOS, iPadOS, and iOS.

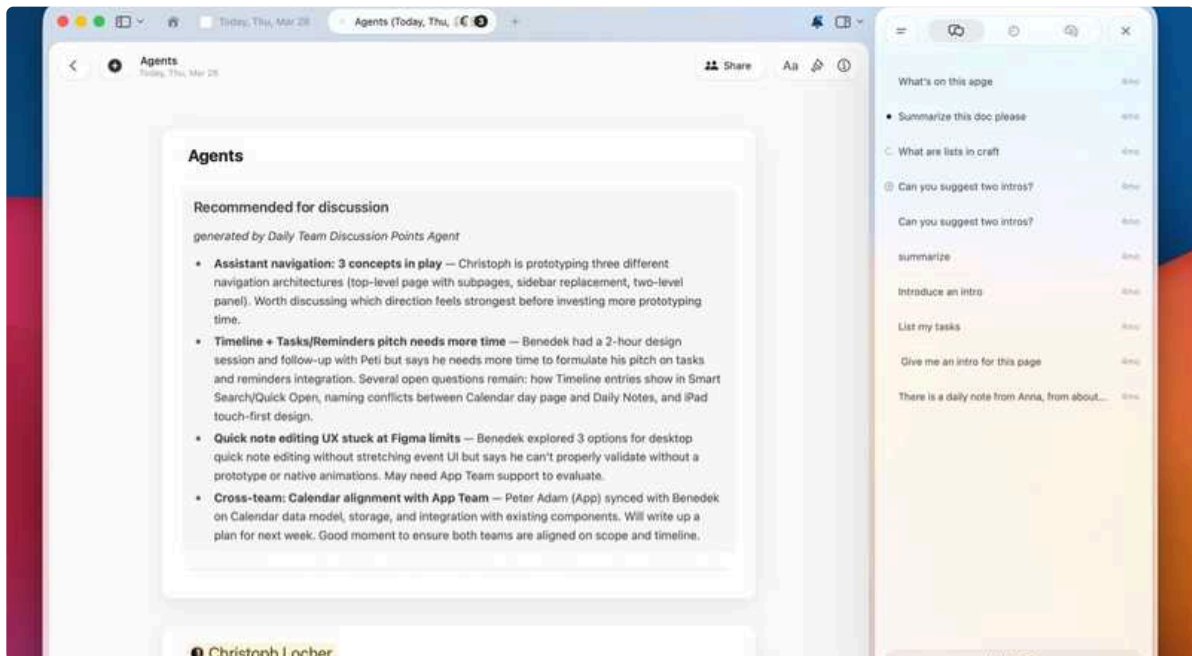
My plan is after I arrive at the same "feature parity" as where I am in Figma, then I'll be able to focus on still undecided topics. I'm happy to pivot from this plan, but my approach is that I'll have a more realistic foundation where I can figure out interaction and navigation (like how quick add feels, how UI size and gestures feel, what is missing from our current plan, etc).

Christoph Locher

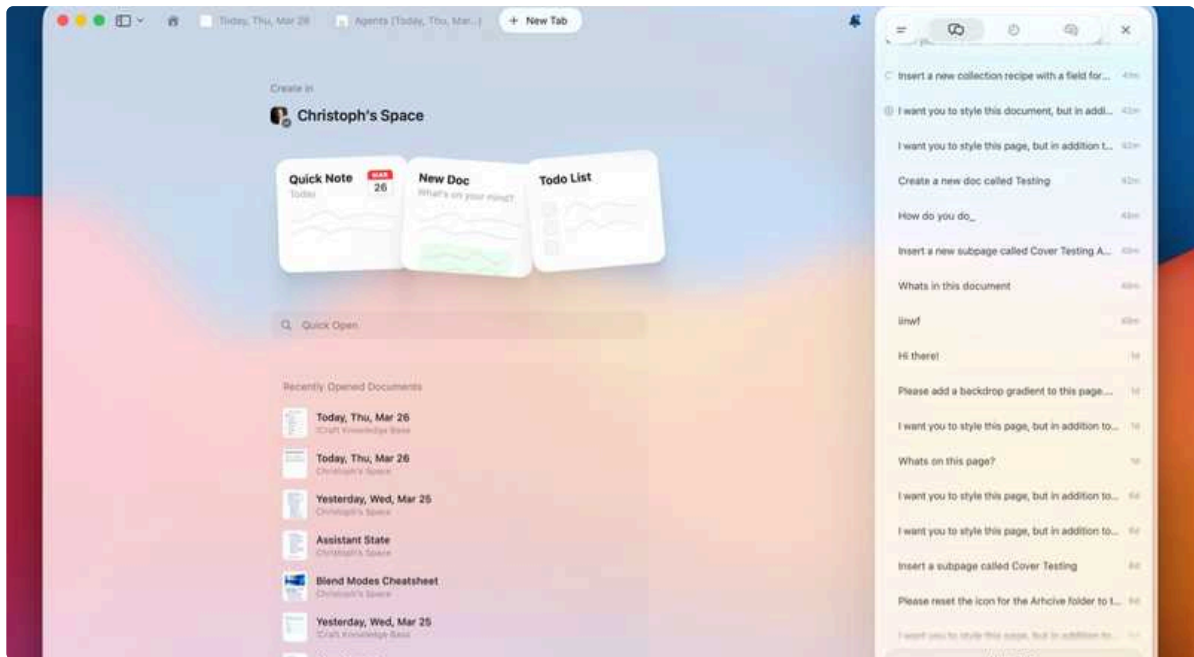
- Investigated a user issue with complex prompt
- Tested different assistant UI setups
 - Brought back the pinned sidebar, with 2-level navigation inside the sidebar



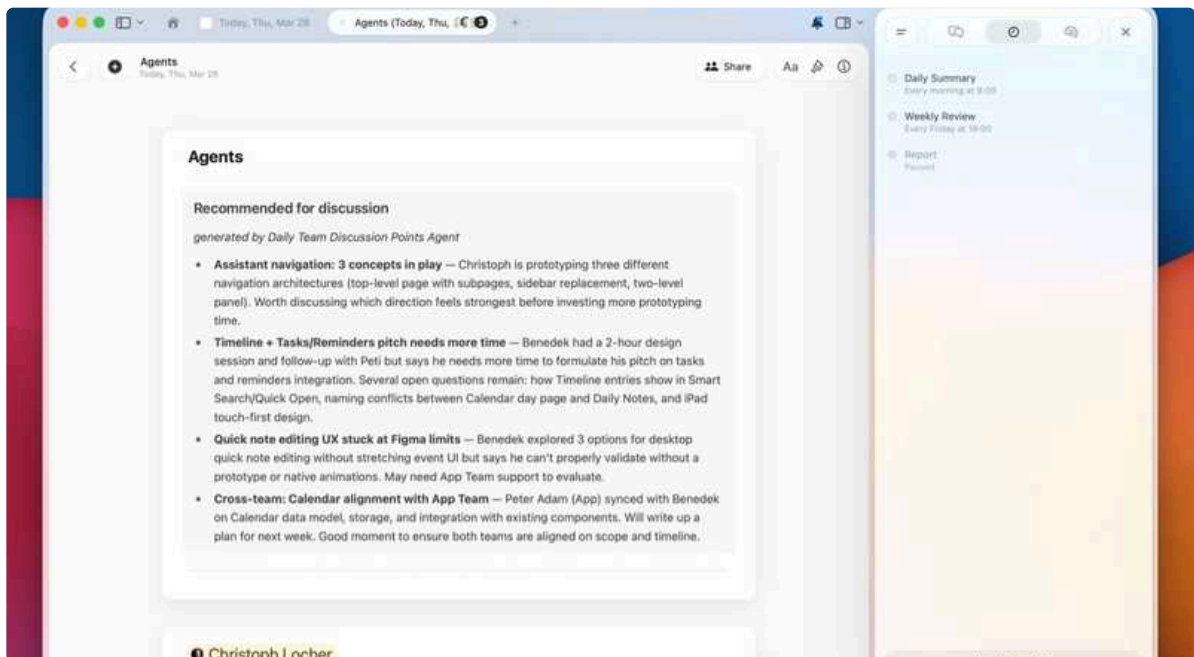
Chat view – now with a back button at the top



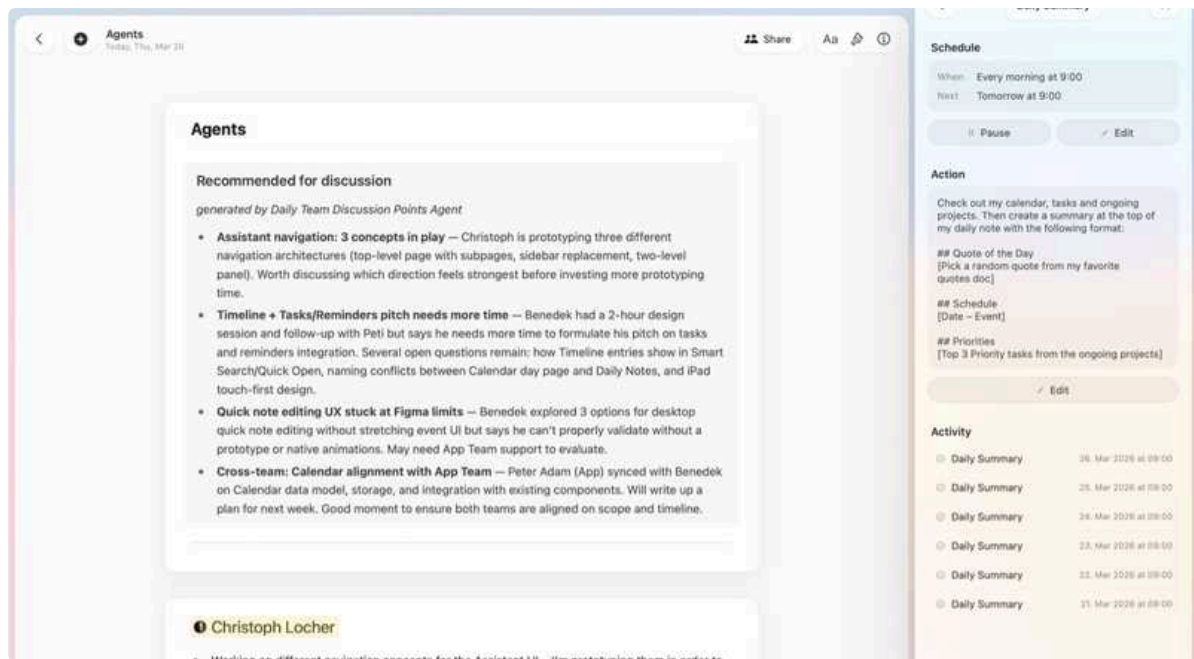
Top level – Segmented Control to select page, the chats page shows all threads with the different states



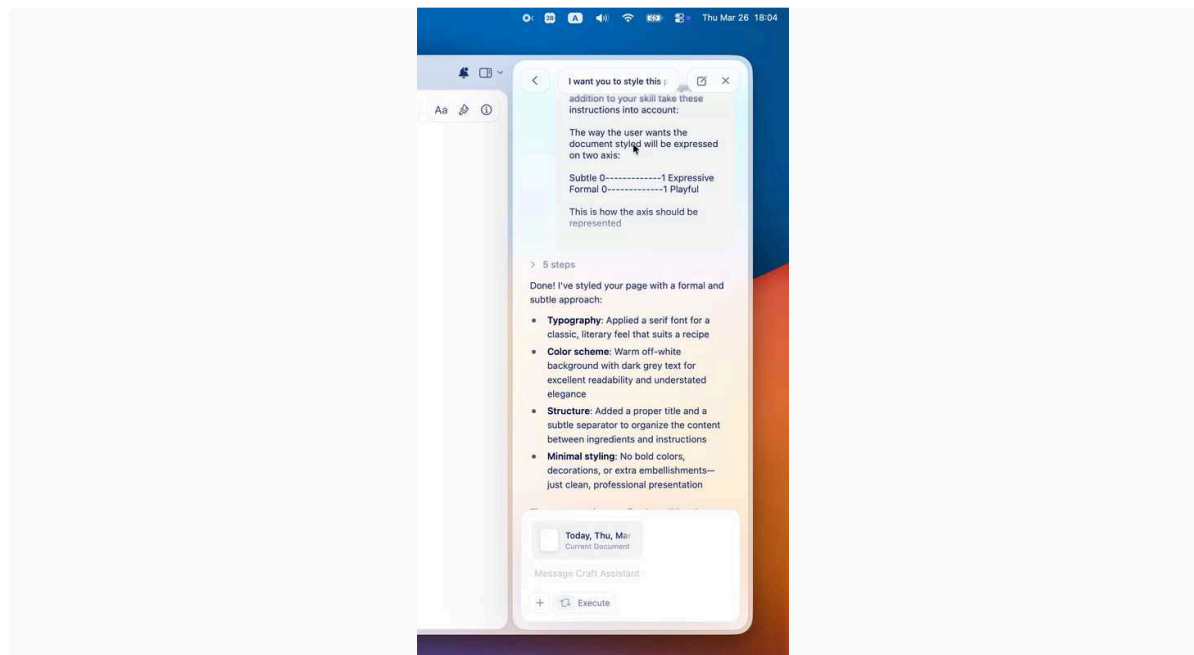
In a different space with a long thread list



Top level – Automation page

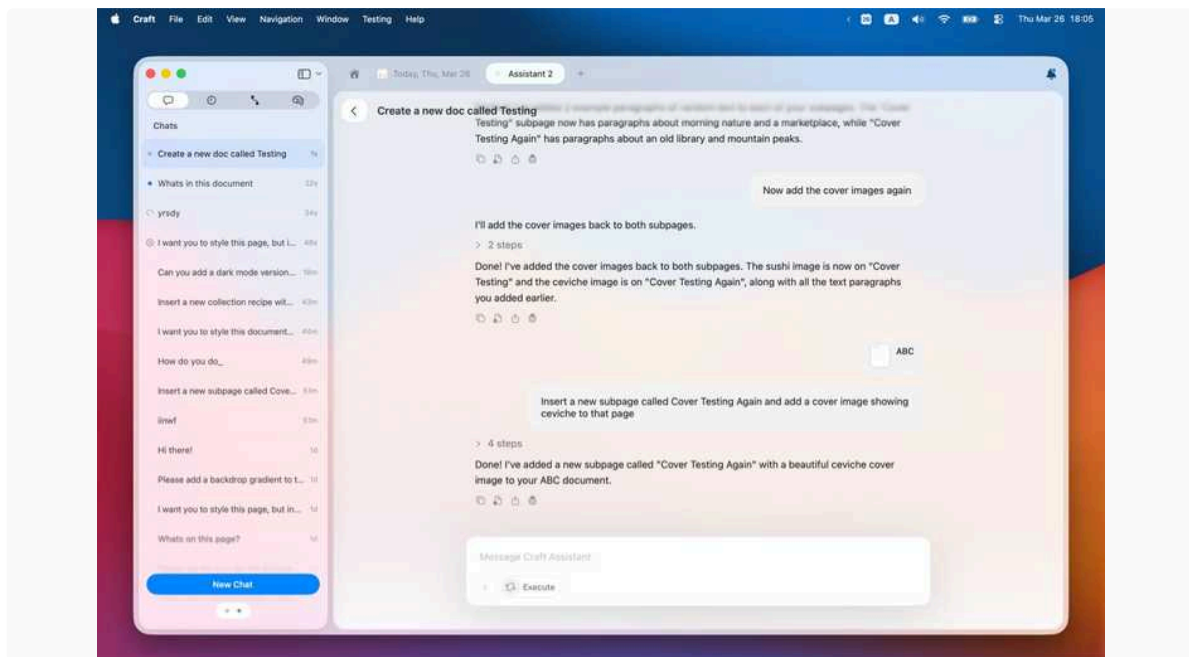


Sub level – automation detail

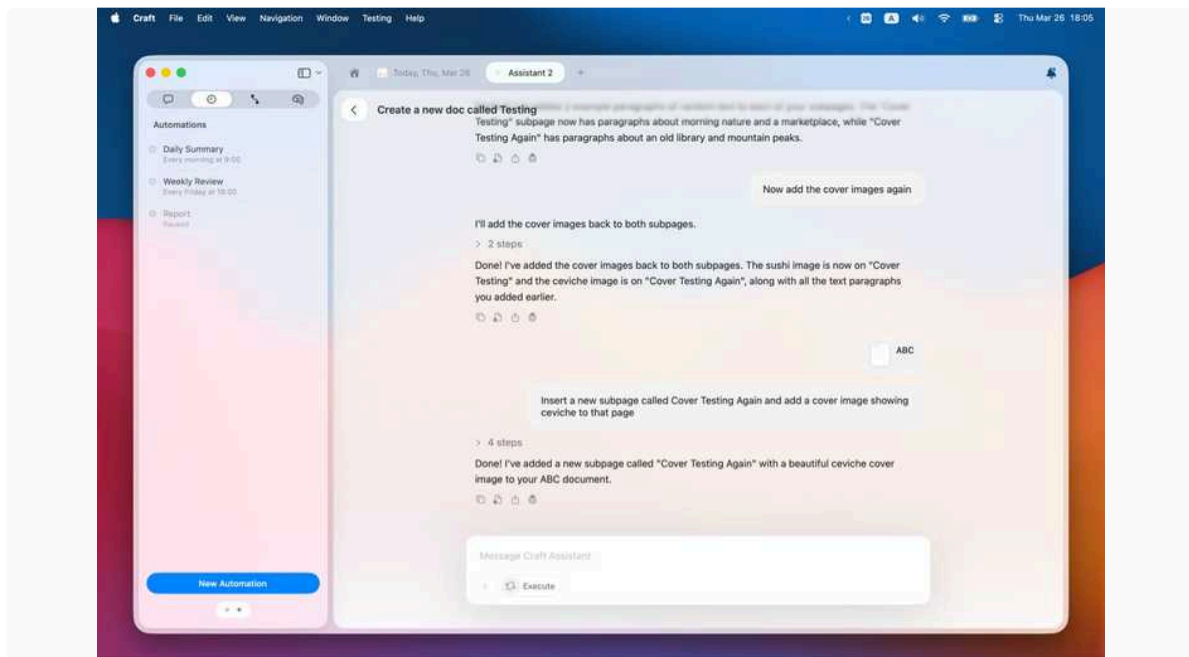


Video of the navigation flow

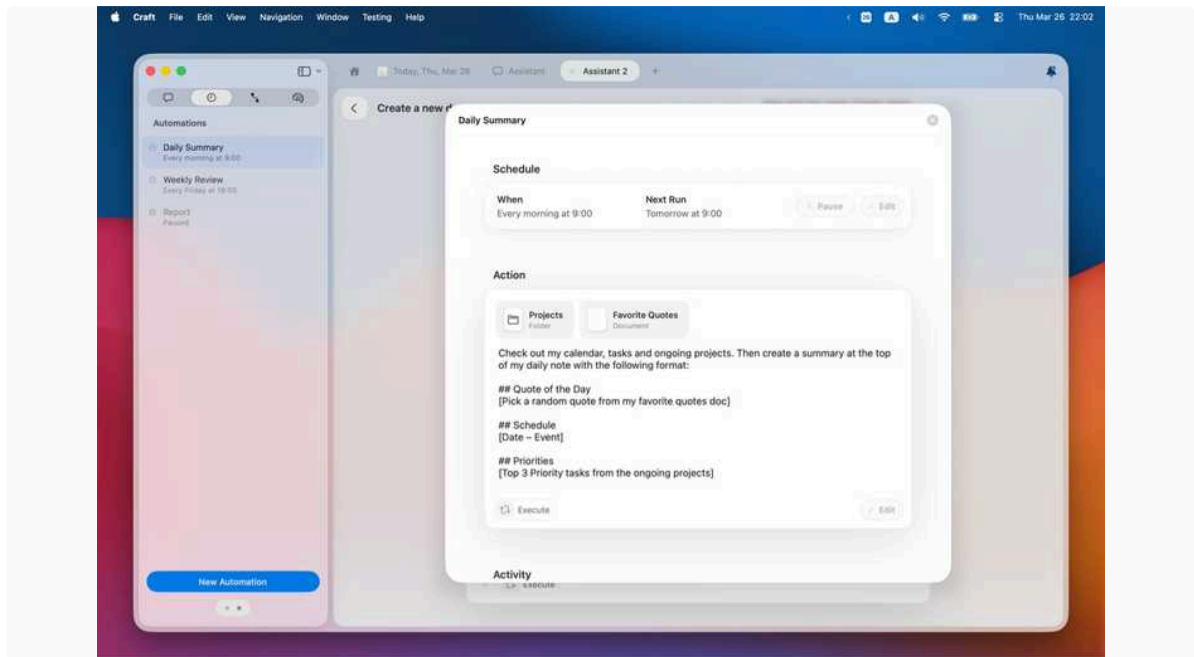
- Top level page, accessed from an assistant item in the main sidebar. Like in the document, this page displays a special left sidebar (which can be swiped back to the folder one)



The standard view of the Assistant page: Chat sidebar on the left, thread detail on the right



The automation view only switches the sidebar, the main view stays on the chat. Automations open as a modal



Tomorrow

- Reflect and compare the three IA variants
- Create a writeup of the variants and the plans for the assistant/requirements for implementation

v0.8.1 shipped — 15 fixes, Docker Compose, and a 100k error flood squashed

Release v0.8.1

- Subprocess error flood fix — `uncaughtException` handler was looping on broken stdout, flooding sessions with ~100k error cards
- Web UI thumbnails — `thumbnail://` protocol fallback for browsers via `readFileDataUrl` RPC
- Project-level skills resolution — `SkillInfoPage` now receives `workingDirectory`
- Remote workspace stability — fixed splash hang, reconnect banner flicker, and added reconnect UX for disconnected workspaces
- Claude OAuth routing — fixed for remote workspaces (channels were LOCAL_ONLY but verification ran remotely)
- Headless server OAuth fix from tester feedback
- Docker Compose for easier server deployment + GHCR image support
- Artifact naming aligned to `Craft-Agents` (plural) across build, workflows, and release scripts
- Unified label ordering across all UI label surfaces
- Image validation, startup error classification, and cross-platform sharp binaries

Infrastructure & Testing

- LiteLLM supply chain vulnerability audit — all clear
- Tested remote Linux server deployment extensively

Research

- Investigated computer use capabilities for Craft Agents

Next steps

- Spike: share workspace resources (sources, skills, automations) with remote workspaces for easier local-to-remote migration
- Continue infrastructure spike for user registration

 Mar 26., Thu

Assistant

- **Automations**
 - **System Prompt Handling**
 - Created tests for server side system prompt generation
 - **AI Agent Automation Service**
 - Created new lambda based service for scheduler handling
 - Created automation and automation run schema in Postgres DB2
 - The schedule schema is similar to our task schedule options (monthly, weekly, daily) with the extension of hour schedule as well
 - The scheduler service runs periodically and looks for due automation schedules
 - Due automations are queued and executed in SQS
 - The job calls the agent-v2 service via the internal `generate` REST API
 - Created `completion` endpoint, that the agent-v2 service can use to report when execution is finished
 - **AI Agent V2**
 - Created automation management REST APIs to create, update, delete automations
 - Create internal generate REST API, that can be used by the automation lambda
 - For automation execution we cant use expiring user tokens, so switched to API key based authentication
 - We also call the AI gateway with API key for these scenarios
 - After the automation execution the `completion` endpoint is used to notify the lambda about the job completion
 - Added local infra to test the services with the scheduler in a non lambda dev environment

Other

- Investigated user report about "wasted credits"

 Mar 27., Fri

- Continue stability improvements + background session execution / automations support

↑ Daily Notes - 2026 March

Mar 26., Thu



Agents

Recommended for discussion · generated by Daily Team Discussion Points Agent · Assistant navigation: 3 concepts...

↑ Mar 26., Thu

Agents

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Assistant navigation: 3 concepts in play** — Christoph is prototyping three different navigation architectures (top-level page with subpages, sidebar replacement, two-level panel). Worth discussing which direction feels strongest before investing more prototyping time.
- **Timeline + Tasks/Reminders pitch needs more time** — Benedek had a 2-hour design session and follow-up with Peti but says he needs more time to formulate his pitch on tasks and reminders integration. Several open questions remain: how Timeline entries show in Smart Search/Quick Open, naming conflicts between Calendar day page and Daily Notes, and iPad touch-first design.
- **Quick note editing UX stuck at Figma limits** — Benedek explored 3 options for desktop quick note editing without stretching event UI but says he can't properly validate without a prototype or native animations. May need App Team support to evaluate.
- **Cross-team: Calendar alignment with App Team** — Peter Adam (App) synced with Benedek on Calendar data model, storage, and integration with existing components. Will write up a plan for next week. Good moment to ensure both teams are aligned on scope and timeline.

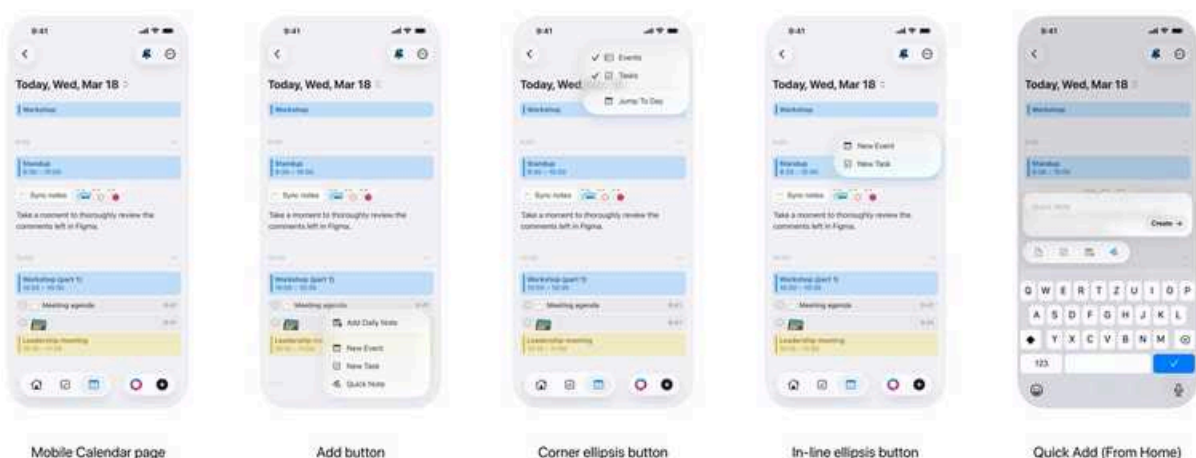
- Working on different navigation concepts for the Assistant UI – I’m prototyping them in order to be able to compare how they feel in practice
 - Top Level Page “Assistant” with subpages (Chats, Automations, Connections, Memory), each with its own subnavigation. Assistant Panel for Chat inside documents
 - “Assistant” item in the main sidebar, which opens an Assistant page. In this page there is the chat/thread view as the main view, and the left sidebar is replaced with the Assistant sidebar, which has Chat, Automation, Connections and Memory subpages within. (Similar to how entering a daily note changes the sidebar to the calendar sidebar)
 - Two-Level Navigation in the Assistant Panel. Top level with pages for Chats, Automations, Connections, Memory. Clicking on a thread in the chats view opens the thread in the sub level. When opening inside a document, it opens in the thread view – you can navigate back out to go back to the top level
- Long design session and follow-up discussions, talking about the Assistant information architecture and the Timeline + Tasks/Reminders ideas

Meetings & discussions

- 2-hour design alignment session
 - We went through Christoph's Assistant progress
 - I showed the Timeline and we had a long discussion on how Tasks and Reminders could be integrated to the Calendar while solving some problems related to them. I'll write an extensive summary of my plan, but I'll need more time to figure out the pitch.
- Integration talk with Peti
 - We went through all corner cases he could think of. The most important discussions were about:
 - Using Comments editor for quick notes with its simpler formatting options (bold, italic, strikethrough, link, code only for now)
 - Needing to figure out how the new Timeline entries will show up on Smart Search, Quick Open, and finding different names for Calendar day page, and Daily notes (as both makes sense to just name the date like Wed, Mar 18).
 - iPad support. It's layout will look like macOS, but it should be designed with touch in mind.
 - Preparing for Assistant integration with context on quick notes, attachments, events, etc.
 - Collaboration. We decided that all actions are shared except calendar events and reminders. Tasks and quick notes will be shared.
 - Extension support for quick notes(1.1): Share sheet action, Lock screen complications, Control centre action, etc.

Design

Mobile view

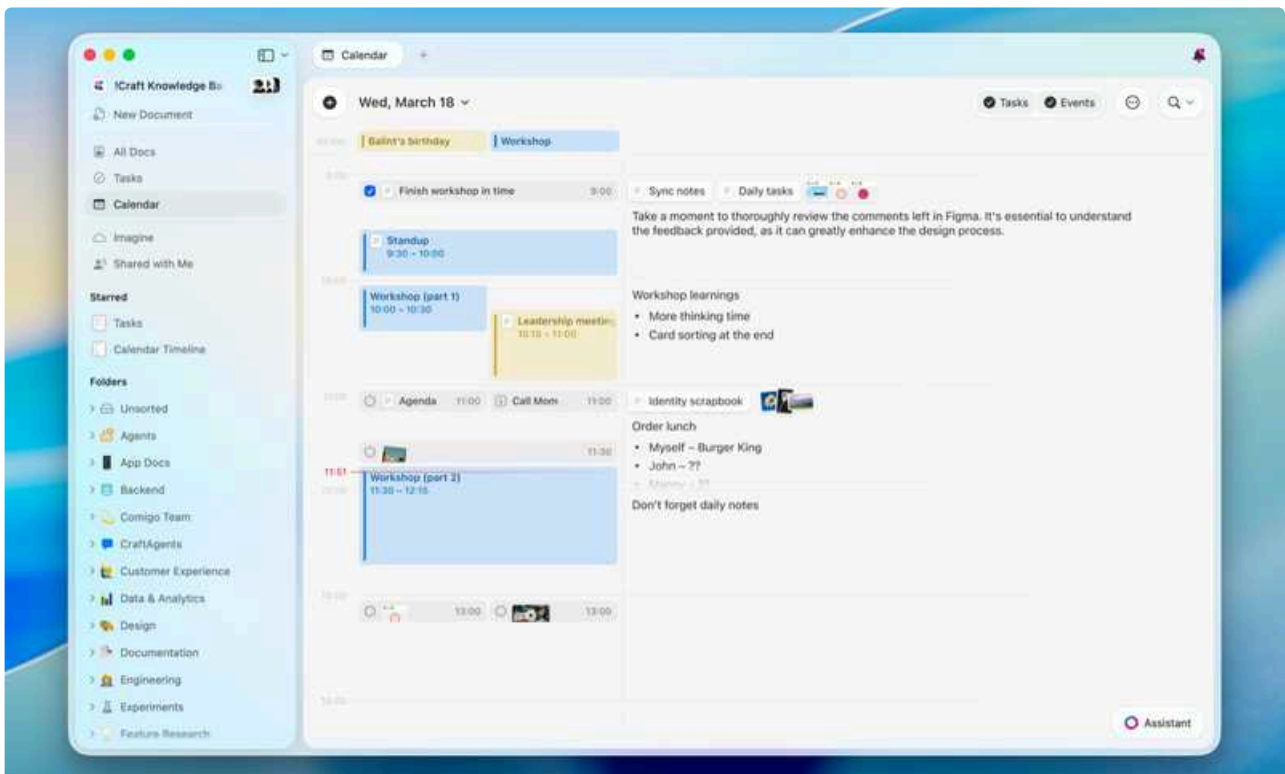


I started updating the iOS / iPadOS navigation items to support the new Calendar. This way I have an idea on how we'll add new events, tasks, or quick notes from the UI.

I'm still designing how editing and deletion will work.

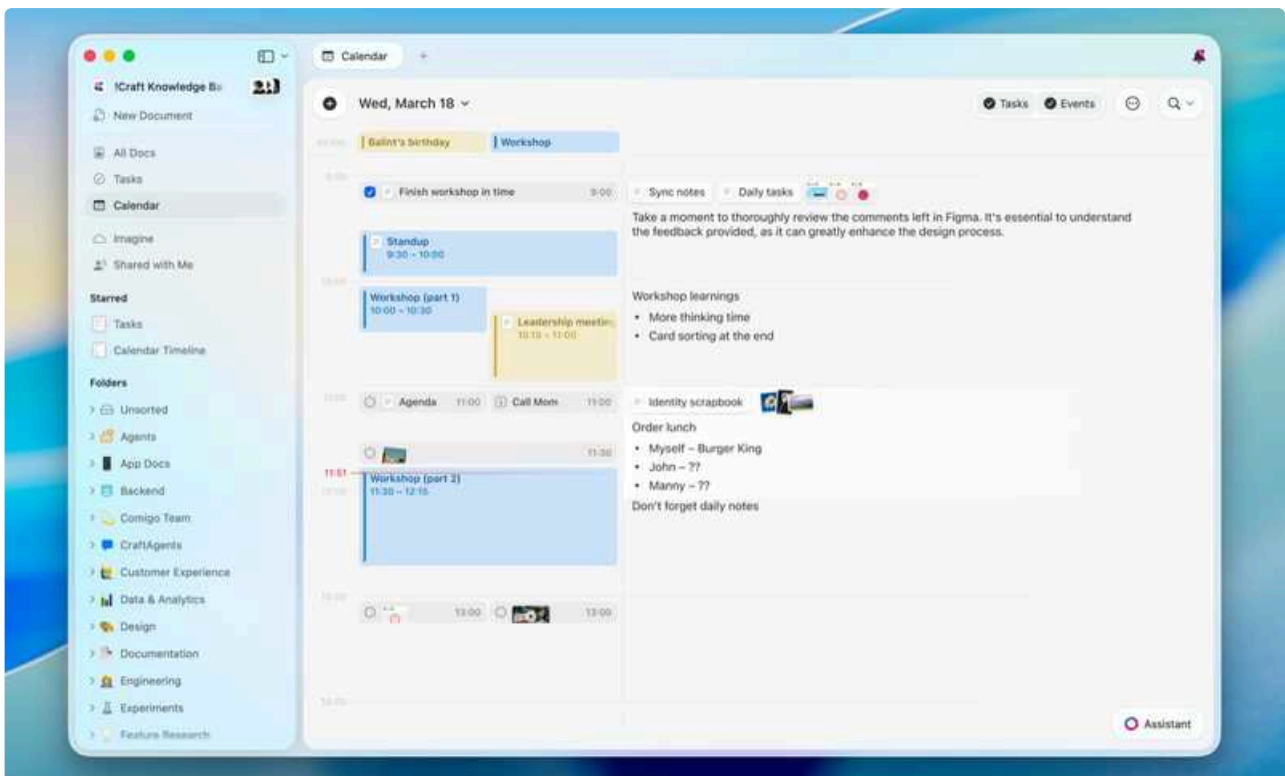
Desktop stretch

I want to figure out how we can keep quick note editing without stretching the UI of events. I have yet to nail it, but here are my attempts:

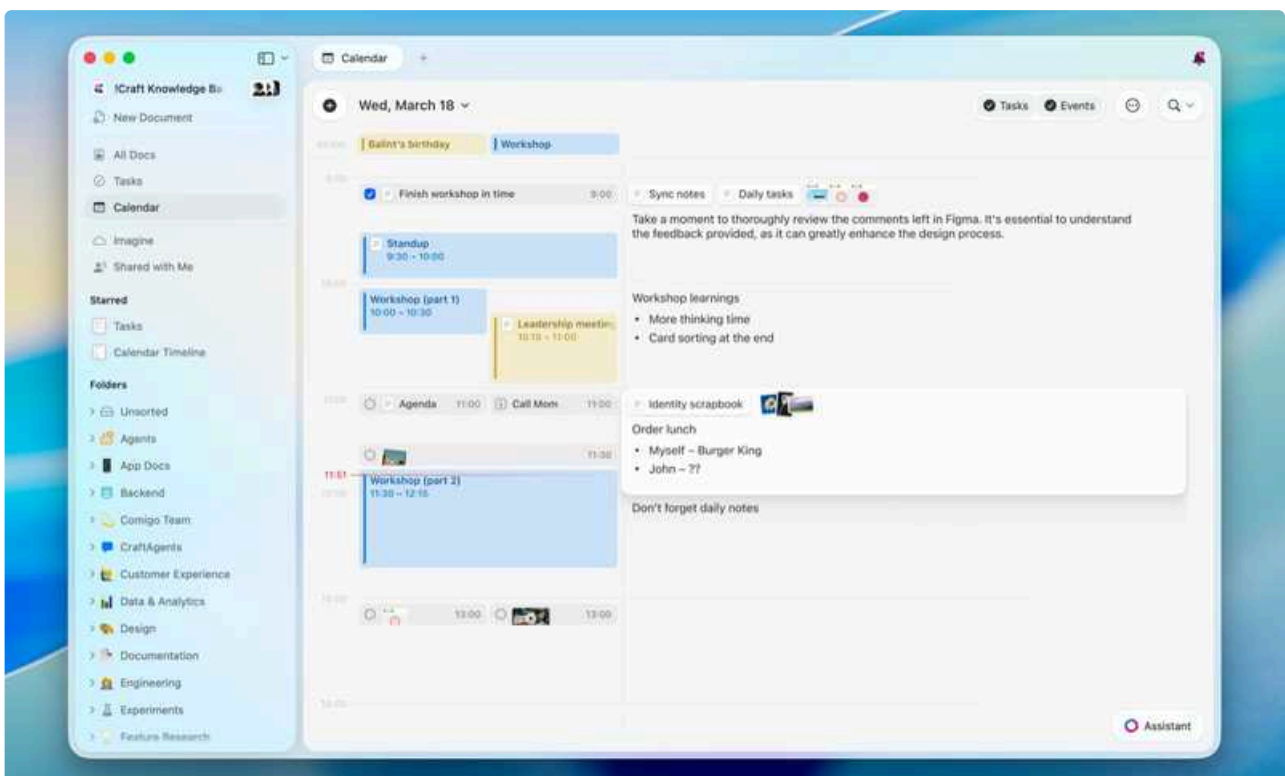


By default we could hide excess notes and only make it visible once user clicks into it.

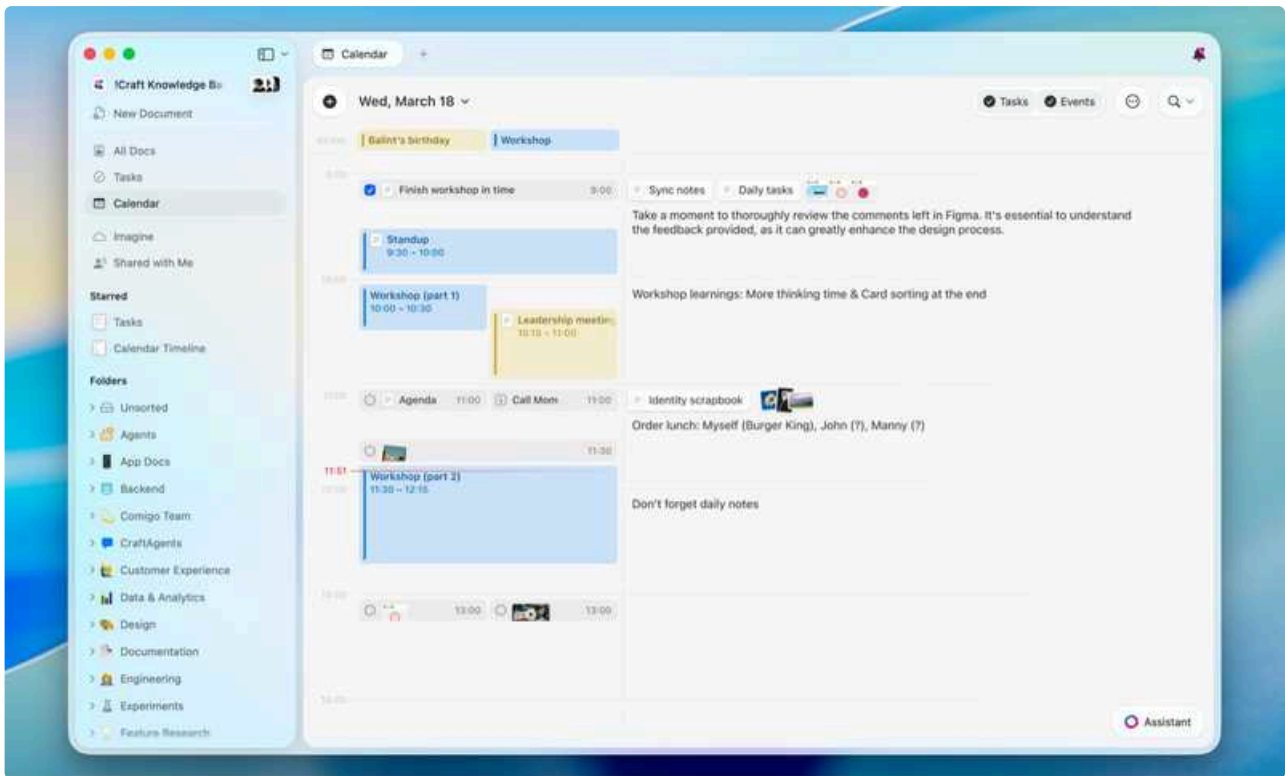
Option 1 could be to increase the size of the cell and highlight it with a lighter bg.



Option 2 could be to elevate the cell into a platformed popup with infinite writing space.



Option 3 is to disable line breaks and incentivise users to write longer rows.



I can't properly validate these without a prototype at least (or the animations we have on native), so I'll not go further with Figma.

Mar 26., Thu

- Formulate my thoughts on tasks and reminders, then write a pitch
- Figure out interaction of events, tasks, and attachments both on mac and iOS
- Find a good looking way to show current time on mobile and some way to highlight the current hour on all platforms



Released 0.8.0

- 0.8.0 is out

Next steps

- Hotfix release because new code path was not added to the allow list so webui is not in the oss repo
- Continue with the infra

Assistant

- **Tool Feature Parity + Tests**

- Fixed some tool schema disparity issues
- Calendar tests broke after yesterday due to the date section grouping changing, changed the tests to enforce stable grouping
- Align block_move error handling behavior between platforms
- Align folder description handling in tree view
- Aligned supported task list filtering options
- Aligned user info credits counting behavior

- **Background session execution**

- **Sticky sessions**
 - Implemented sticky session on the load balancer
 - Adding the threadId to the session cookie
 - The load balancer uses the value to route to a mapped instance
 - When the client changes thread it has to reconnect the socket to allow the sticky header to change and get a new mapping
 - Logging sticky mapping state for end to end verification

- **Automations**

- Started creating the foundations for automations
- **Headless execution**
 - Created REST endpoint to start a generation session without a client
 - In this mode we always rely on the server tools
 - Enforce thread caches to be tied to the userId + threadId
 - Also added REST endpoint to fetch a running session, that can now be polled to read incoming messages
 - Later on some kind of notifications can be pushed to the client for real time updates on automated sessions
- **System Prompt Handling**
 - Decided to refactor the entire sytem prompt handling in order to supported the future capabilities
 - The current system prompt generation happens entirely on the client side, which became a significant limitation
 - In order to run headless sessions the server needs to be able create the necessary base prompts with potential extensions, like highlighting the non interactive nature of the session
 - Currently a proper session cannot be started without the client as the system prompt is missing

- Alternative option could be, that the client precreates the system prompt for a scheduled automation when the automation is created, but this approach is not sustainable long term, because the prompt has many dynamic parts, that potentially need to change over time (eg current date, embedded document titles, user details, or just the static prompt itself would go through iterations of new versions, that precreating during creation would break)
- Centralizing the complex prompt creation on the server is also better to reduce complexity on multiple clients (web later) and reduces attack surface of prompt injection by removing the client side full control on the system part of the prompt, where its much easier to instruct the LLM to try to break out of the session sandbox
- I moved the prompt generation to the server by extracting all the variables that in some way alter the final prompt and requirinig the client to pass these configuration variables in generate or later for automations
- For now Im not refreshing document and userinfo on the server side, but it can be improved later on with direct unified API access in the agent service

 Mar 26., Thu

- Continue stability improvements + background session execution / automations support

↑ Daily Notes - 2026 March

Mar 25., Wed



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Tasks vs Reminders differentiatio...

↑ Mar 25., Wed

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Tasks vs Reminders differentiation in Timeline** — Benedek is unsure whether/how to visually differentiate task-reminders from non-task-reminders in the events timeline. He flagged this needs ideation in today's design discussion with Christoph and Lisa.
- **Timeline integration sync with App Team** — Benedek starts morning discussions with Peti (App Team) on integration of Timeline features (tasks, reminders, hourly quick notes). This is a cross-team dependency — ensure alignment on data model and block-level interactions.

- **Gyula half-day Thursday** — Gyula will be offsite Thursday afternoon. Remote setup core functionality is wrapped up but still has rough edges. If any review or feedback is needed before polish, today is the day.
 - **Infrastructure direction for remote agents** — Gyula evaluated CF Worker + per-user Durable Objects as a hosting approach with embedded client in the app. Also noted Computer Use is available but slow. Worth a brief share-out on findings and next steps.
-

 Mar 24., Tue

Figma

- In the morning I was focused on the Kanban Collection. Lisa collected a ton of feedback from her time using it and we had a session figuring out what we can do to improve the experience. I prepared for a while for it but the session was successful.
- I then had a meeting with Christoph about the Timeline. We discussed two important areas: tasks + reminders and the mobile layout.

Tasks + Reminders

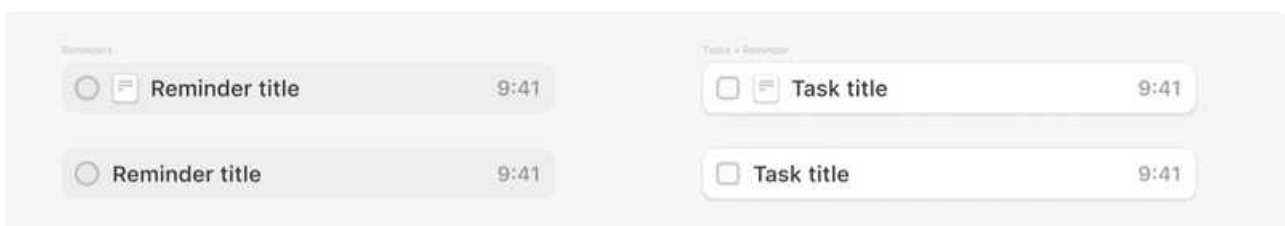
Christoph helped me clear up the confusion I had about the Tasks integration:

- Tasks can get a schedule and a deadline, both date only.
- Reminders are modifiers on any block that send a notification about it at a given date + time. They are not connected to tasks, so non-task blocks could be reminders.

This way:

- Tasks will only appear in a daily list of tasks scheduled for the given day or overdue ones that are yet to complete. In short, they won't appear on the timeline.
- Reminders will be added to the events timeline for precise positioning.

The only question: How could I show if a reminder is a task or not a task? Do I need to differentiate them in the first place or I'm looking at it too much from the system perspective?



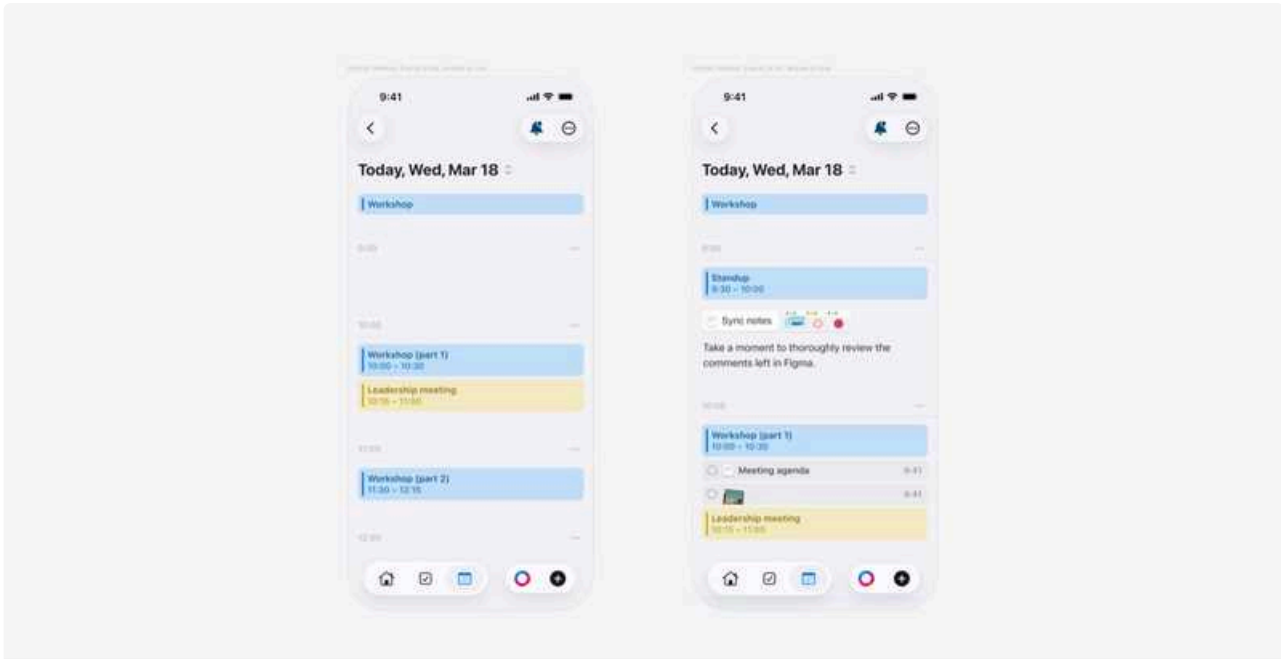
I'm not sure yet and I'll need some ideation in tomorrow's design discussion.

Mobile view

I've arrived at a solid iteration for the mobile view yesterday, so I iterated on it. My goal was to figure out:

- How would the layout look empty?

- How would users add and edit events, attachments, or quick notes?



Above you can see the newest iteration. Here's how it could work:

- Each hour has a generous min height (this could be reduced a bit if we want to save some space) and they can individually grow based on the content.
- Tap the quick note, or an empty part of the layout to edit the quick note of an hour.
- Tap the ellipsis buttons of a given hour to add an event, a task, or an attachment.

Not designed yet, just thought about it:

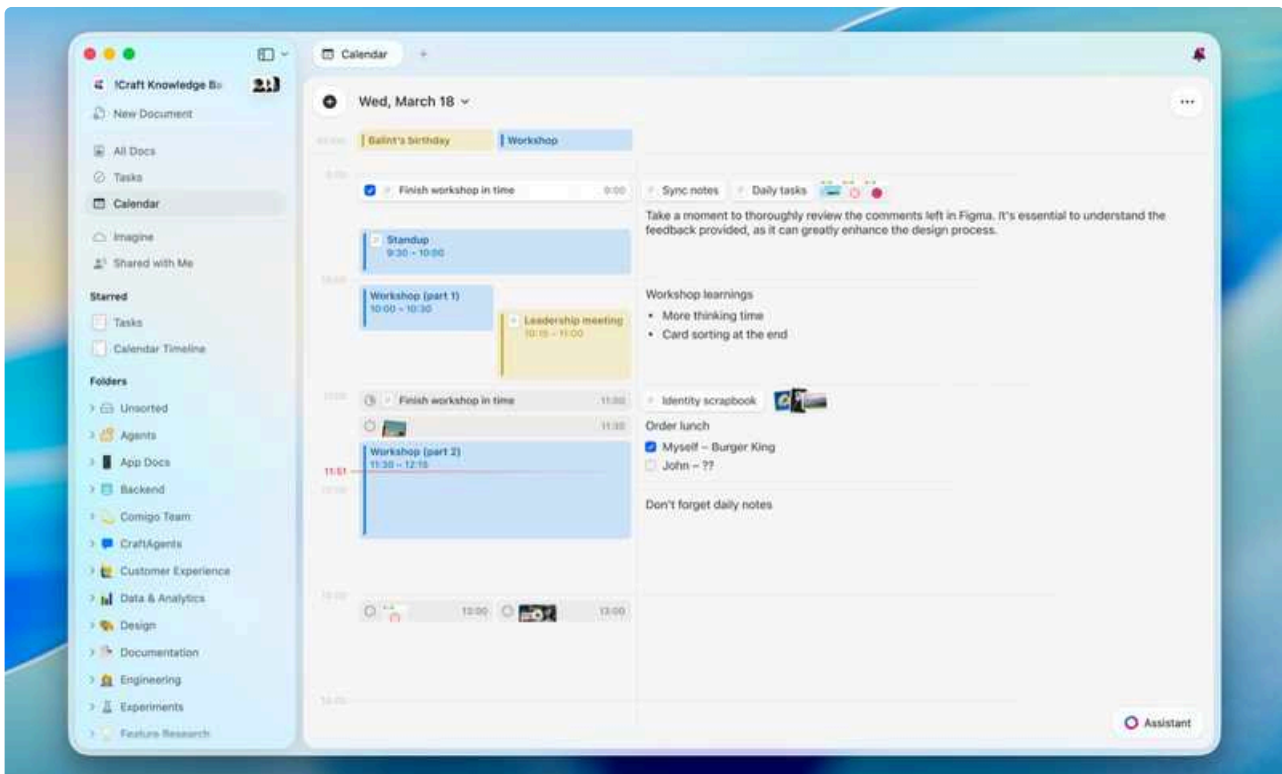
- Tap on an event, reminder, or a task to open their configurator popup.
- Swipe left of a note to enter multi select. Only notes will be enabled, events, attachments, tasks could be disabled (maybe only events) with a lower opacity.

Attachment + task visuals

I was playing around how we could make attachments and tasks playful:

- Website links could get a browser image / favicon preview without showing the title. This could make multiple websites more compact next to each other. Similar types of attachments could be grouped together for easier scannability.
 - Some sites like Github, Linear, Youtube, etc could get special treatment because individual links contain important info like the ticket number.
- Tasks, Events, Reminders could show possible Subpages / Event notes on the items themselves. Clicking them would send users directly to those pages.
- Reminders of non-text blocks (link, image, etc) could utilize the same preview used in the Entries column.

Since the entries column quick notes will be block based, it will be super easy to drag blocks or attachments over from the entries column directly to the events column to make blocks reminders. This way organization of a given day could be super fast and powerful.



Integration

I'll have chats with Peti during the mornings to figure out how we'll handle the integration of each feature:

- Tasks – They'll be kept on the daily layer in a widget-like interface (not yet designed).
- Reminders – All reminders scheduled for a given day will show up in the events view. Users could drag quick note blocks, daily tasks, attachments, and documents (possibly even from the left sidebar) to turn them into reminders.
- Hourly quick notes – Imagine them as mini documents with those few blocks inside them. Text blocks, and attachments like Images, Links, and linked docs will have their own block ID. The order of the attachments will need to be figured out so it doesn't clash with the notes' order.
 - Users will also be able to drag items between hours, or copy, cut and paste them into other days possibly.
 - Each hourly note cell should be ready to be typed or pasted into, possibly years into the future
 - An idea to try: Once a user's hourly note gets too long and it would seem logical to turn the text into a document, then users should be able to highlight multiple blocks (from different hours) and group them into a document. The multi select, grouping animation, etc are not ready in design, just an idea to think about.

 Mar 25., Wed

- Discussion with Peti about integration and his potential feedback on the design
- Discussion with Christoph and Lisa of the Timeline during the design alignment
- Continue design with:
 - Figuring out Task vs Reminder
 - Adding the day level widgets to the design.

 Christoph Locher

- Timeline discussion with Bene
- Created a PR to replace the "Out of Credits" error message with a styled banner with a button. For payed users it leads to the Credit minisite, for free users it now opens the upgrade flow
 - This also helped me look deeper into the message types, formats an conversion flows through an discussion with the agent and having it draw up various schematic charts
- Worked on a sidebar/panel-style implementation of the Assistant pages, to compare to the full screen one

 Mar 25., Wed

- Continue with sidebar/panel-style implementation of the Assistant pages

Remote setup on review

- Wrapped up the functionality there are still rough edges but core functionalities are tested
- Posted some teaser videos
- Main areas:
 - Session transfer
 - Mobile compatible WebUI
 - Server setup

Infrastructure study

- If we do not host the instances, but the users, we can provide tunneling making it easier to access for the users (no need to set up domains, proxies, tunnels)
- Did some exploration and the CF Worker + per user DO looks like a valid approach with an embedded client in the app
- I evaluated the new features of Claude Cowork and Co and they are aligned with what we do
- Computer Use is a new thing not sure if it comes in the SDK, but we can add it similar to Browser Use, however so far i am not 100 % sold in its current form(will be good, but bit slow still)

Next steps

- Polish the solution in prep for release
- 📅 Mar 26., Thu will be only half a day as i will be offsite in the afternoon, if any urgent bugs i can fix in the morning or later night (no release happened so i dont expect critical issue unless Anthropic comes up with something)

Assistant

- **Tool Feature Parity + Tests**

- Fixed issues with the new rate limiting interfering with test execution
- Disable tools on server side that the current client only supports behind feature flags (eg code execution tool and json space tools)
- Reran the review again and again as it keeps finding new feature parity issues, that it hasn't find before, so I keep fixing these
- Previous server daily note creation tool created a fake daily note, not an actual metaobject linked document, so implemented proper daily note + object link + meta object creation
- Previous server task listing tool had hardcoded structure, didn't support proper task grouping, so implemented those
- Server calendar tool also didn't support all filtering options, added the missing ones
- Fixed hardcoded comment thread resolved state in server tools
- Fixed server find blocks tool not including root block
- Fixed server user info tool not including some of the user data
 - Also the AI credit balance check was using direct postgres connection, which this service does not and should not have to the gateway's database, so I replaced the SQL query with direct REST API usage of the gateway credit balance API
- Wrote new tests for all the adjusted functionality that pass on both platforms

- **Background session execution**

- **Session sync**
 - Automatically syncing session history when reconnecting
 - Improved stamp + last message date based mechanism to determine if session is stale
 - Stale S3 session state also caused issues, so made sure, that active sessions always rely on the in memory or local disk data over S3 data
- **Socket reattach**
 - When we rejoin an existing session its not enough to download the missed content, as the session is now has a disconnected socket and running in full background more
 - Created a session reattach mechanism, that is able to reconfigure the toolprovider to use the new session and also inject the session object into various other components that are involved in the execution
 - With this the rejoined session continues transmitting messages live to the new client without losing any data
- **Load balancing**
 - Reattaching works well locally, because there is a single server instance, but in AWS we will require implementing sticky sessions in the load balancer to make sure the client can find the instance where the rejoined thread is currently executing, so I'm going to work on this next

📅 Mar 25., Wed

- Continue stability improvements + background session execution

↑ Daily Notes - 2026 March

Mar 24., Tue



Agents team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Timeline mobile design — team al...

↑ Mar 24., Tue

Agents team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Timeline mobile design — team alignment on option 4** — Benedek explored four mobile layout approaches and favors option 4 (list with hourly dividers). Does the team agree? The desktop distortion issue remains unresolved — is the list approach viable there too?
- **"Out of credits" messaging ↔ CX team** — Christoph is updating the error to a styled upgrade banner today. CX Team's Tamas just aligned with Christian, Christoph, and Viktor on AI credit complaint handling. Sync on the user-facing copy and how this fits the new complaint playbook.
- **Remote session transfer — demo timeline** — Gyula fixed critical bugs (LLM connection lock, sdkCwd path mismatch, workspace creation flow) and has WebUI working on mobile. What are the remaining blockers for the e2e demo? Is the relay server approach ready to pilot?
- **MCP URL normalization fix** — Gyula stopped force-appending `/mcp` to all HTTP URLs, which was breaking custom-path servers like Adzviser. Are there other known external sources affected by this pattern?

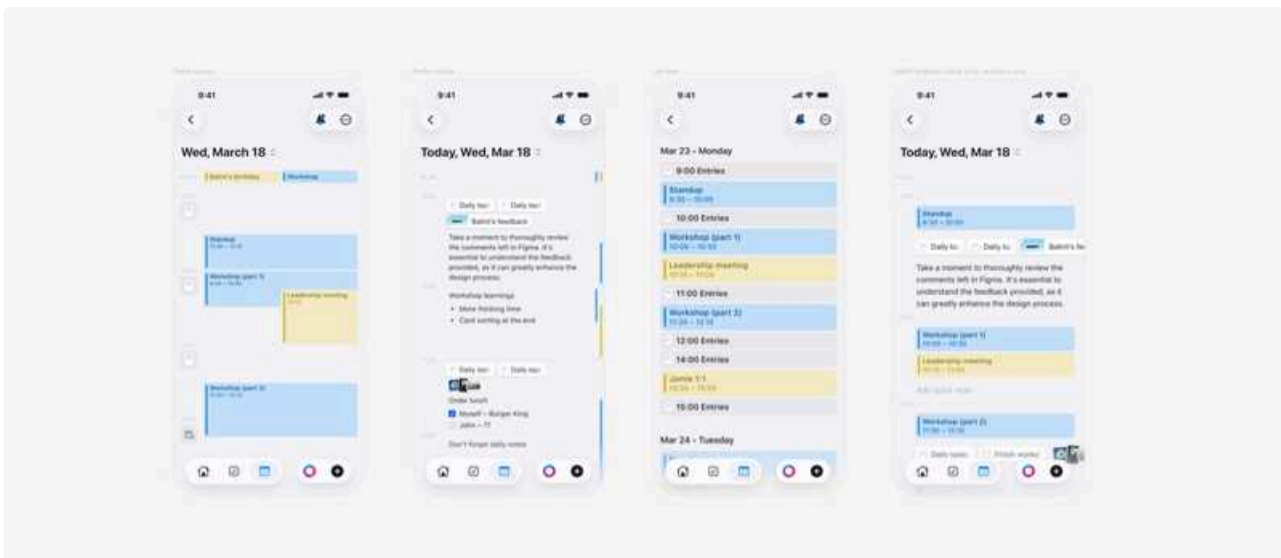
Kickoff with Peti and Lisa

I shared the current state of the Timeline with them. I'll have daily talks with Peti on where the project is going so he can mentally prepare with the integration. He might also have insights on what works and won't work on an implementation level.

Timeline

Today was not my most efficient day. I was often staring at my screen just thinking of new ways to solve how this information dense Timeline layout will fit into a mobile screen while providing in-line editing, a nice overview of the day, and avoiding the issue where too many rows of notes will stretch the events in a given hour.

I made four new versions of the mobile view, each with a different approach.



1. Prever events

Entries and attachments are grouped behind an hourly quick note container, so any amount of content won't distort the calendar.

2. Prefer entries

Events are super subtle, only used to give some grounding to the entries.

3. List view

Events and entries are both listed in order of event start. I passionately dislike this list type view, my eyes just cannot understand what I'm seeing.

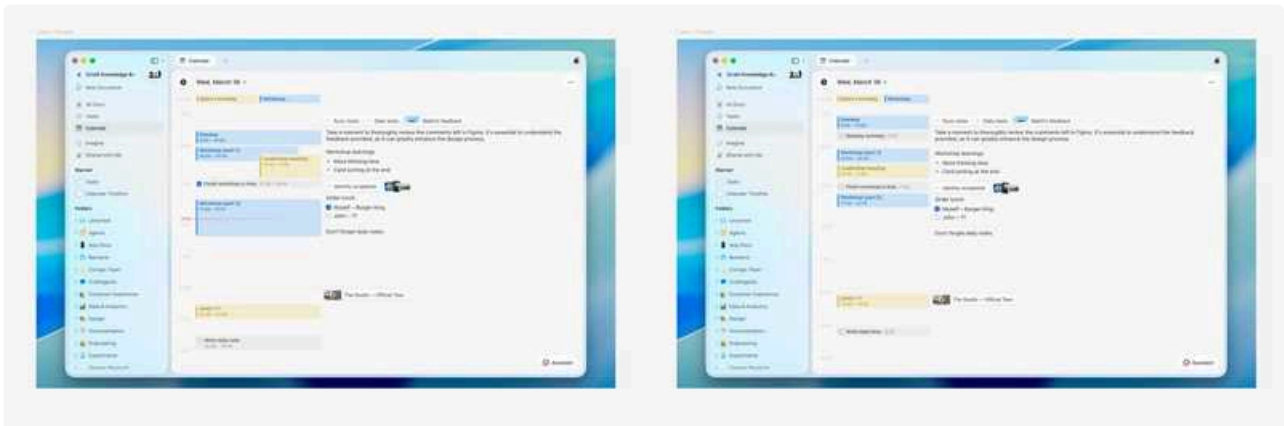
4. List view

Events and Hourly entries both visible in one view. Events are listed like in the list view, then you see an in-line version of the attachments. For some reason the list confusion doesn't kick in here for me. Separation and the timeline-like hourly dividers help me understand the timeline.

My favorite one is the 4th one so far. It is clear, in-line, and doesn't suffer from distortion like a side-by-side layout would have.

I started thinking whether this list type view could work for the desktop (to avoid the distortion issue there as well) but for me it doesn't work.

(please disregard the location of tasks, that is another rabbit hole)



Mar 24., Tue

- Iterate on the 4th option for the mobile view to see how empty states, long and short notes, etc will look like, just as a stress test.
- Focus on integrating tasks

Christoph Locher

 Mar 23., Mon

- Some issue investigations
- Sync with Viktor
- Continued working on the prototype, lots of little refinements, additional flows



Github

Directory not found

<https://github.com/lukilabs/Luki/tree/christoph/assistant-vision-prototype>

 Mar 24., Tue

- Update "Your out of Credits" error message to styled upgrade banner
- Continue with the prototype
 - Further improve Create Automation Flow
 - Tweak Chat Messages/Turn Group

The great Bedrock excavation + remote transfer gets closer

LLM Infrastructure & Bedrock Fixes

- Fixed Bedrock model ID resolution — bare Anthropic model names caused `400 ValidationException` on AWS; added static mapping + discovery fallback
- Fixed stale `CLAUDE_CODE_USE_BEDROCK` env var leaking into Pi SDK sessions
- Ported Bedrock auth normalization from OSS repo — IAM/bearer/env auth branches, connection validation
- Added 1h prompt cache TTL toggle — app-level setting applied by network interceptor
- Fixed 1M context beta header still being sent when disabled — SDK subprocess now respects the flag via env var
- Strip `cache_control.ttl` when extended caching disabled to prevent SDK from overriding user preference

Remote Session Transfer

- Fixed critical LLM connection lock bug — forked sessions no longer preserve source server's `llmConnection/connectionLocked`, preventing SDK auth hangs
- Fixed `sdkCwd` path mismatch on cross-machine transfer — always regenerate from target workspace root
- Fixed remote workspace creation flow — switched from overloaded `TEST_CONNECTION` to direct `invokeOnServer` RPC
- Designed branch-fallback summary approach for remote transfer — mini completion to condense context instead of raw history copy

Pi Agent Stability

- Guarded Pi SDK compaction race condition — `waitForCompaction()` polling before `prompt()` to prevent concurrent `AbortController` crash
- Added branch fork fallback — summarize parent conversation via mini completion on fork failure so agent retains context
- Set up Cloudflare tunnel for remote server access via `agents.craft.do/remote/ghalmos`

MCP & Source Integration

- Fixed MCP URL normalization — stopped force-appending `/mcp` to all HTTP URLs, breaking custom-path servers like Adzviser
- Fixed MCP schema validation — strip `$schema` from external tool `inputSchema` so Pi agent AJV doesn't choke
- Set up Adzviser marketing data source (30+ platforms)

Remote Setup

WebUI

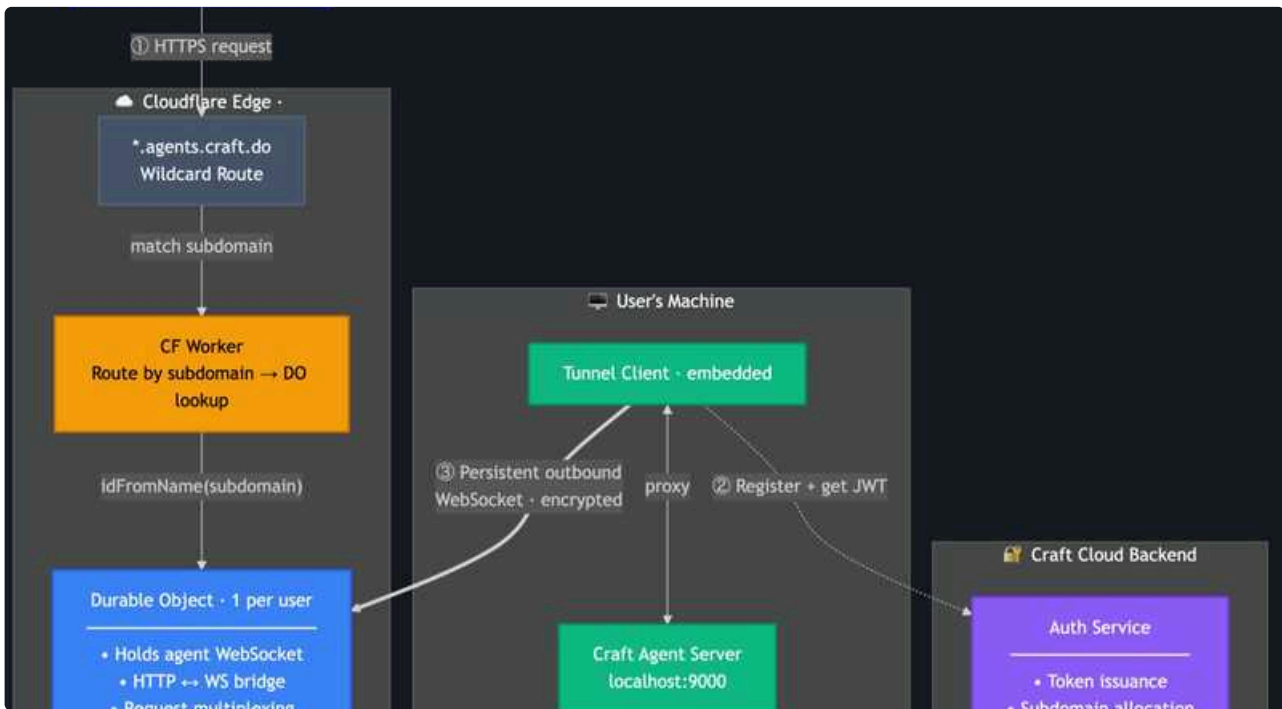
Mobile

- Made WebUI responsive for mobile viewports — CSS overrides for sidebar/navigator/chat panel
- Fixed mobile "Back to list" button navigation and integrated into PanelHeader
- Addressed settings overflow and cascading menu issues on phone screens

WebUI

- WebUI is working nicely it is fully functional except OAuth this needs to be streamlined Claude OAuth worked, ChatGPT Plus not yet, AWS Bedrock and API keys also work fine

Infrastructure for the future tunneled setup



Next steps

- Stabilize remote session transfer — proper history handling, LLM connection resolution (works with Claude Agent SDK, still issues with Pi Agent)
- Pilot relay server approach for remote connectivity
- Hopefully i can do an e2e demo and teaser as well :P

Assistant

- **Tool Feature Parity**
 - Added tool execution timeouts
 - Extracted edit folder validation logic to server side wrapper
 - Updated tool input/output schema definitions to fully be identical
 - Added tool schema comparison tests
 - Some additional feature parity tweaks
- **Performance**
 - Managed to find a way to speed up the mac harness tests
 - Instead of running all test files in parallel have a constant thread count (5)
 - The largest test files are executed first and finished threads are continued with additional test files to keep the thread count consistent
 - Each thread has a dedicated mar harness process, thats started and kept alive throughout the execution instead of terminating after each harness command
 - By eliminating the harness bootstrapping the performance increased 10x (from ~20 minutes to ~1.5 mins)
- **Release**
 - Reviewed Chris's change about gateway rate limit handling for the claude agent SDK
 - Pushed out the current release version with these changes and other recent release improvements
- **AI Tools Service review**
 - Started reviewing and productionizing the service code
 - Strictly tying all tools and actions to spaceld or userId
 - Enforce ownership check for S3 exportable URL generation
 - Propagate trace/request header for correlation
 - Added API key based authentication for AI tools
 - Updated prod instance count to 2 for redundancy
- **Background session execution**
 - Added tool fallback path in tool wrapper, that falls back from client side tool to server side if the client becomes unavailable
 - Continue generation when socket is disconnected
 - Also automatically reject pending tool calls to avoid the generation getting stuck
 - Immediately persist session changes to S3 when generation finished
 - The client can fetch the remote changes when rejoining the session

- Continue stability improvements + background session execution



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Gyula's Remote Workspace doc...

↑ Mar 23., Mon

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Gyula's Remote Workspace doc needs review from Balint** — Gyula flagged the "Remote Workspace in Craft Agents" document as needing review before he can continue. Check if Balint has had a chance to review over the weekend, and unblock if not.
- **Remote session transfer edge cases** — Gyula noted unresolved issues: LLM connection transfer when moving a Claude Max session to a remote without Claude Max configured, and source availability mismatches. Discuss how critical these are for the initial release vs. follow-up.
- **Benedek working with App Team on Timeline prototype** — Cross-team work today: Benedek plans to collaborate with the App Team on the first prototype/implementation of the timeline layout and interactions. Confirm alignment on scope and approach.
- **Christoph's week goal: complete screen/flow coverage** — Christoph aims to have all desktop and mobile screens/flows covered by end of week. He specifically needs to build the automation creation flow and schedule edit flow, which involves an agent extracting schedule from user input (needs loading/error states). Discuss any design dependencies.

 Benedek Gobolos

 Mar 20., Fri

Timeline



Event Interaction

This doc highlights how the event configurator popup and Event item states · Event states · Image · By default eve...



Mobile layout

This doc explores how the Timeline could look on small screens. · The goal is to keep the usefulness of the two co...


 Mar 23., Mon

Work with the App Team to make the first prototype/implementation of the timeline layout and its interactions in the main app.


 Christoph Locher

 Mar 20., Fri

I continued working on the prototype to explore more of the behavior and flows of the various sections. I recorded a video walkthrough with some commentary:

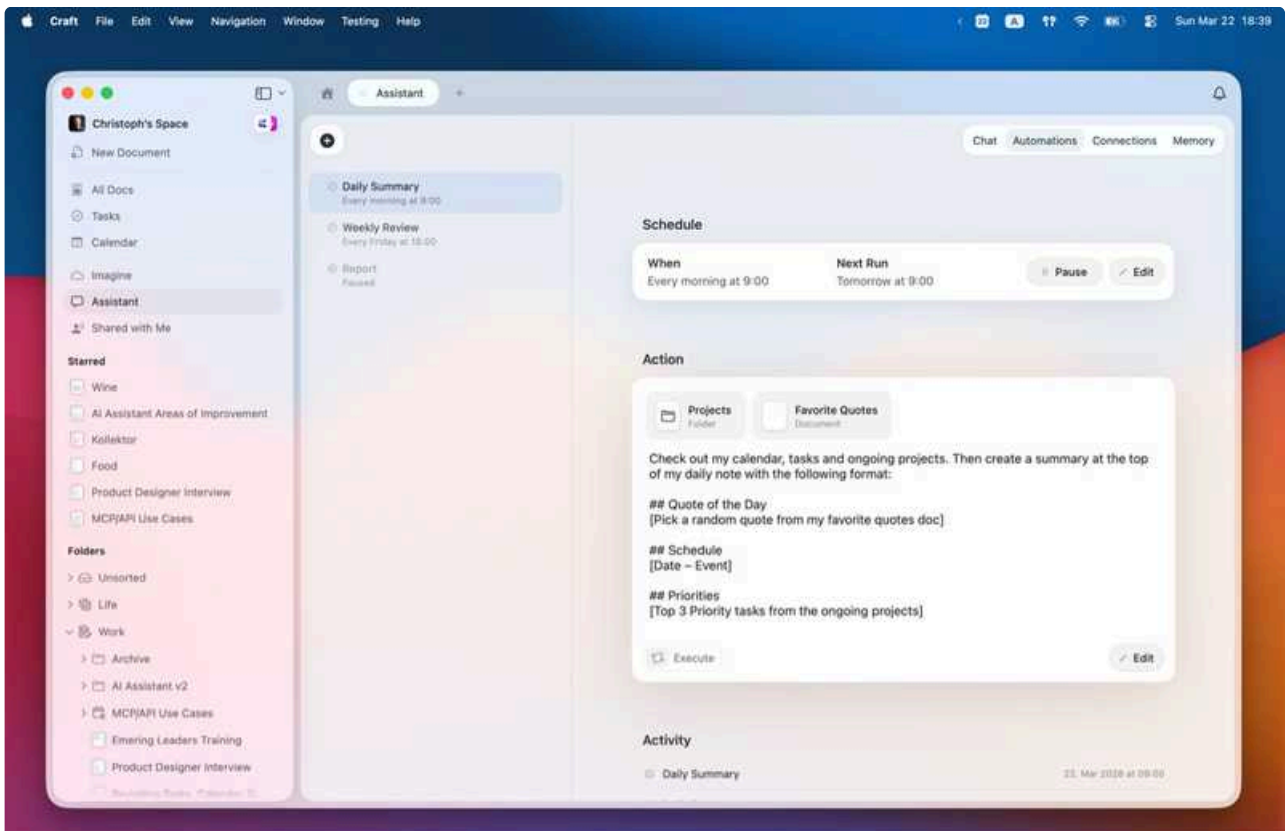


Chats

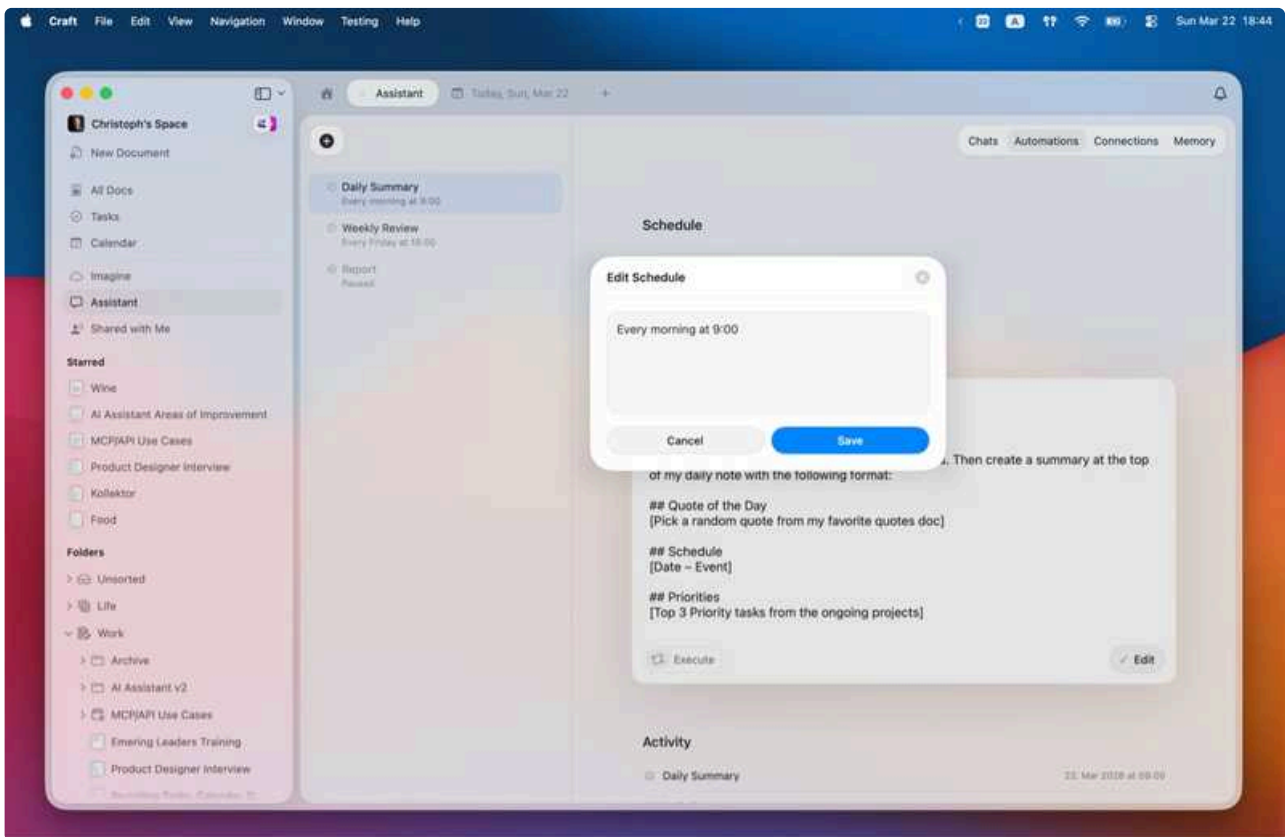


Because this is using the real thread data, this is where I did the least changes so far, but I am narrowing down the UI states threads can have and will continue to extend this more.

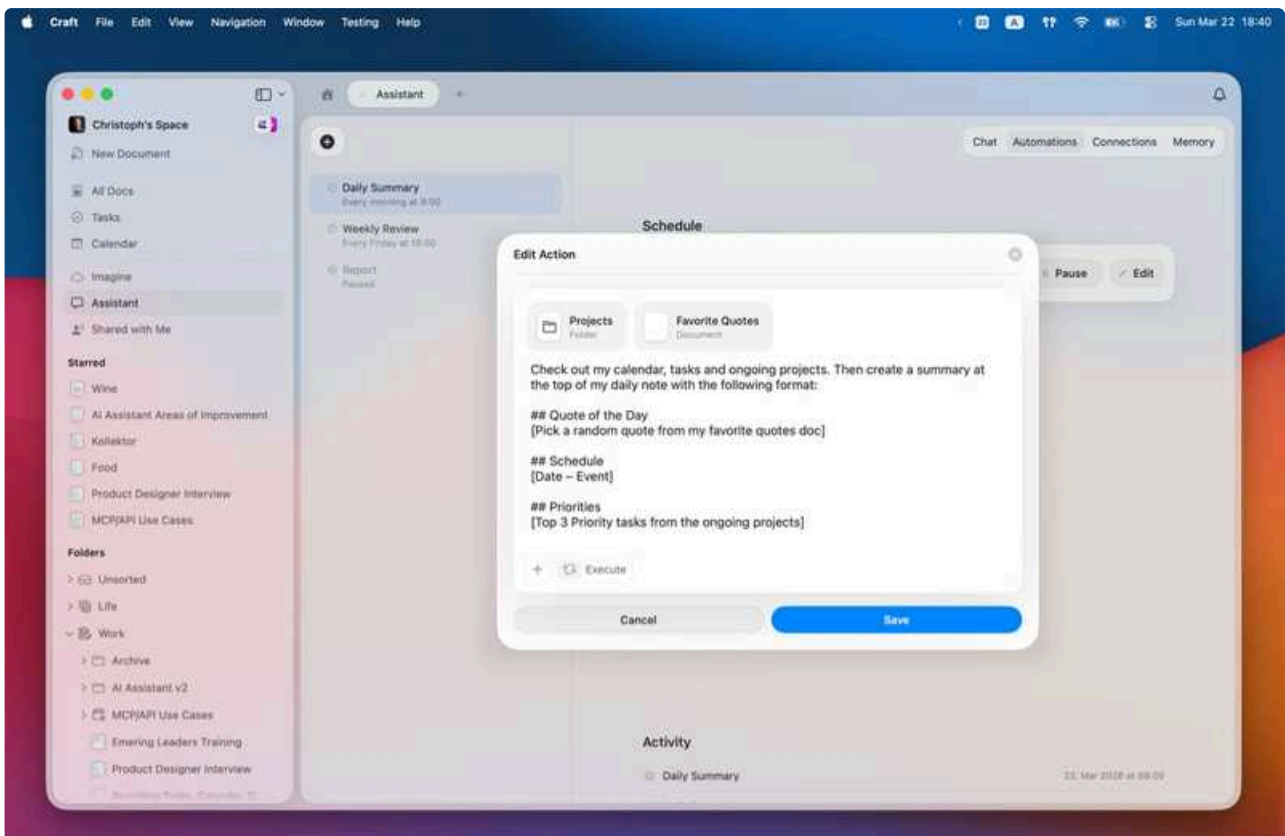
Automations



Automations are displayed showing their schedule and the prompt they trigger. I like using the chat input to show the automation content here, because it symbolizes that an automation is essentially just an automated new thread with a regular prompt, being created and run at a specified time.



Changing the schedule happens in a panel. I need to visualize the flow here further, because essentially we need to run an agent to extract a schedule from the user's input here and this needs a loading state/error state.

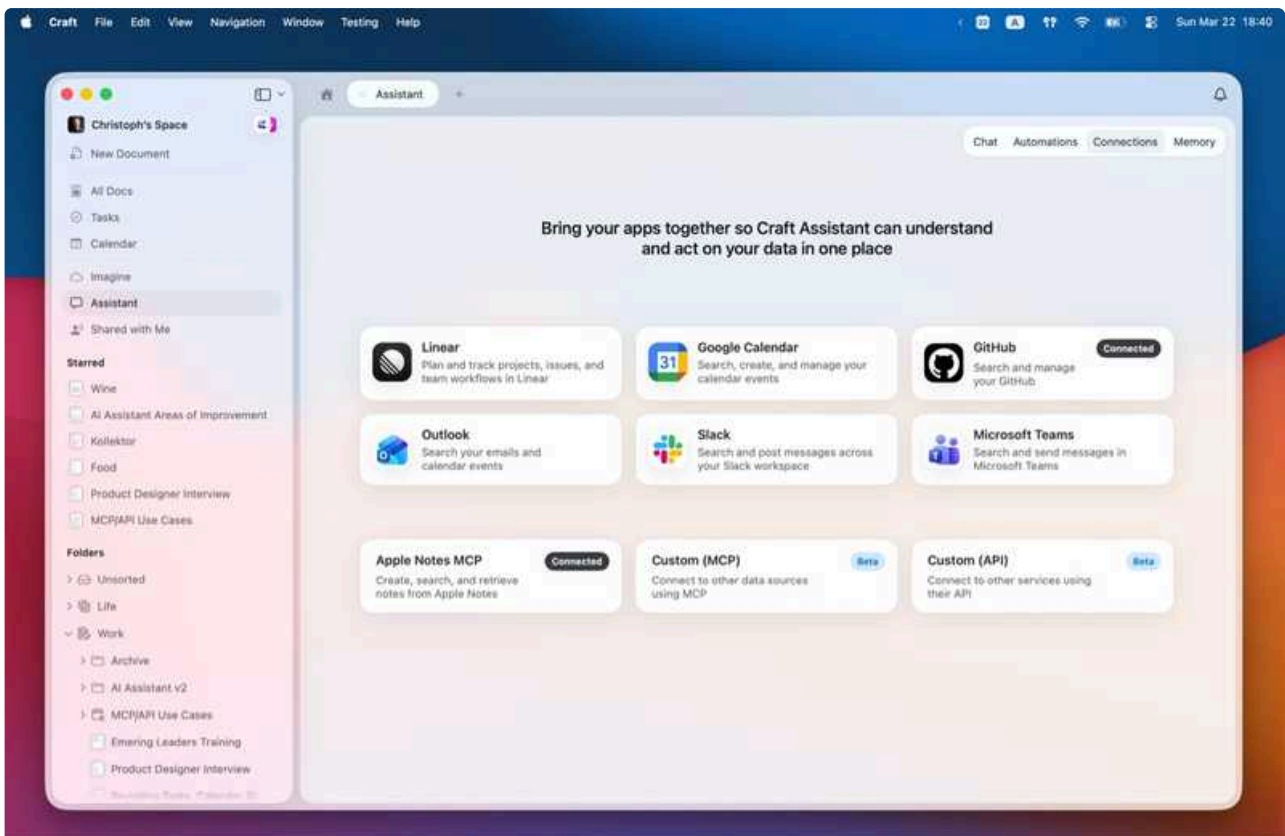


In this view, users can edit the automation prompt, set the mode and model and add context.

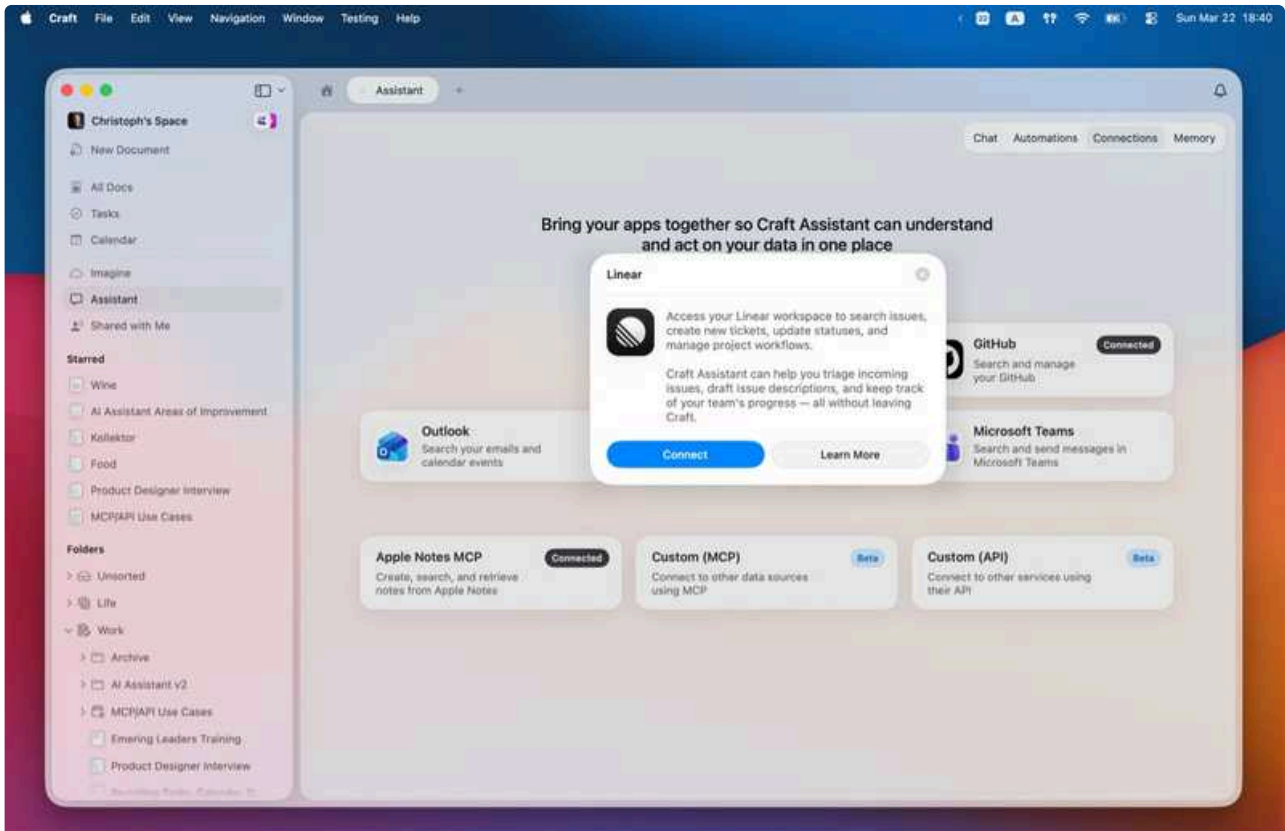
Below this section there is an activity log, linking to the threads created when this automation ran in the past.

I need to look into the flow of creating a new automation from scratch next.

Connections

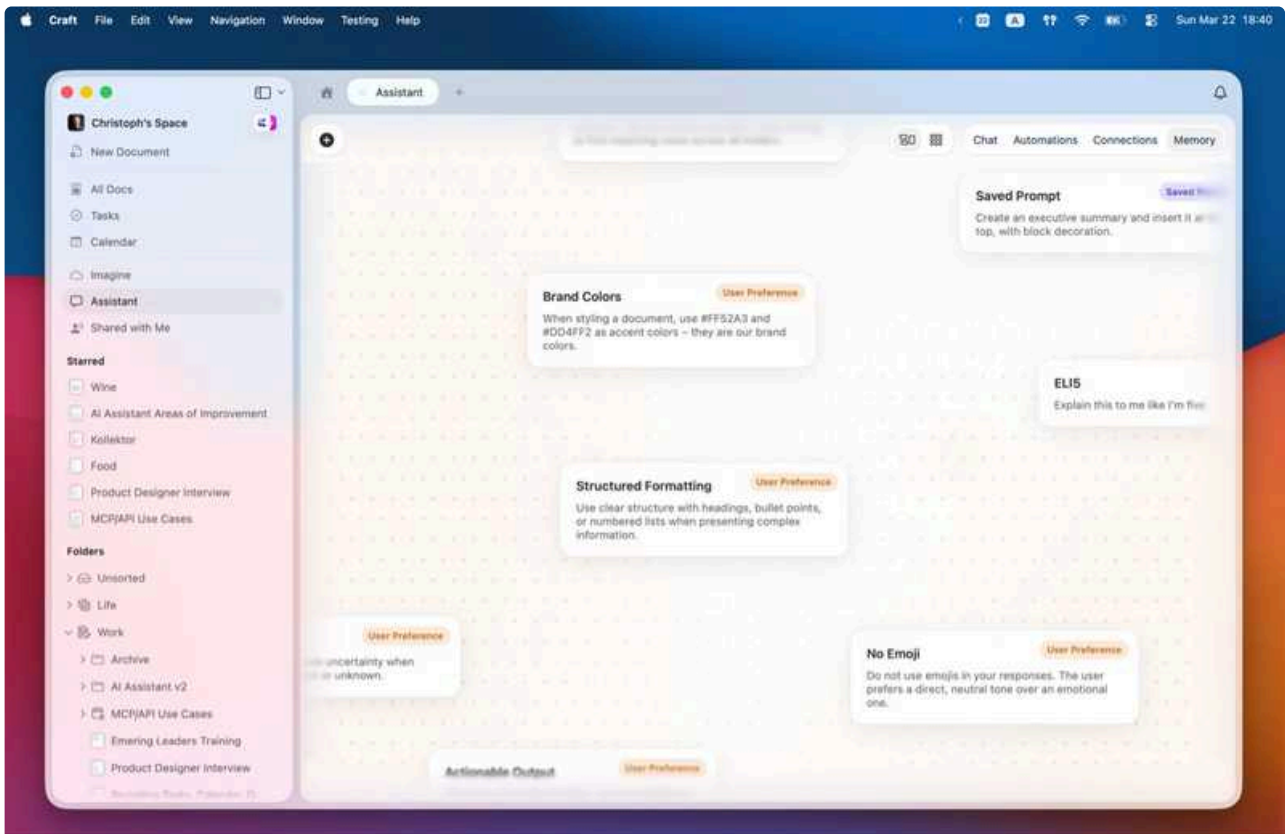


For each connection there is a panel with additional information and the actions to connect/disconnect it. This is placeholder content, but the flow works already:

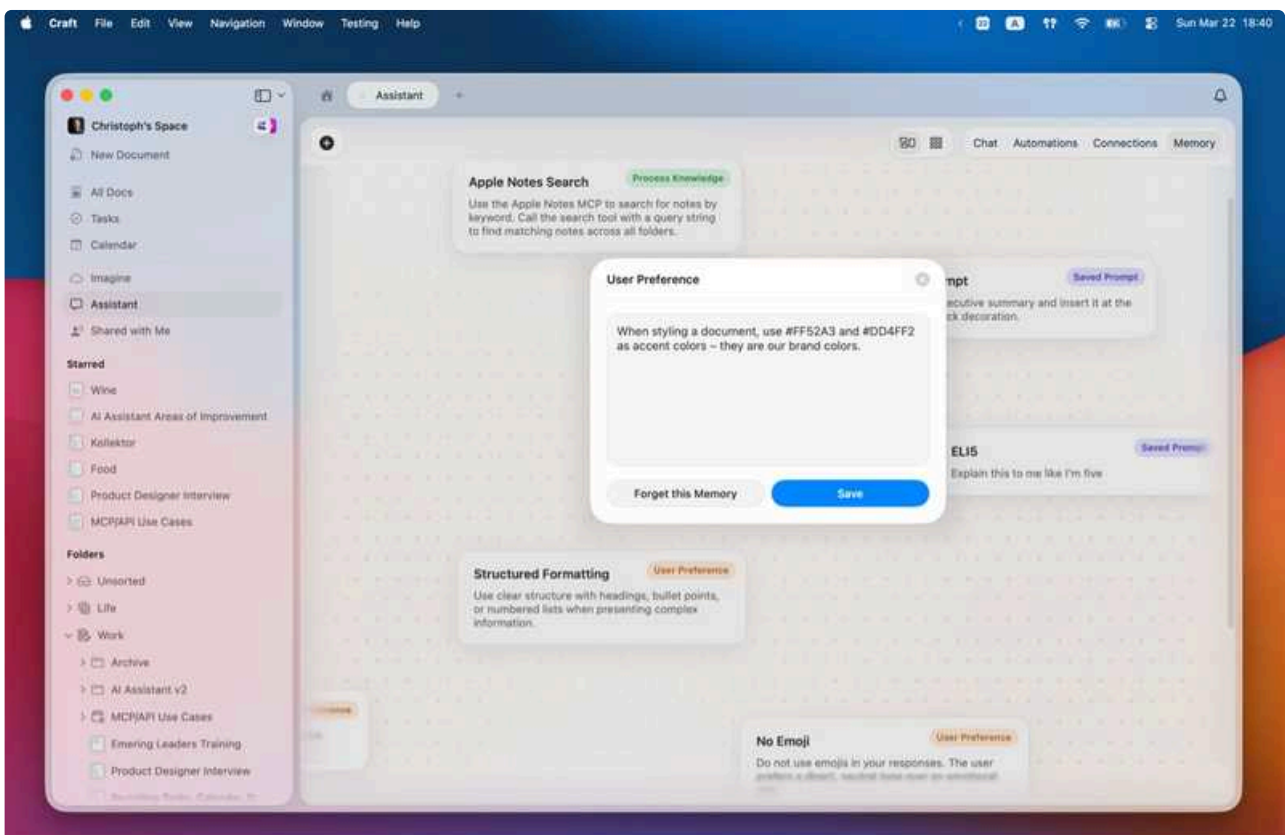


Memory

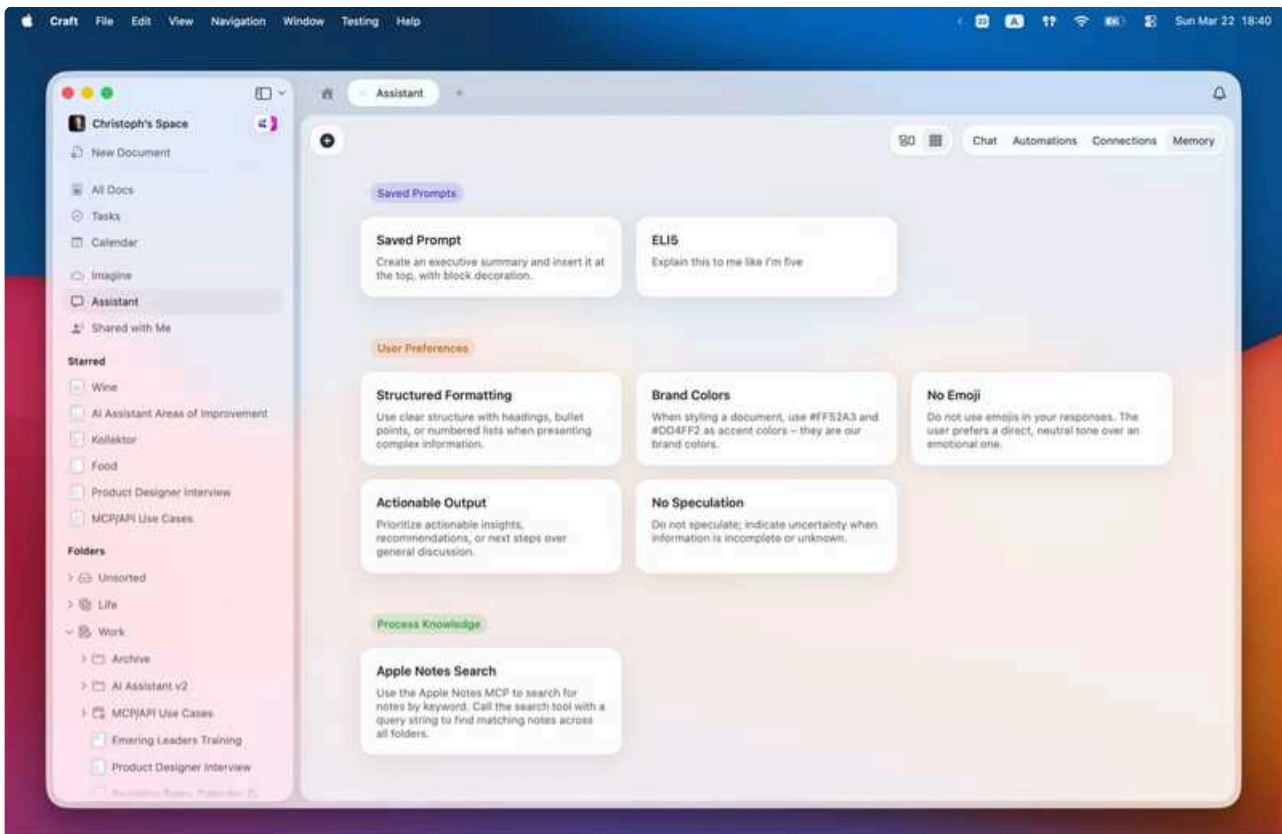
Memory is the most explorative part. It has a floating view where the memories are spread out spatially, there is some movement in the back and they float around.



You can tap each item to see detail, edit or delete the item:



Because we know that the floating view, while fun and immersive, gets a bit hard to oversee with many items, there also is a more traditional grid view to switch too.



Mar 23., Mon

Goal for the Week

Have complete coverage of all the screens/flows on desktop and mobile.

- Continue with the prototype
 - Chat View: Visualize the different chat states, start looking into extending the message display
 - Automation: Build the automation creation flow and schedule edit flow

Goal of the week

- Finalize Remote and release either internally or to everyone depending on the stability and capability
- Spin up remote server that can be used
- Bugfixes and update to latest Pi (needs thorough testing)
- Start looking at it how to set up the infrastructure and user management for remote offering (like Nabu Casa with Home Assistant)

 Mar 20., Fri

Release & Bug Fixes (v0.7.10 → v0.7.11)

- Released v0.7.10, kicked off v0.7.11 with three targeted fixes
- Fixed model name truncation (#444) — custom endpoint models like `gpt-5.4` displayed as just "Gpt"
- Fixed context window override (#448) — custom endpoints hardcoded to 128K, user config now flows through
- Added per-workspace 1M context toggle for Opus 4.6

Remote access

Remote workspace and session transfer is shaping up -  Balint Orosz if you can check on the doc below much appreciated and i continue from there on Monday

Remote Workspace in Craft Agents [NEEDS REVIEW]

Remote Workspaces · Connect a remote workspace · It is possible to connect a remote headless server's workspace...

- Designed and started building browser-accessible Web UI (`apps/webui`) reusing Electron renderer components
- Cookie-based JWT auth with security hardening
- Fixed remote WebUI auth bug — Secure cookie flag coupled to RPC transport, breaking plain HTTP logins

Infrastructure

- Designed Docker multi-platform build pipeline for headless server (amd64 + arm64, GHCR, CI smoke tests)
- Reviewed server binary deployment architecture

 Mar 20., Fri

Assistant

- **Tool improvements**
 - Fixed tool hook parameter handling issues
 - Fixed page style handling issues
 - Added parity tests + fixed for save document properties
 - Fixed server-tools ID mapping for block operations
 - Unified code path in tool wrapper between mac and server tools to avoid future issues
 - Full tool input and output schema review and parity fixes
 - Added missing output schemas to server tools as well
 - Mac harness test stability improvements to reduce flakyness
 - Did also a performance review why the mac tests take so long
 - The mac harness has to be started multiple times during a test to run the harness operations and currently the harness takes 1-2 sec to load and it starts a new process each time (to support heavy parallelization)
 - It's not an obvious fix to make these faster at the moment, so I left them for now with more parallelization and higher timeouts
 - Fixed a bunch of feature parity, validation, schema handling, input/output issues
- Sadly even with ~400 tests and very high line coverage I could still find issues with manual end to end testing
 - But I'm also adding more tests for these during the fixes and making sure both mac and server tools pass
 - Also 1 review is nowhere near enough, but I was doing a loop of **reviewer thread reviews mac vs server tools parity → fix parity issues → review and fix again and again**
 - During the middle of the week my codex limits were running out heavily so I only did reviews with codex and fixed them with claude and I realized how inaccurate those fixes were that came back later several times to cause problems.
 - So I decided to **rely only on codex** for the fixes as well as the results are **much much better**. (note: now I subscribed to another personal chat GPT pro for a 2nd batch of codex usage limits that I can juggle with the other one throughout the week to not get blocked by the quota as it was depleting way too fast)
 - The reviews are becoming better and better each time, but they are also going deeper when the basics are fine. Eg now it found that the folder icon editing is not as thorough on the server, as we are validating the icons from the mac bundle, so I will have to move that to the server side (and unified) with the list of valid icons

 Mar 23., Mon

Weekly Goal

- I believe the reviews and parity fixes should still continue with some occasional end to end testing too see that everything is fine. These repeated review + fix + run all tests loops take significant time, but I believe the core part of the code is stable enough to allow parallel work, while these are happening, so both server productionization and new feature development could happen in the meantime
- **Parallel work**
 - **More parity improvements + testing**
 - **Taking full ownership of AI Tools service**
 - As we discussed the Agents Team will own the AI tools service as well
 - Review and understand the current code on a much deeper level, not just the tool implementation, but the entire infra to make sure authentication, authorization and stability is ready
 - Productionize whatever necessary based on the reviews, add logging, monitoring and dashboards for proper visibility
 - **New feature development**
 - We can pick a few of the most important new feature ideas from [Christoph Locher](#)'s list
 - Also with the server tools we can now implement the tool hot swap and fallback mechanism for full background session execution, that doesn't require an active socket connection
 - **Priorities**
 - Start with background execution
 - Look at memory implementation
 - User preferences first
 - Maybe dynamic memory after
 - Automation

↑ Mar 23., Mon · Agents Team

Event Interaction

This doc highlights how the event configurator popup and Event item states

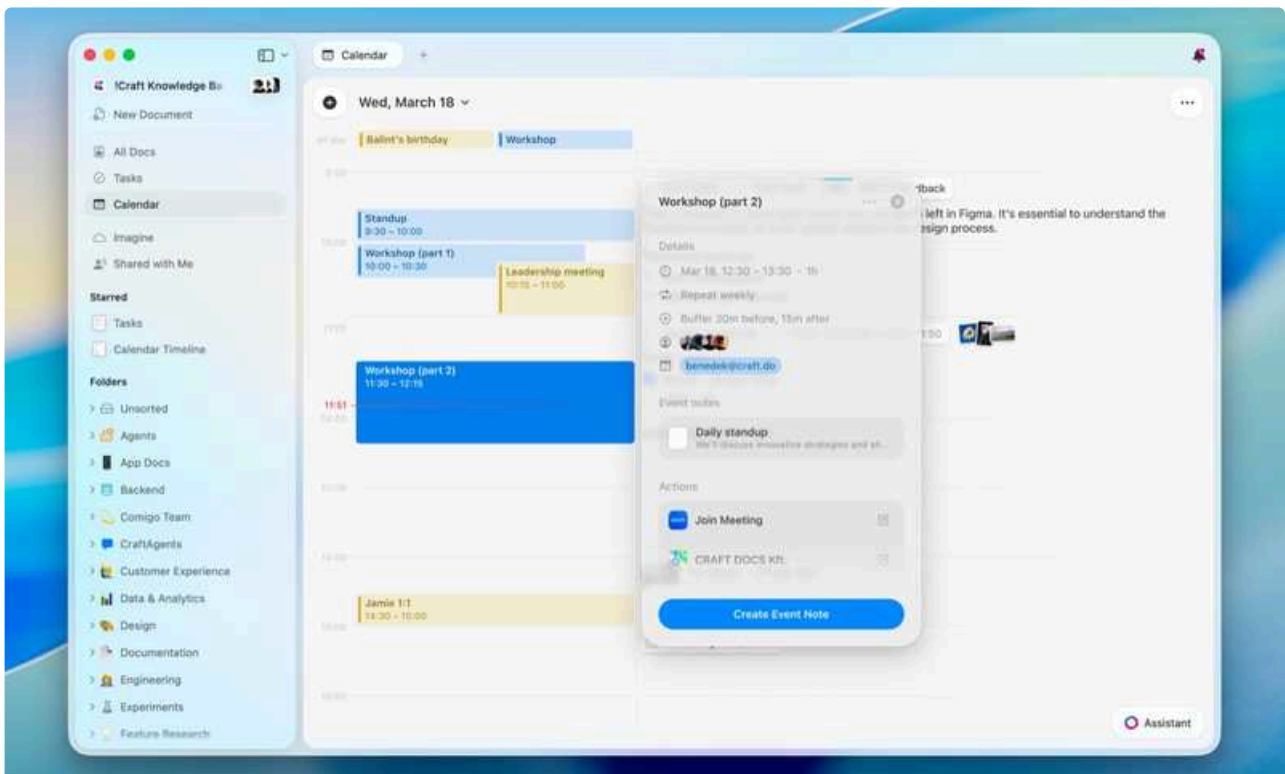
Event states



By default events have a subtle background and no hover states. Clicking on them or dragging them will change their backgrounds to be fuller.

Events could get a buffer area before and after an event as preparation, travel, focus blocking, etc. This visual is not changed when event is interacted with.

Event configuration



This is how the event popup will open. It can have a lot of information, so its max height is capped off at 550pt with scroll masks.

Options

The configurator popup has options for the core things a calendar feature could offer:

- Calendar type (show available calendars from connected accounts)
- General
 - Event title
 - Add call
 - Add location
- Schedule
 - Start date & time + End date & time
 - All day toggle

- Repetition
- Buffer (time blocking, Craft specific)
- Guests
- Description

Features absent from the flow:

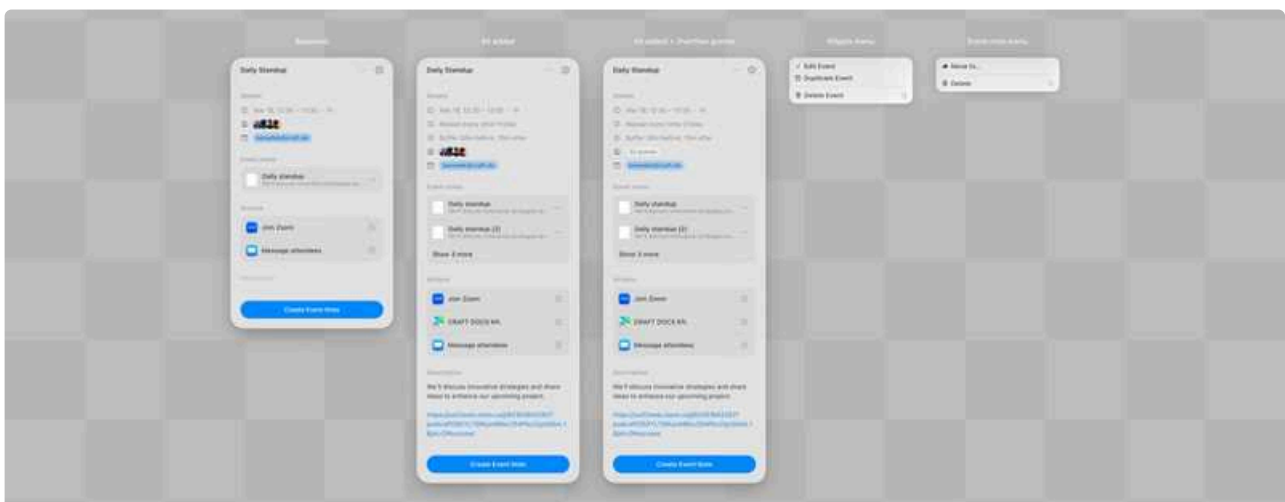
- Invite acceptance – Not shown
- Custom repetition configuration – Handled if there is external custom config

Preview

This is how clicking on an event would look. We could provide a preview of all the details, instead of throwing users to the edit view instantly.

User can do a handful of actions in this preview mode:

- Edit, Duplicate, or Delete the Event
- Open, remove, and move event notes to a different folder.
- Open call
- Open location on Maps
- Message all guests via Mail

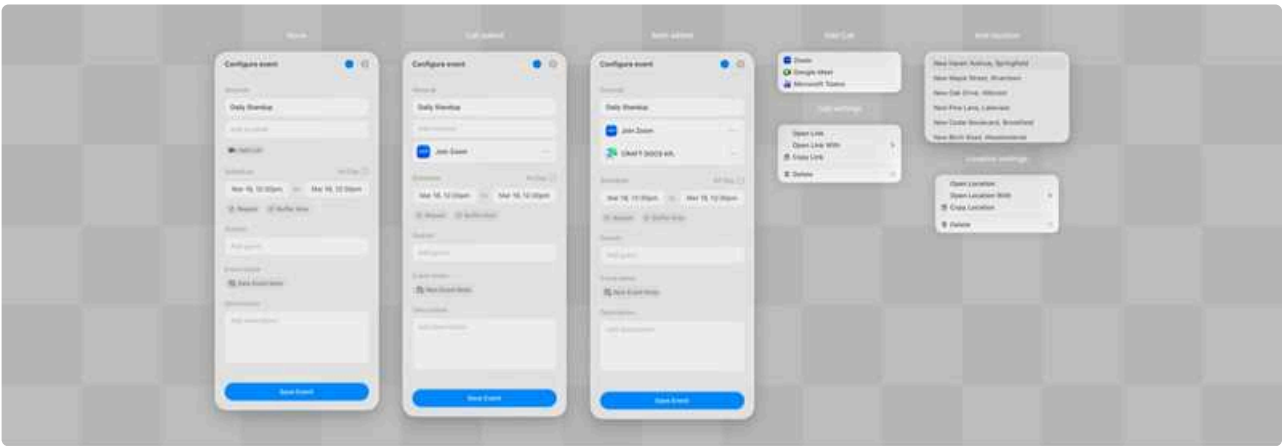


Edit view

I was aiming to provide core options without adding tons of new UI elements. I only used custom lists for search purposes, even then I repurposed the space search results list.

General

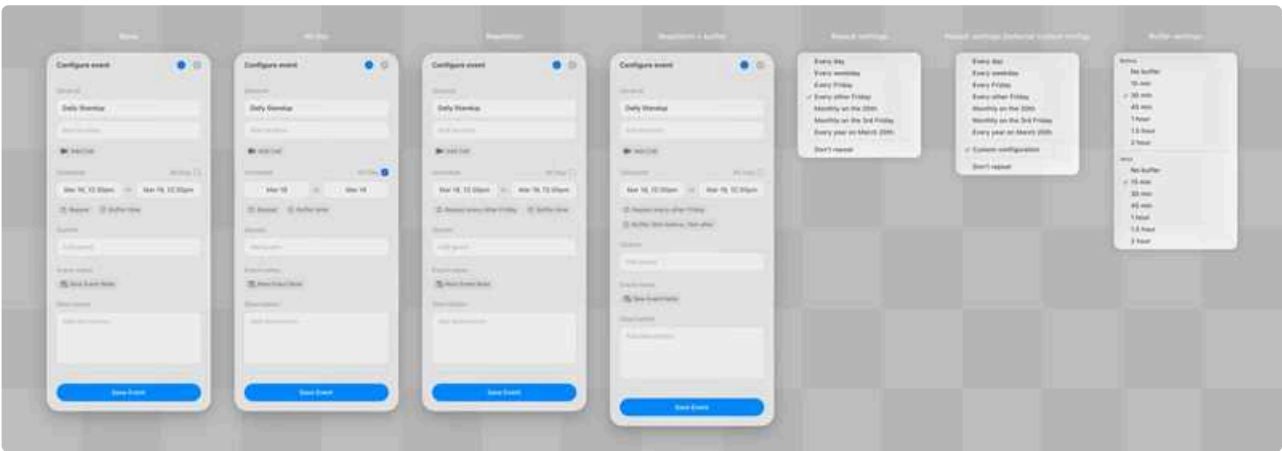
Event title, location, and call attachment is in the general tab. Location search uses custom list, others use OS native pickers.



Schedule

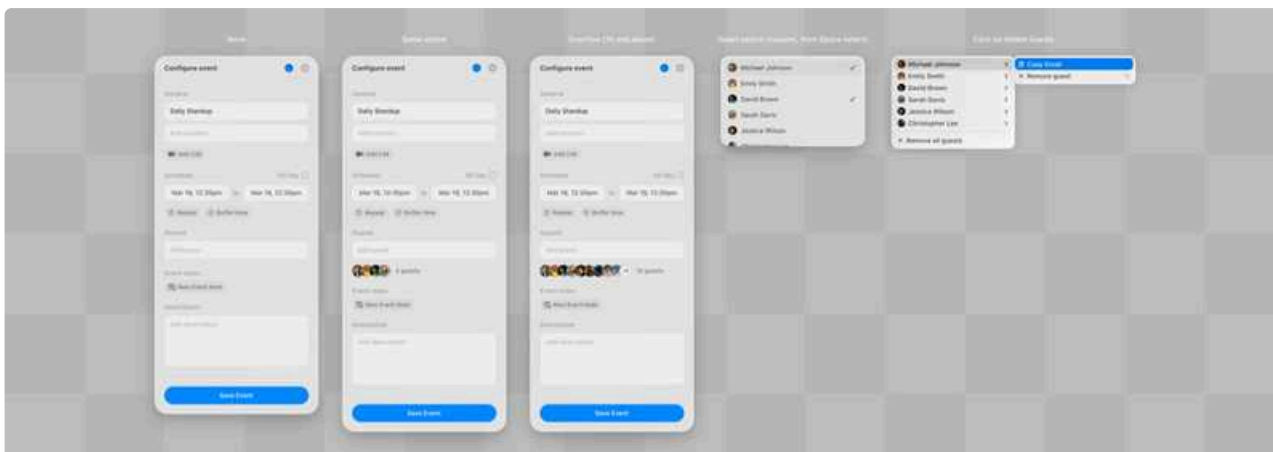
Supports Start + End dates, All day toggle, Simplified repetition options and Buffer times. Date & time picker uses Craft's existing popup, while other fields use OS native options

Is an event has custom configured repetition added, then we'll display it in the picker as a "Custom configuration". Selecting another configuration will remove the custom one and it won't show up again for the event.



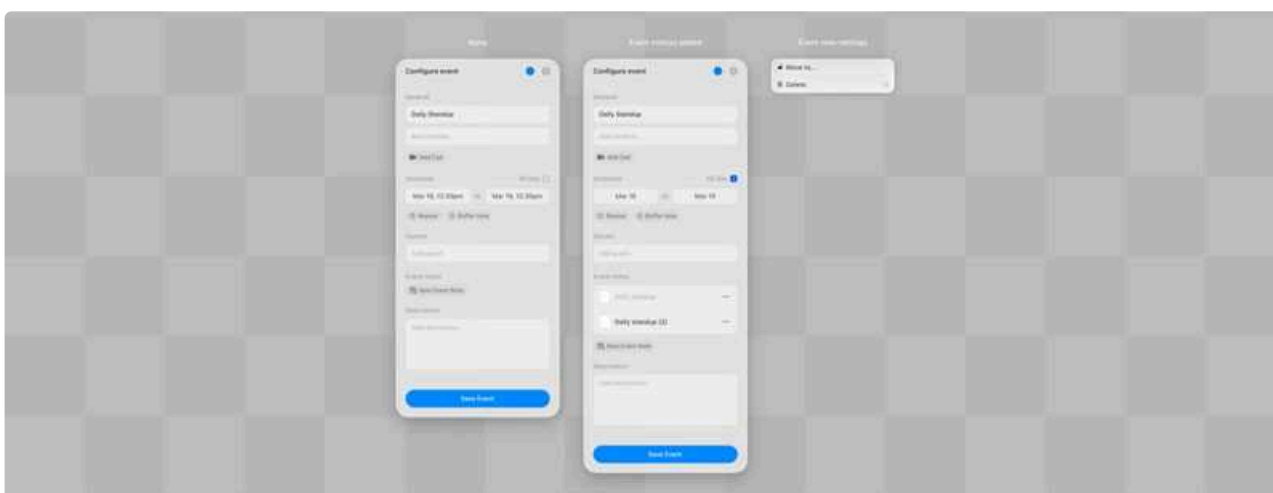
Guests

Users can add guests, but invite acceptance is not visible. Already added guests can be deleted by clicking the guest list, then removing one or all members via the OS native menu.



Event Note

Users can add event notes inside the edit menu as well. Options to delete and move notes to a folder are also available



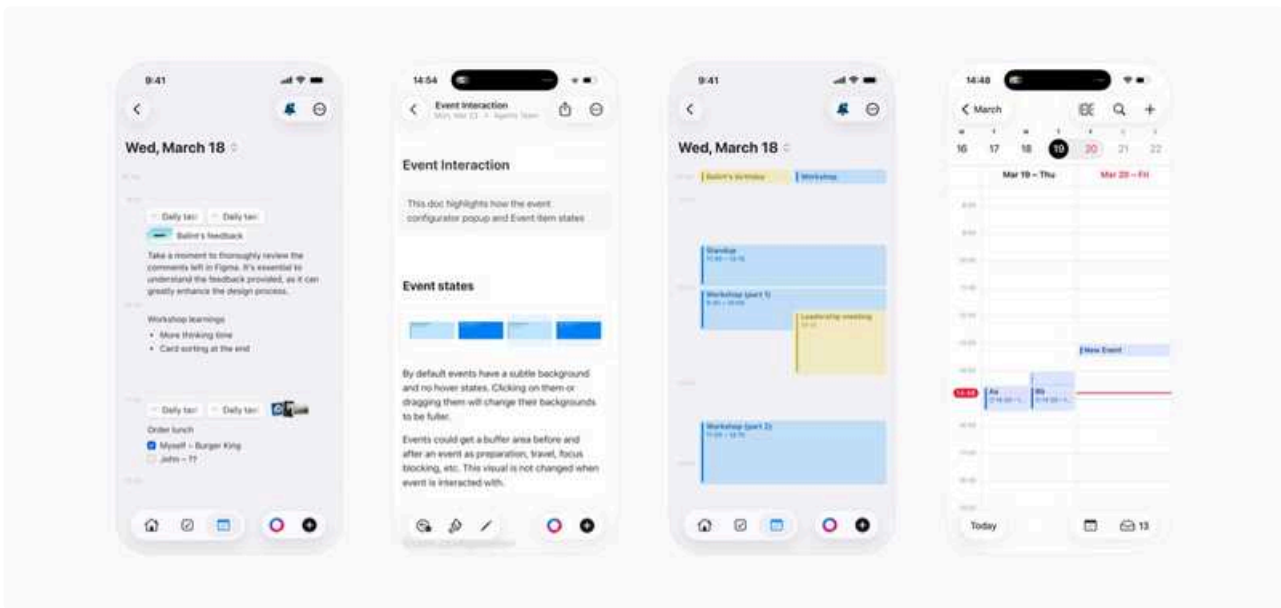
↑ Mar 23., Mon · Agents Team

Mobile layout

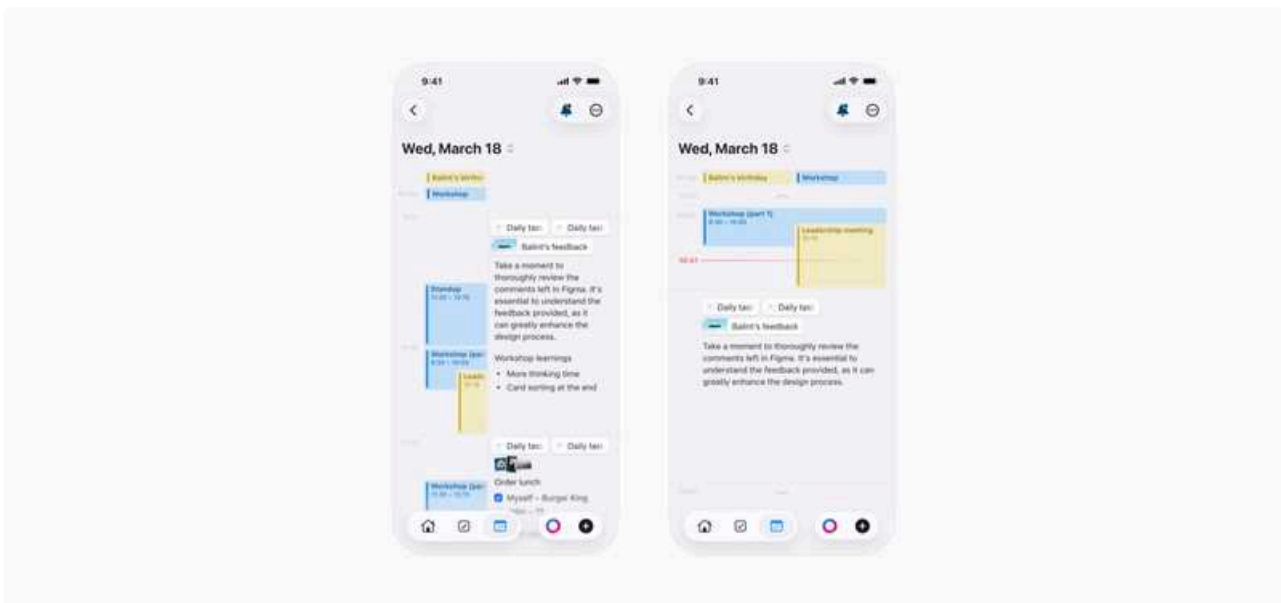
This doc explores how the Timeline could look on small screens.

The goal is to keep the usefulness of the two columned timeline while accounting for the super compact screen. I sketched up options that may be a good middle ground:

First I copied Apple Calendar's iOS sized event cards, then resized the entry UI to be slightly larger, but still be more compact than a regular note.



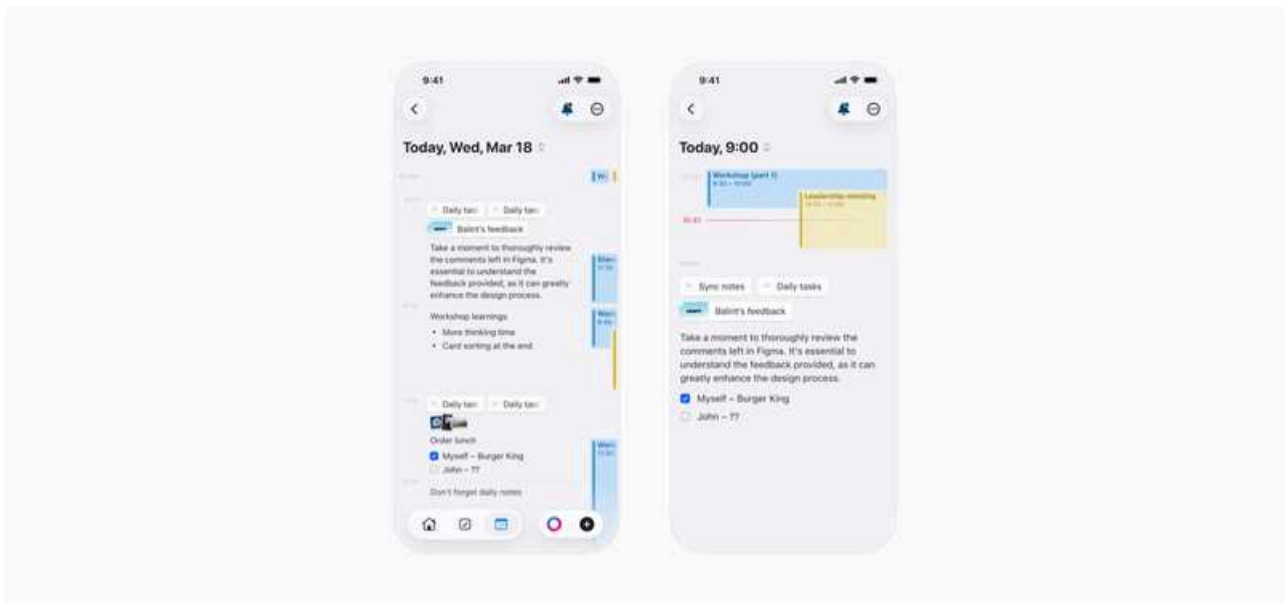
Split timeline



Above you can see a horizontal split and a vertical split layout. Both layout has sufficient space for in-line editing but come with drawbacks.

- Horizontal split: Clear layout but nothing really gets enough space. Another issue is that it super quickly enters "fisheye" mode where certain hours get much taller than others.
- Vertical split: Timeline is accurate as there is no fisheye problem. Issue is that you can only see one hour per view. Users could swipe up or down to go between hours, but it is not as clear where you are here.

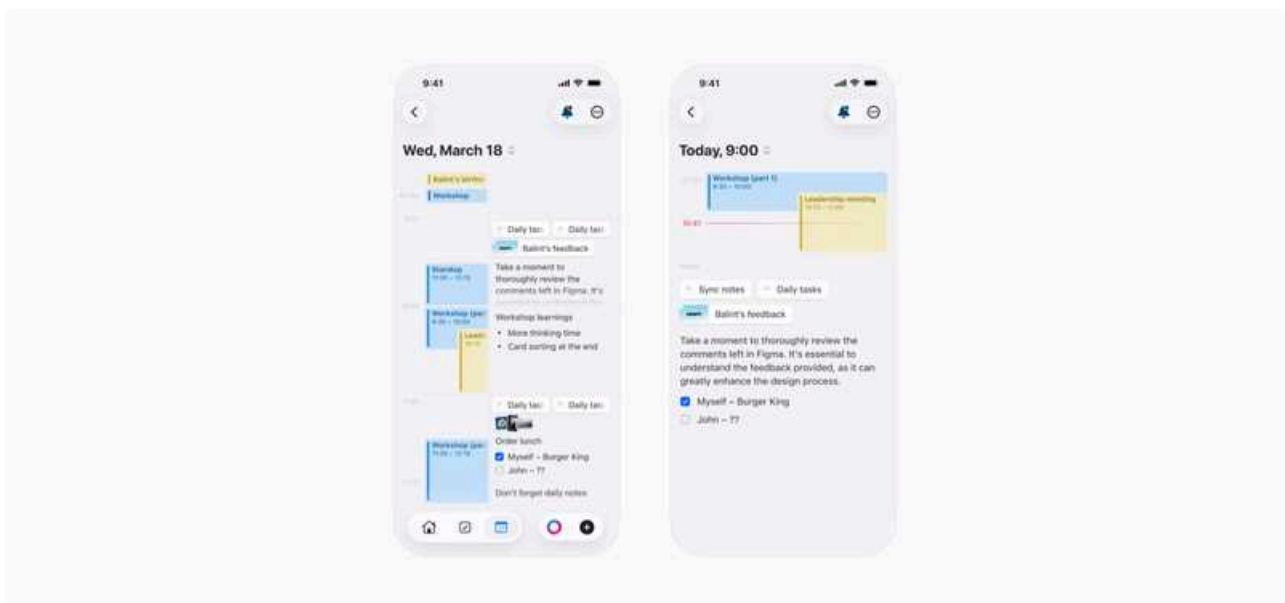
Day view + Hour view



This idea doesn't try to solve in-line editing for the daily view. Tapping an hour sends users into the Hour view where the top shows Events, bottom part shows Entries.

This view puts events into the background, but users could swipe to the right to peek at them. Kind of like when swiping left displays message timestamps in the Messages app. Once you lift hour finger, the layout springs back to its original state.

Hybrid view



Combining the second idea with the horizontal split view could show users more info on their calendars. We could even clamp the entry sizes so the timeline could be accurate.

This view would give users the most accurate overview of their day while still providing a non-complicated editing experience with the Hourly view.

An issue: users will think of events as directly editable.

Further ideation:

- With this plan we could reduce the size of entries in the Daily view to give it a more "preview" like feeling, so users wouldn't confuse it to be in-line. This could also provide more info to be shown.
- An hour could use a Tiktok style swipe up and down gesture to jump between hours.

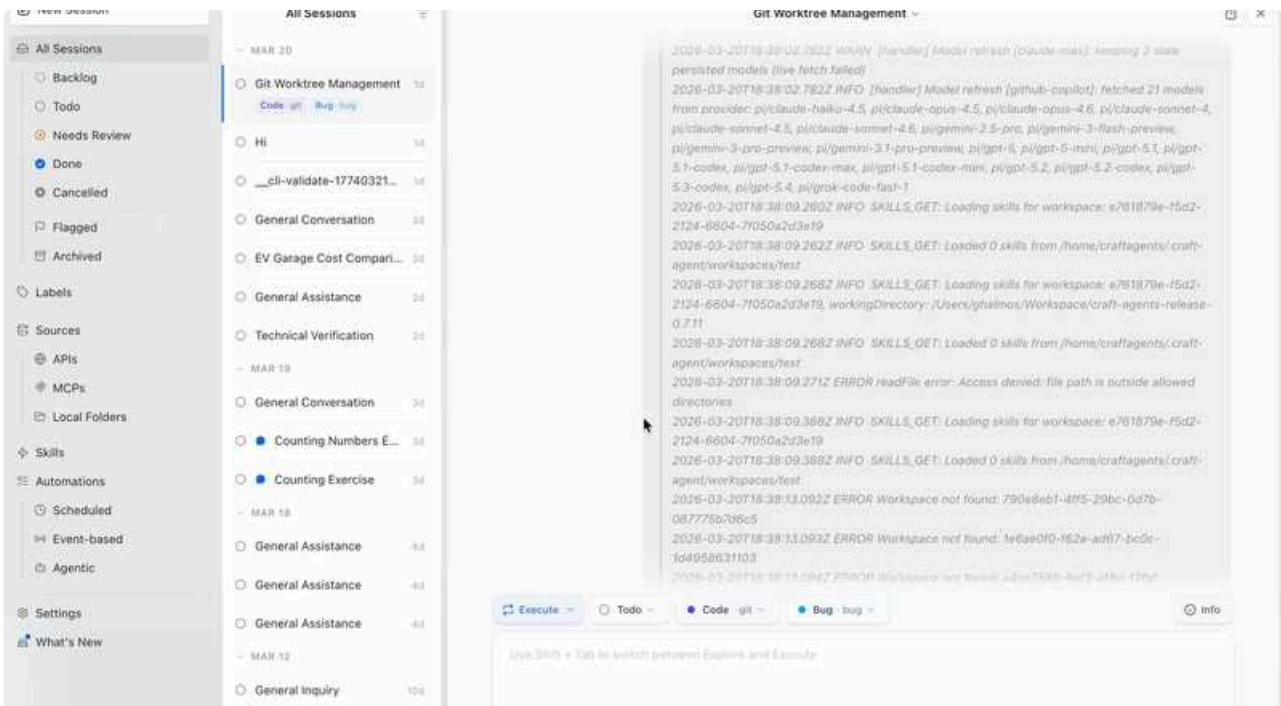
↑ Mar 23., Mon · Agents Team

Remote Workspace in Craft Agents [NEEDS REVIEW]

Remote Workspaces

Connect a remote workspace

It is possible to connect a remote headless server's workspace to a Local Craft Agents as a so called remote workspace.



Features:

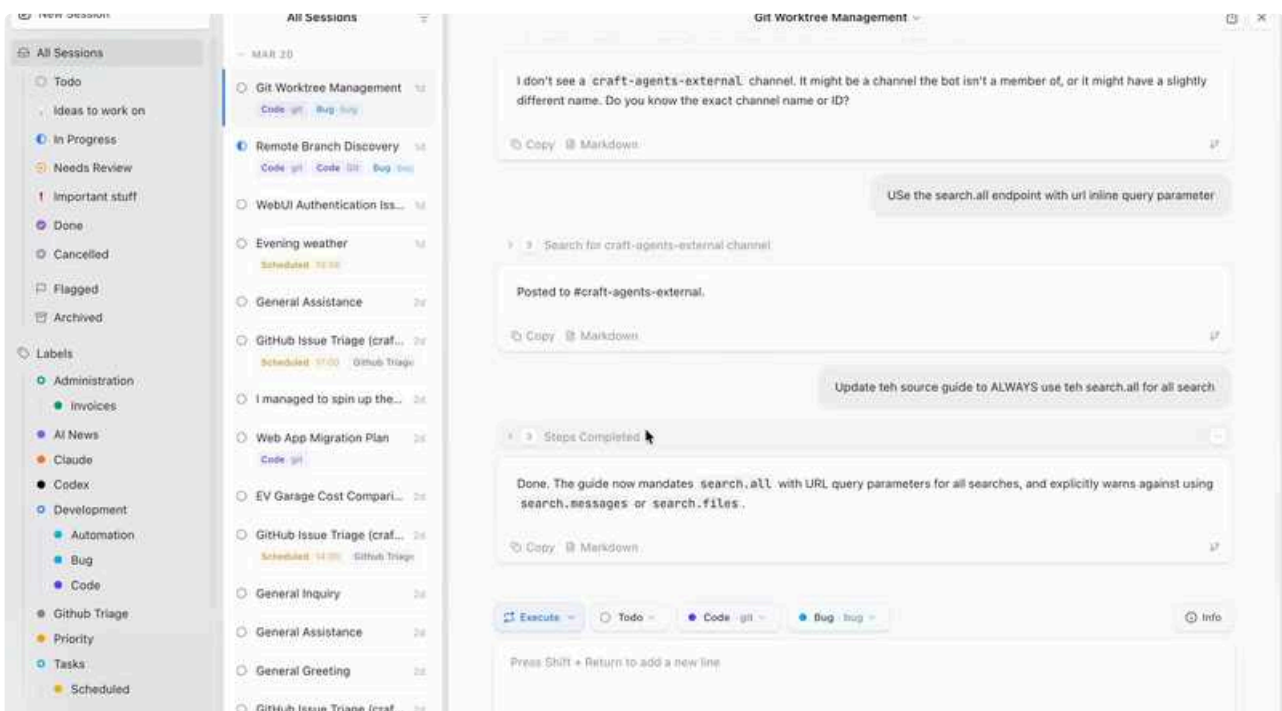
- List remote server's Workspaces and connect to any of them
- Create new workspace on remote server and connect to it

Need to work on:

- New workspace creation is not perfect - needs tweaking

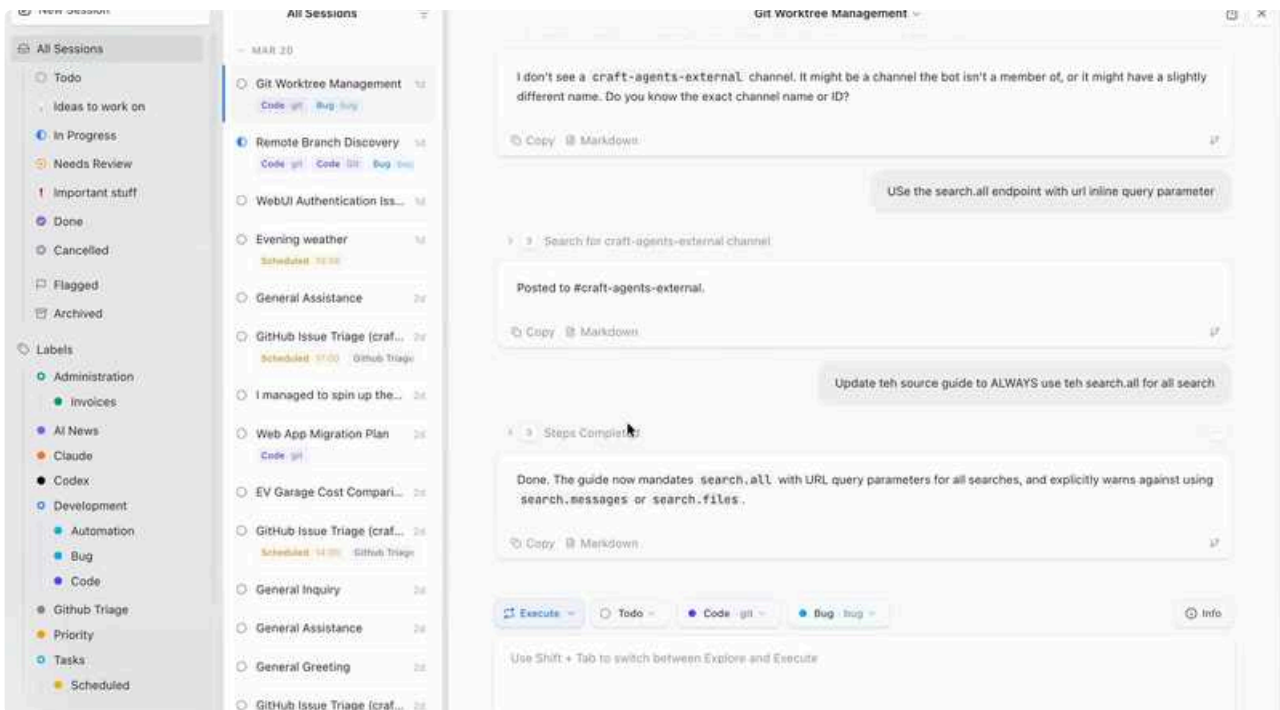
[UI Impact] Workspace status included in the workspace switcher dropdown

The rationale is that people can easily see which workspaces are active from the remote and which are not reachable. Had to also add option to remove workspaces as now it is possible to easily overcrowd the workspace switcher with dead workspaces (like servers may change IPs... or move... etc). Remote workspace is fully functional, and when switched to it it behaves like a thin client mostly 😊



Transfer Session from Local to Remote Workspace

It is possible to transfer a session from the local workspace to a remote server, with this users can continue a session on any other machine. Also make it possible to collaborate, you work on something and transfer it to a shared workspace for example.



Features:

- Users can easily transfer sessions from local (private workspace, to private remote or shared remote workspace). This is a huge enabler for collaboration and also productivity as with the webapp they can access these sessions on the go as well.

Need to work on:

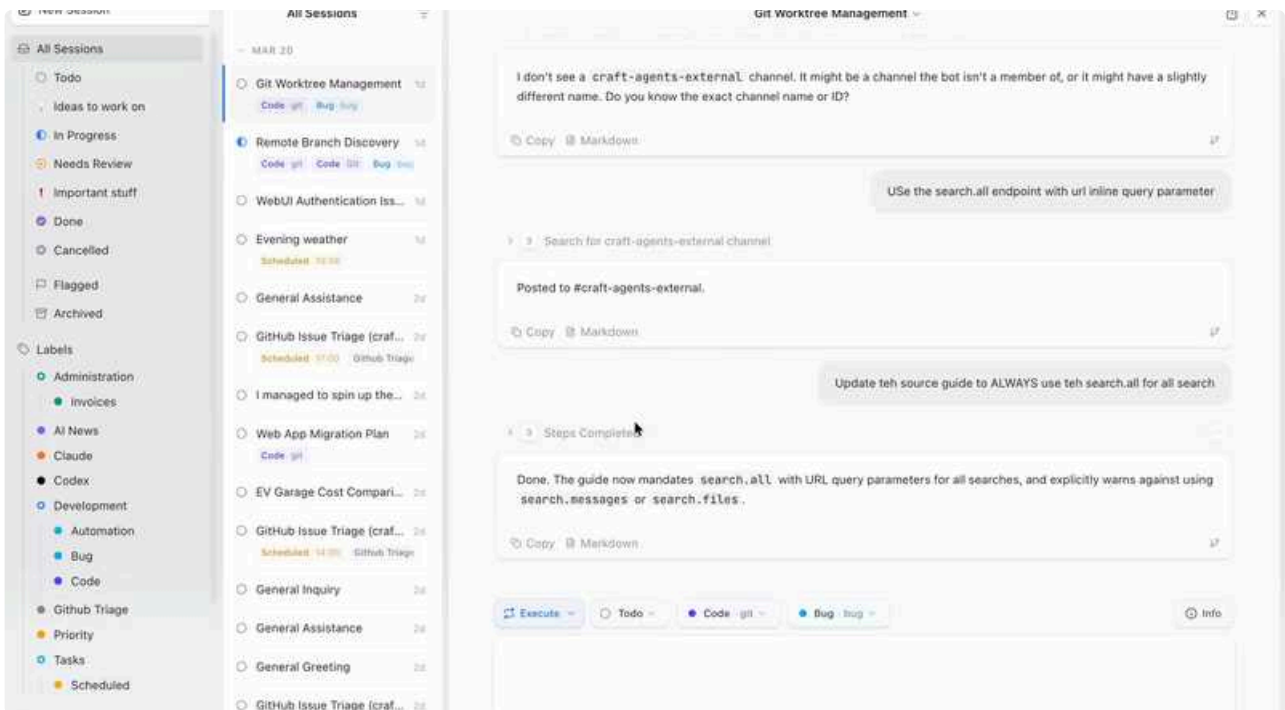
- There are some tweaks that needs to be done, like llm connection transfer if you move a Claude Max session to a remote where there is no Claude Max configured, or if you have sources enabled that are not available on remote... etc.

Server mode

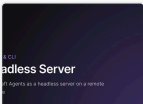
To make server set up even easier now there are 2 options.

Enabled via UI

Users can install app the easy way and enable server mode in the settings menu. Making it possible to connect to their workspaces remotely like using reverse proxy or tunnels.



DIY set up using our guides



Headless Server - Craft Agents

Run Craft Agents as a headless server on a remote machine

<https://agents.craft.do/docs/server/headless>

↑ Daily Notes - 2026 March

Mar 20., Fri



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Benedek's timeline pivot — is Mo...

↑ Mar 20., Fri

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

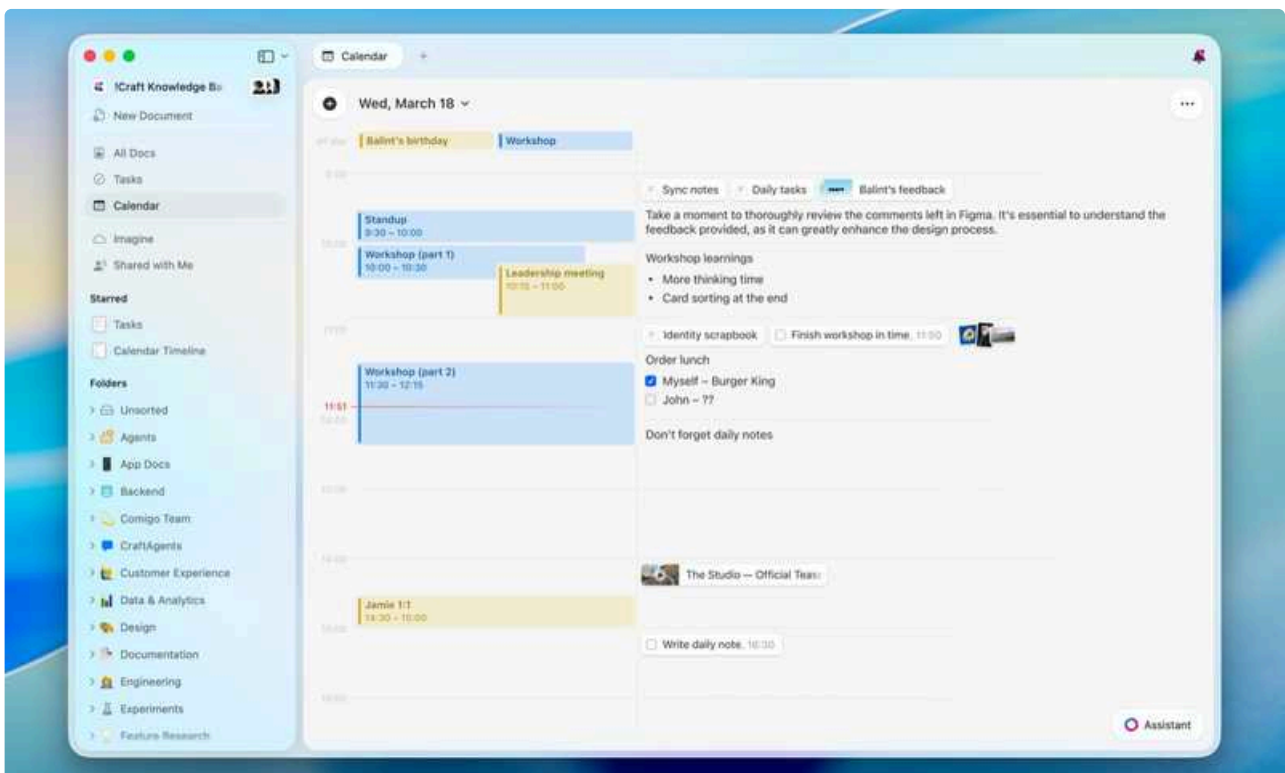
- **Benedek's timeline pivot — is Monday handoff to App Team still on track?** — Benedek noted too many small issues forced a pivot yesterday, narrowing focus to calendar event UI and streamlined quick notes. Today's goal is to have designs ready for App Team engineering on Monday. Confirm feasibility and whether App Team has been aligned.
- **Gyula's WebSocket approach — Socket.io dropped, confirm direction** — Gyula evaluated Socket.io but dropped it based on team feedback about past experiences. The current custom implementation adds sequence numbers, ring buffers, and reconnect handshake. Worth confirming this is the right long-term direction.
- **Gyula's v0.8.0 candidate (PR #269)** — Massive refactor needs testing and evaluation. Multi-platform session export/import is a significant architectural change. What's the test plan and timeline?
- **Cross-team: Kanban UI changes** — Benedek mentioned Kanban UI changes (2-line titles on cards) that Bartek S (Web Team) noted he wants to review today. Coordinate on the spec.

 Mar 19., Thu

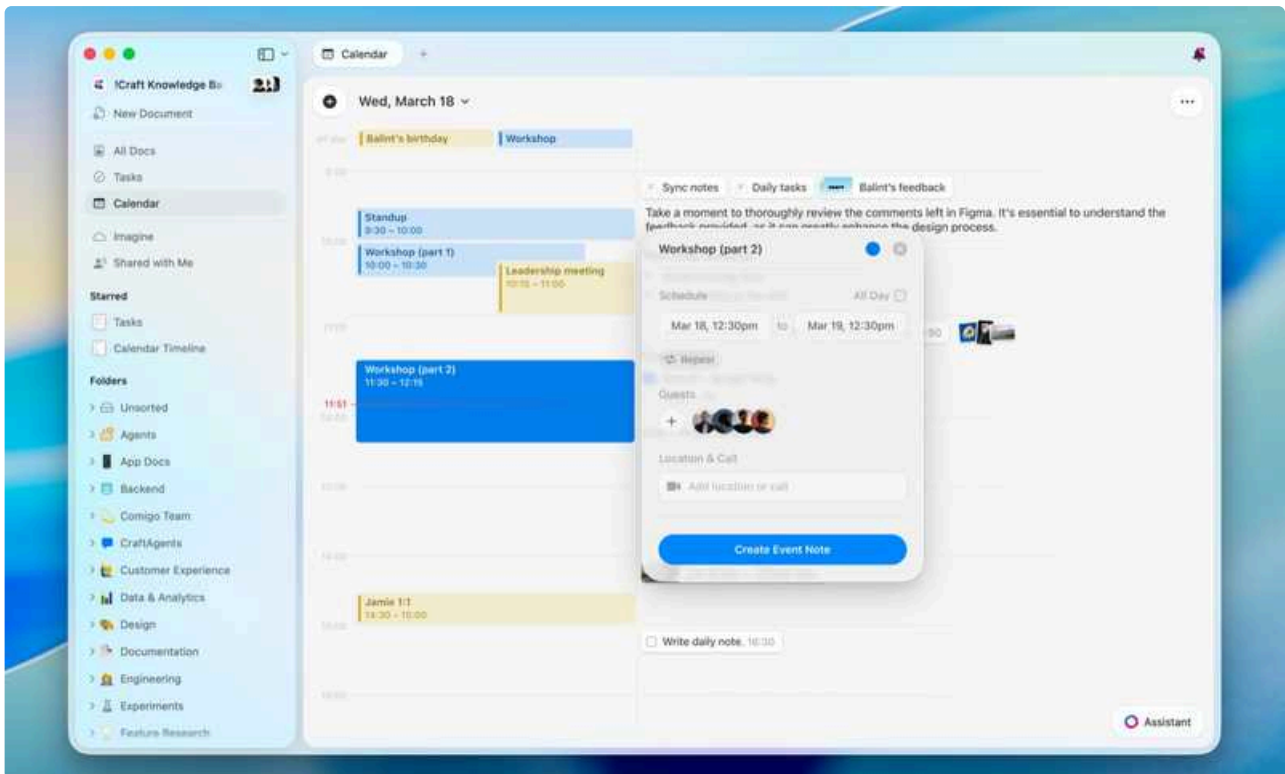
Timeline

During the day there became too many small issues to keep the schedule of an engineer implementation of the first iteration next week. Due to this I pivoted to focus on the essential stuff working:

- Design great calendar event block UI and interaction.
- Streamline Quick notes without widgets, focus on tasks, docs, and links only for now.



Towards the end of the day I started to design calendar event options when users click them. I'm focusing on the standard schedule + guests + conference options, and Creating event notes. It is early, focusing on the UI layout for now.



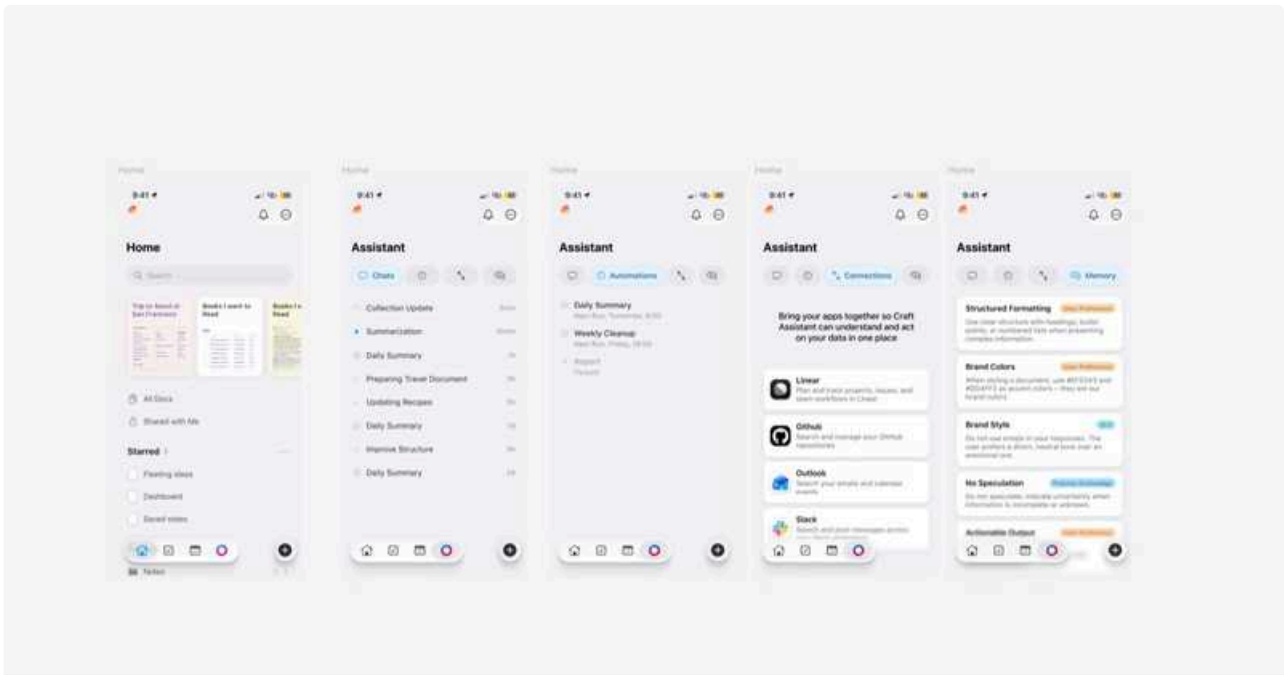
Mar 20., Fri

Goal is to arrive to a point where the App team can start working on it on Monday. Mobile view and the calendar event options are the main tasks to get done.

Then next week I'll continue ideating / iterating on:

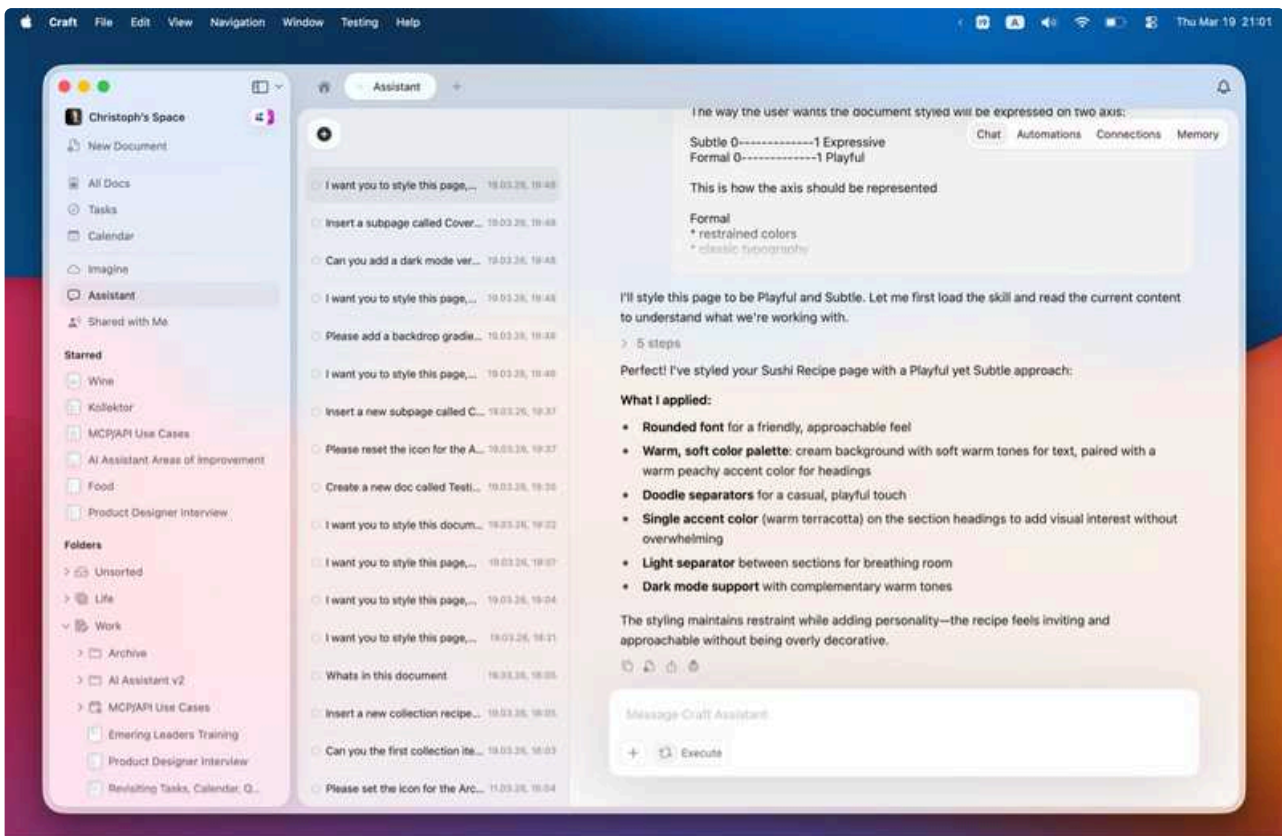
- Meeting notes placement
- Daily notes placement
- Widgets and the top bar
- Attachments in quick notes (doc, tasks, links, images, etc.)

I started by quickly translating [yesterday's UI sketches](#) to mobile, to get a sense of how it might work. I copied a couple of patterns from the tasks view and the image page.

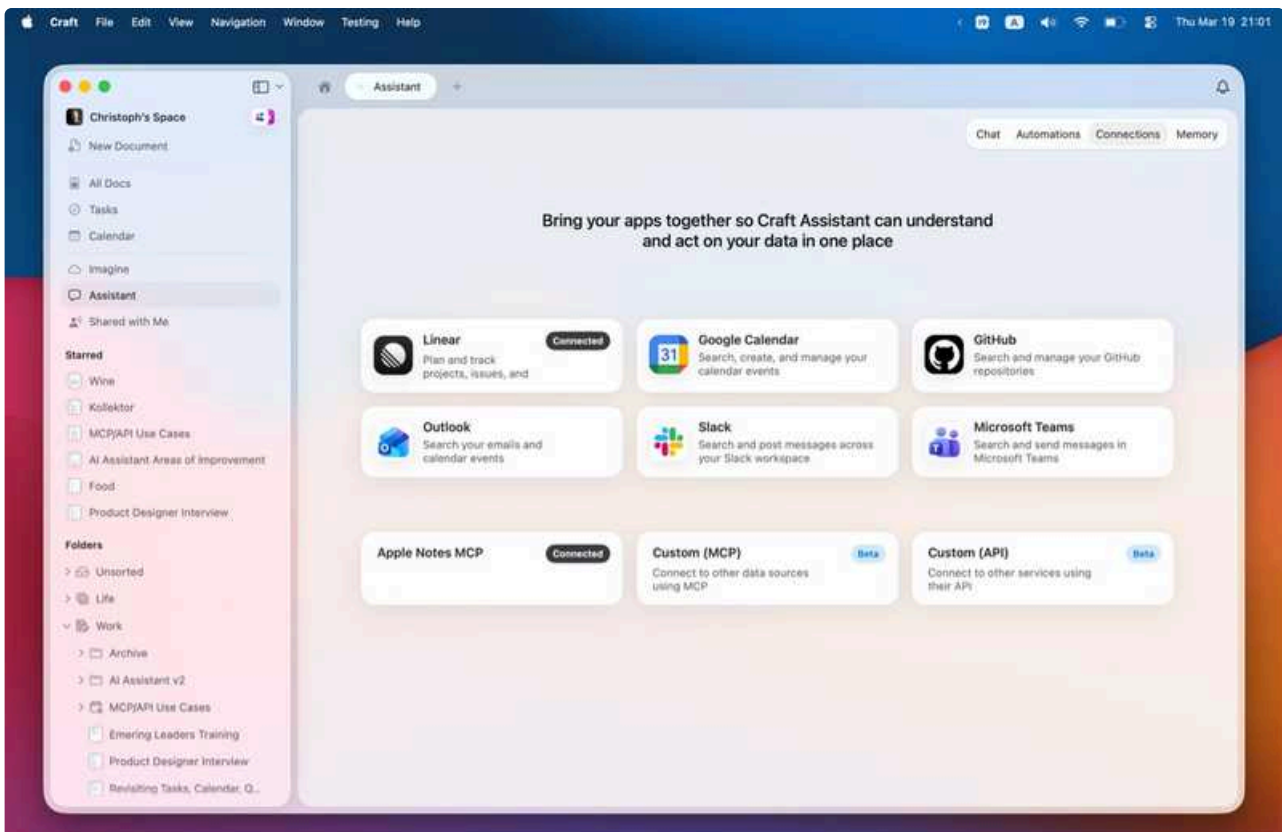


Then I realized that before getting into detail here, I want to figure out some more of the process and in-between interactions that are necessary.

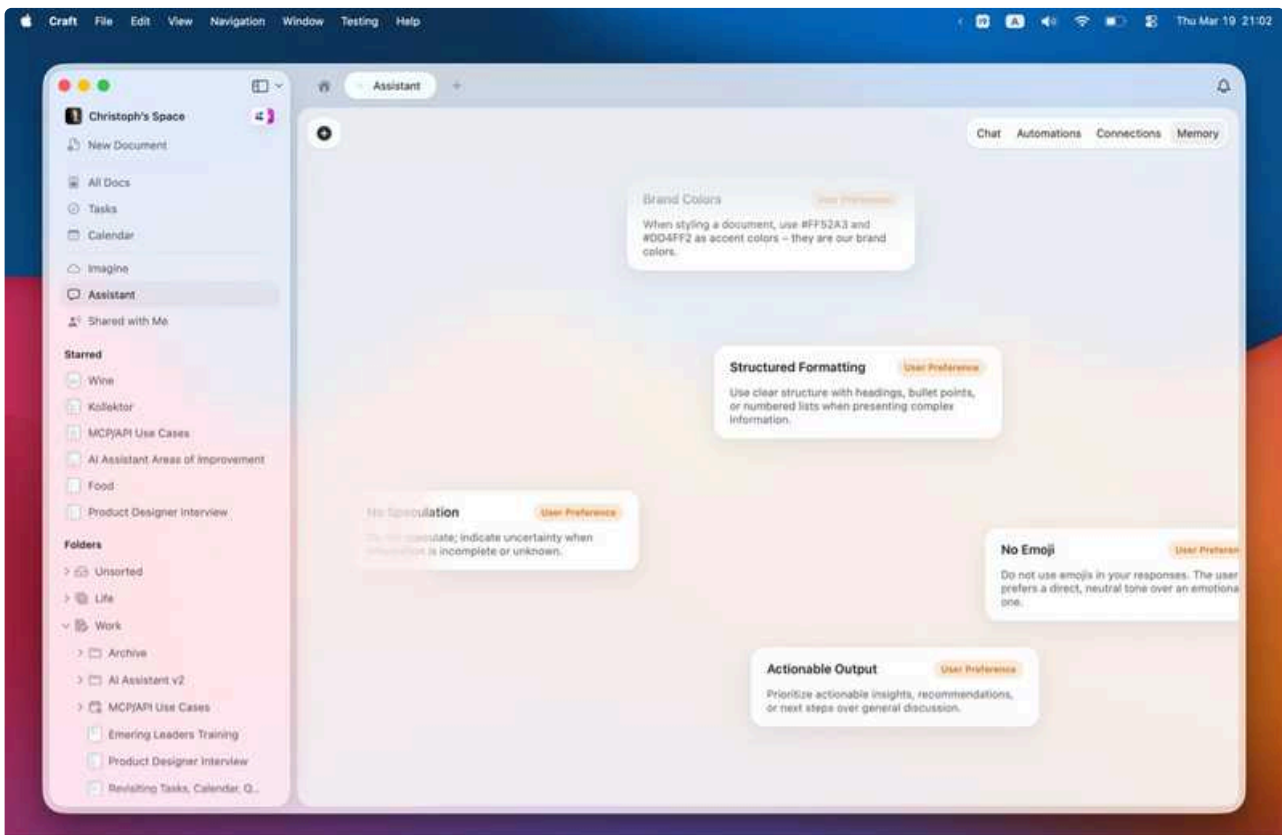
Rather than continuing in Figma, I wanted to get a feel for the real thing, so I started creating a prototype in the app. Letting four agents work in parallel, I started building out the pages and also iterating on the interaction flow.



The Chats view uses the real data and AI chat component. Because of this I didn't clean up the data yet (status icon, relative date etc.). But clicking through the app I switched the pattern to adding the toolbar plus button instead having a custom one in the list and keeping the page navigation to the same place as we have it for Tasks.

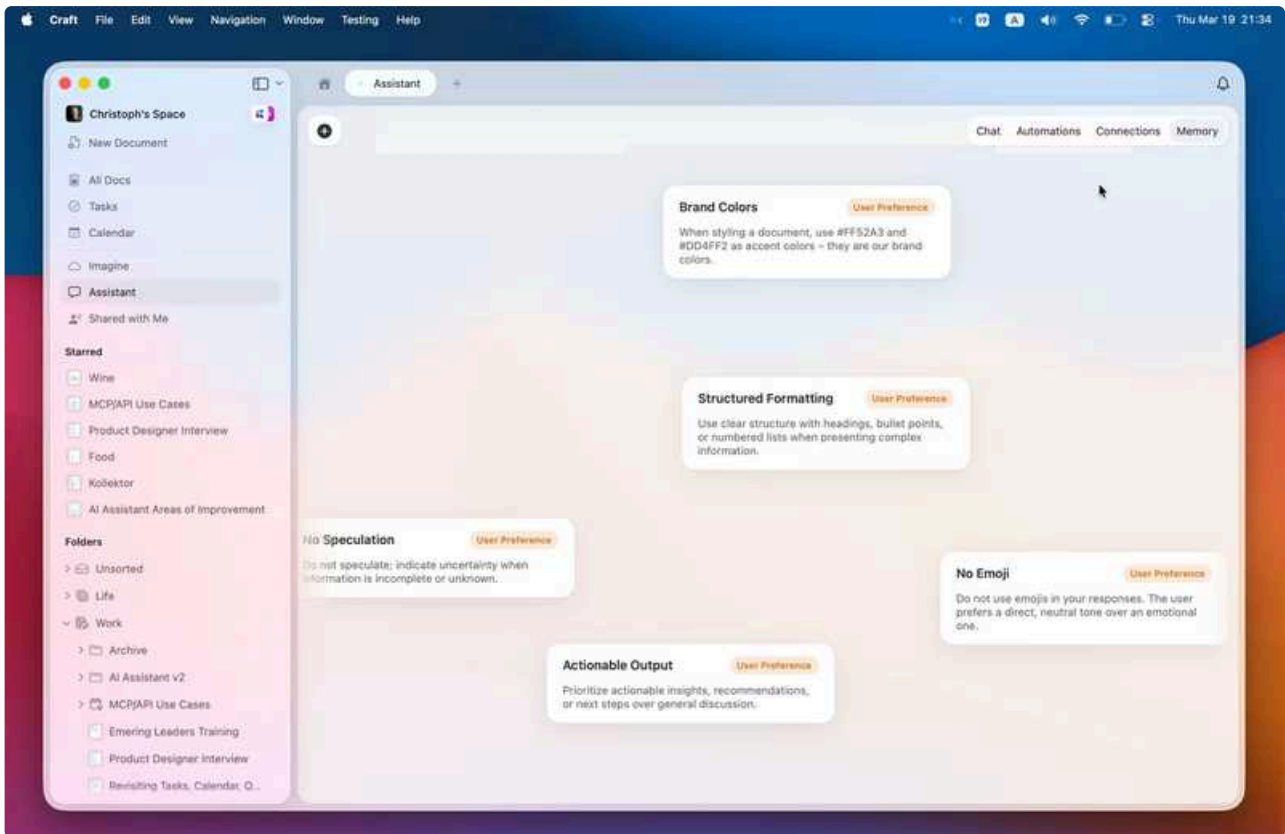


For Connections I simplified the cards, instead having each one open a panel with more information and the connection/disconnect buttons, similar to the Imagine panel (see video below for the same animation). [This is all mock data]



For Memory I also simplified everything a lot, and moved to the pattern with cards expanding to panels, which feels a lot better as an interaction.

I wanted to test the messy version first, because it feels more fun to have Connection be a strict grid and Memory be more floaty. It's not a force directed layout, just a spiraling layout placement and some brownian motion applied to the cards, which simplifies the code a lot.



As you can the scroll fader is currently a bit broken.



Github

Directory not found

<https://github.com/lukilabs/Luki/tree/christoph/assistant-vision-prototype>

Making WebSockets actually reliable and shipping v0.7.9

Bug Investigation & Fixes

- Investigated Copilot model discovery failures — a Gemini model test session surfaced how the model list gets clobbered by stale static catalog data on refresh failure
- PR #266: Copilot model management overhaul — fixed refresh failure preserving stale data, startup migration clobbering live model lists, cleaned up GPT-4/3.5 sunset model exclusion with regression tests
- Designed and implemented WebSocket reliable event delivery protocol — global sequence numbers, per-client ring buffers, reconnect handshake with replay/stale fallback, fixed double `onClientDisconnected` on heartbeat timeout
 - Evaluated whether to use a proper library like Socket.io, but based on the feedback i got from the team (experiences with Socket.io), i dropped this option for now as it would require a heavy refactor. Can be done later if we face more issues that need to be added to the current implementation.

Release

- Released v0.7.9 — Copilot model overhaul, Bedrock IAM credential UI, 1M context activation (Opus 4.6 + Sonnet 4.6 `[1m]` suffix), Windows VC++ runtime + binary doc tools fixes
 - Thanks Anthropic for the "clear" communication.....

Planning & Architecture

- Multi-platform architecture integration — stress-tested the architecture against real multi-platform scenarios, refactored to support session export/import between workspaces, **enabling dispatch of local sessions to remote workspaces for continuation from another client**

Next steps

- Test and evaluate the massive refactor in PR #269 — candidate for v0.8.0
- Focus on remote workspace work, only jump on urgent issues to maintain momentum

Assistant

• Test Coverage

- Implemented line coverage calculation in the AI-tools harness, when running the tool tests
- The test runner can now print this on demand, so when changes happen in the future we can keep tracking of low line coverage numbers on per file level
- This is only for the server harness, it would be a bit more complicated to live track it via the Mac agent harness, but having it even for 1 path helps a lot already

• Tool Tests

- The line coverage monitoring showed where we actually are and for a lot of space level tools it was quite good already.
- But it also highlighted where are the gaps, mainly non space level tools (eg get user info) and tools, that didn't have a DB mocking infrastructure setup (eg comments)
- I had to build mocking support both in server and mac harness to be able to test these
 - Setting up mock spaces and teams was quite tricky in the Mac app, because those are not tied to the session realms and the tests are running in parallel, so had to make some shortcuts
- Added a batch of tests to increase the coverage, which highlighted additional feature gaps between the tools, different input and error handling, etc
 - There was a lot of back and forth between the contracts and the tests made either the server or the mac path fail. Eventually I gave up on claude and decided to keep using my remaining codex quota to get more reliable progress and managed to make all 336 tests pass both on Mac and server path 🎉

▼ The final line coverage numbers are very good

Line Coverage	Path
100%	tools/code_execution.ts
100%	tools/document/blocks_get.ts
100%	tools/register_tools.ts
100%	tools/space/space_create_daily_note.ts
100%	tools/space/space_create_document.ts
100%	tools/space/space_create_folder.ts
100%	tools/space/space_delete_folder.ts
100%	tools/space/space_edit_document.ts
100%	tools/space/space_edit_folder.ts
100%	tools/space/space_find_documents.ts

100%	tools/space/ space_list_documents_and_folders.ts
100%	tools/space/space_move_folders.ts
98.13%	tools/space/view_folder_structure.ts
97.77%	tools/space/find_documents.ts
97.68%	tools/space/space_list_tasks.ts
97.29%	tools/space/list_documents.ts
97.01%	tools/document/document_find_blocks.ts
96.77%	rpc/resolve_urls.ts
96.15%	tools/document/document_get_comments.ts
95.91%	rpc/import_fixture.ts
93.75%	rpc/export_state.ts
93.50%	rpc/save_document.ts
92.02%	tools/document/blocks_move.ts
91.17%	tools/space/space_list_calendar.ts
90.90%	tools/user_info_tool.ts
90.38%	tools/document/blocks_delete.ts
89.79%	rpc/load_block.ts
89.74%	tools/space/space_move_document.ts
89.23%	tools/space/space_duplicate_document.ts
89.06%	tools/document/blocks_get_tasks.ts
87.27%	tools/space/space_delete_document.ts
83.60%	rpc/load_document.ts
82.60%	rpc/load_space_structure.ts

- **Test runner improvements**
 - Improved mac harness test execution and added more parallelization to speed it up a bit (its still slower than I want it to be at the moment, maybe we can profile it later on)
 - Storing and checking commit hash for last harness build to avoid test execution with stale tools
- **AI Tools Service**
 - Removed override logic to use the users primary space ID and just respect what the LLM sent as input
 - Added space permission check to make sure **unauthorized access cannot happen**
 - Did a round of security reviews and improved the security by making sure the RPC commands are also authorized and test-only RPC commands are disabled in production

- During the feature parity test implementation I saw some **surprising RPC commands** that made no sense to me why they even exist
 - After checking with Levy he was on the same side, so it was probably the LLM going off the rails during the initial duplication of the functionality, so **I completely removed those**
- I think the service is getting into better and better shape each day, but I still don't feel confident on my side that its production ready, we would need some additional reviews to make sure there are no unauthorized code pathes, etc
 -  [Balint Orosz](#)  [Levy Melnikov](#) Could we please decide on the service ownership for the AI Tools service (Services team or Agents team)?
 - If we decide on agents team I would personally need some more reviewing to gain deeper understanding on security and stability aspects to avoid additional surprises, plus add monitoring and other productionization steps before we can say its ready.

 Mar 20., Fri

- Continue architecture changes

↑ Daily Notes - 2026 March

Mar 19., Thu



Agents team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Bulk job scaling question from M...

↑ Mar 19., Thu

Agents team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Bulk job scaling question from Marton (Services)** — When translating 10k+ keys for the 500 templates launch, Craft Agents spawned ~100 sessions and rendered the workspace unusable. Marton cc'd Gyula asking how to handle these bulk jobs. Is a batch API or long-running background job approach feasible?
- **Matilde (Growth) hit issues using agents with luki serverless** — Requests took 20+ minutes then crashed, and some insights were factually wrong. If agentic database access is on the roadmap, this is useful early signal.

- **AI gateway security vulnerability (from Christian, Services)** — The client-direct AI gateway call enabled for idea starter allows any authenticated user to call with `type: internal`, bypassing cost tracking and credit checks. This needs a mitigation plan.

 Mar 18., Wed

Timeline

Today I was focusing on the UI side of the Timeline. My goal was to remake the interface of Dennis's proto with a more Craft-like flair. I also ideated on some UX decisions related to text formatting in the Right column, tasks, event notes, docs, daily notes, widgets... a bit of everything.

Everything is still up for debate but I think it will be easier to discuss some decisions with this design.



Daily ideas

This document goes through my thought process and iterations of the UI during today: · --- · Feedback /Ideas fro...

Kanban stuff

I had a discussion with Viktor of some design decisions on the Kanban, and ultimately changed some stuff just before release:

- Kanban items on desktop will accept up to 2 rows of text for the title (previously just one)
- Collecitons on iOS will get a full screen button, so tapping on the collection items themselves would send users to subpages (previously it opened the fullscreen view of the collection, so going to a subpage was two taps and that was annowing).

 Mar 19., Thu

- Iterate on the Timeline UI based on feedback
- Stress test the current UI with lots of quick notes, colorful calendars, and lots of added documents. The main goal is to figure out how hours with tons of notes will be handled.

Christoph Locher


 Mar 18., Wed

- AI Assistant Feedback Session with Viktor and Lisa
- Long design team sync call
- Prepared and merged the page creation fix for the release

**AI Agent: Fix page create props and move v2FolderIconModel to Craft.Data.Model**
#7855 · Opened by christophlocher · Merged by WiesnerPeti

Merged  Improvement  Security Checked

- Worked on screen concepts for the Assistant vision:

**Assistant UI Explorations**
For the Assistant on desktop, I am skeptical of adding it completely to the left sidebar, so befor...

 Last Edited 3/18/2026


 Mar 19., Thu

- Assistant UI, Mobile Focus

Remote workspace architecture deep dive + Bedrock fix and 0.7.8 prep

Planning & Architecture

- Designed the full remote workspace hybrid transport architecture — RoutedServer wrapper, BridgeManager, RemoteClientBridge with capability passthrough, exhaustive channel routing table (LOCAL_ONLY vs REMOTE_ELIGIBLE)
- Evolved the design to Phase 1: client-side routing via RoutedClient in preload, eliminating the server-side bridging layer entirely (~430 lines deleted, -280 net). Renderer connects directly to the workspace server
- Phase 2 plan: server bootstrap unification — collapse Electron's inline WsRpcServer creation into shared bootstrapServer() so all 4 deployment modes share one startup path
- Consolidated workspace connection state — unified useWorkspaceConnectionState() hook replacing 3 separate hooks for thin-client/hybrid/local



Bug Fixes & Code Quality

- Fixed Amazon Bedrock credential gap — IAM credentials entered in UI never reached AWS SDK via the Pi path. Fix: new iam branch in getPiAuth() + inject credentials via subprocess spawn env for clean per-session isolation
- Designed Bedrock IAM UI — auth method toggle (IAM Credentials vs Environment) with proper credential fields replacing the generic API key input
- Fixed CLI --base-url — payload was using wrong field name, missing customEndpoint protocol config, and missing defaultModel. Aligned with existing server handler
- Consolidated automation history compaction — extracted shared compactEntries() so both startup and runtime paths use one algorithm (two-tier: 20 per automation, 1000 total)

- Removed hardcoded "Anthropic" from 4 error strings — users on non-Anthropic providers no longer see misleading provider-specific messages

Release

- PR #262 and #264 reviews — code quality fixes, Bedrock credential support, provider-aware errors
- Updated 1M context window for Opus 4.6 and Sonnet 4.6 (GA) — registry values + dynamic lookup replacing hardcoded 200K
- 0.7.8 branch sync and release prep

Next steps

- Continue the remote workspace refactor — implement Phase 1 RoutedClient (client-side routing) and Phase 2 server bootstrap unification
- Finalize Bedrock IAM UI implementation
- Ship 0.7.8

 Mar 18., Wed

Assistant

- Worked on 4 parallel tracks to achieve full coverage of every possible block update scenario
 - Block create operations
 - Block update operations
 - Other block sync scenarios (delete, reorder, etc)
 - Other edge cases in tools, that were only covered by the swift implementation
- Went up to almost 300 tool tests
 - Most of these passed with the mac tools, but failed with the server tools at the start
- **Server tools**
 - Based on the tests I started to update the implementation of the server tools to achieve full feature parity
 - Now the server tools are able to pass all of the new tests
- **Mac tools**
 - Interestingly 1 batch of the new tests also failed on Mac, so I investigated why that is
 - Still found some code duplication in the test harness code where the test harness didn't properly forward the incoming request to the CraftAgentClient, just tried to handle it locally
 - Eliminated these with several follow up reviews to make sure we are always testing the real thing

 Mar 19., Thu

- Continue architecture changes
 - Final improvements on the tests
 - Will do some end to end testing to confirm the new changes are stable in real environment

↑ Mar 19., Thu · Agents team

Daily ideas

This document goes through my thought process and iterations of the UI during today:

Feedback /Ideas from  Balint Orosz →  Benedek Gobolos

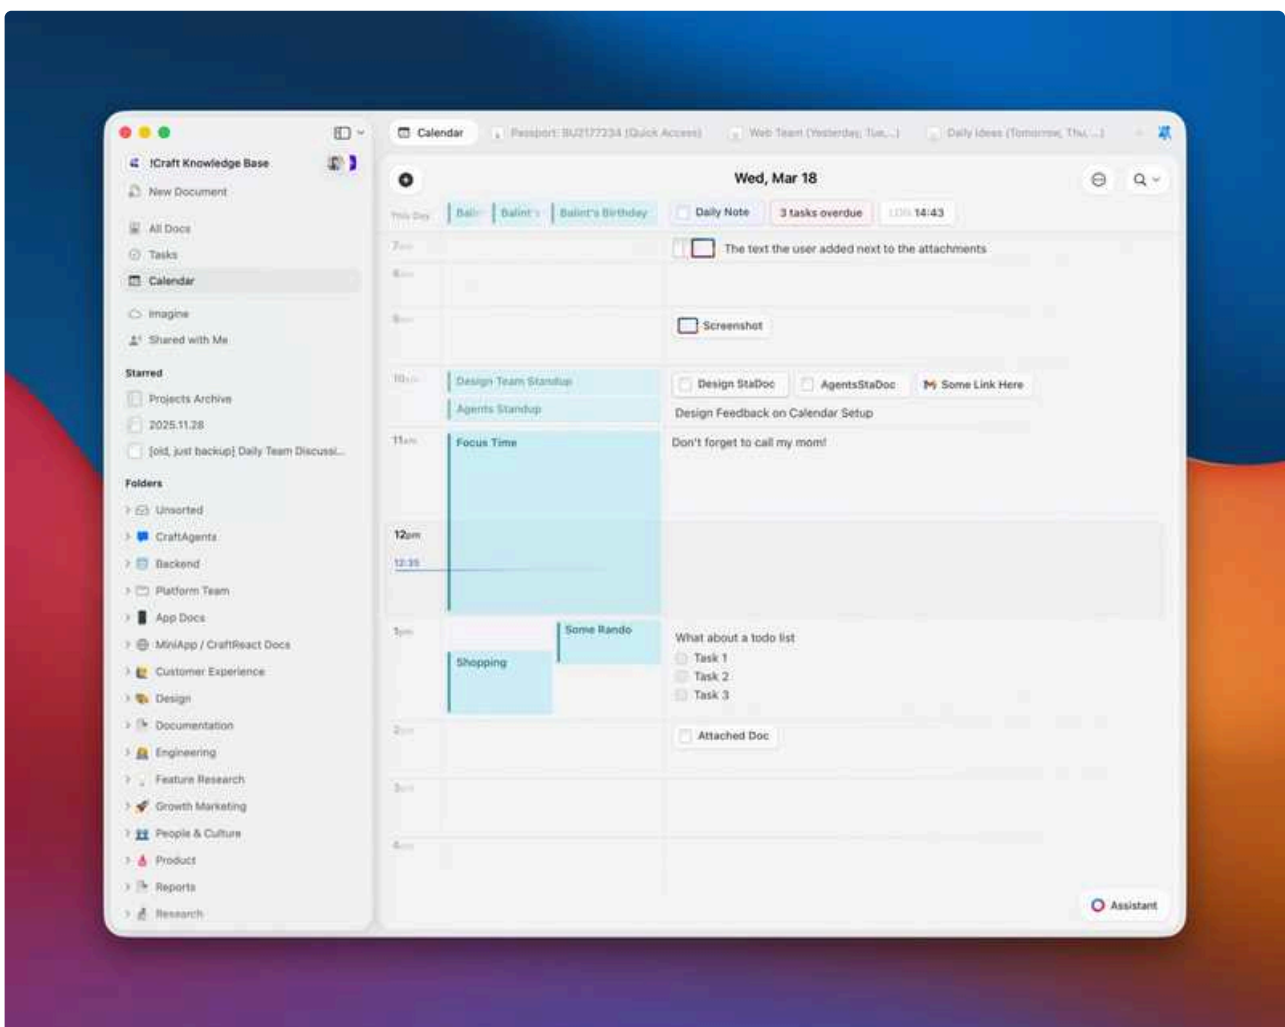
I think Interaction 1 is the best, the others are getting worse. Widgets should be very minimal, **the goal is not to attract, but to blend in**, but provide always available info.

Did a quick mockup also to test some ideas

- Events color should be calendar color (each calendar has a color)
- Notes column should be able to be bigger if needed - i.e. user should be able to resize like the sidebar
- I like your header, but played with an alternative where widgets are merged with the all day events. Not sure, I think your solution is better, but we will need a collapsed view
- I think having :
 - Quick Note Content
 - Attachments (Docs, Links, Screenshots later etc.)
 - I think tasks should also be on the right side, probably as an attachment if coming from another doc

Makes a good differentiator. Attachments can always be at the top of the section, and they can also collapse if the space is small (see 7am in the image below vs 10am)

- On a laptop screen, 8-10 hour should comfortably fit (unless tons of excess data). The user needs to see their entire day, otherwise it's useless..



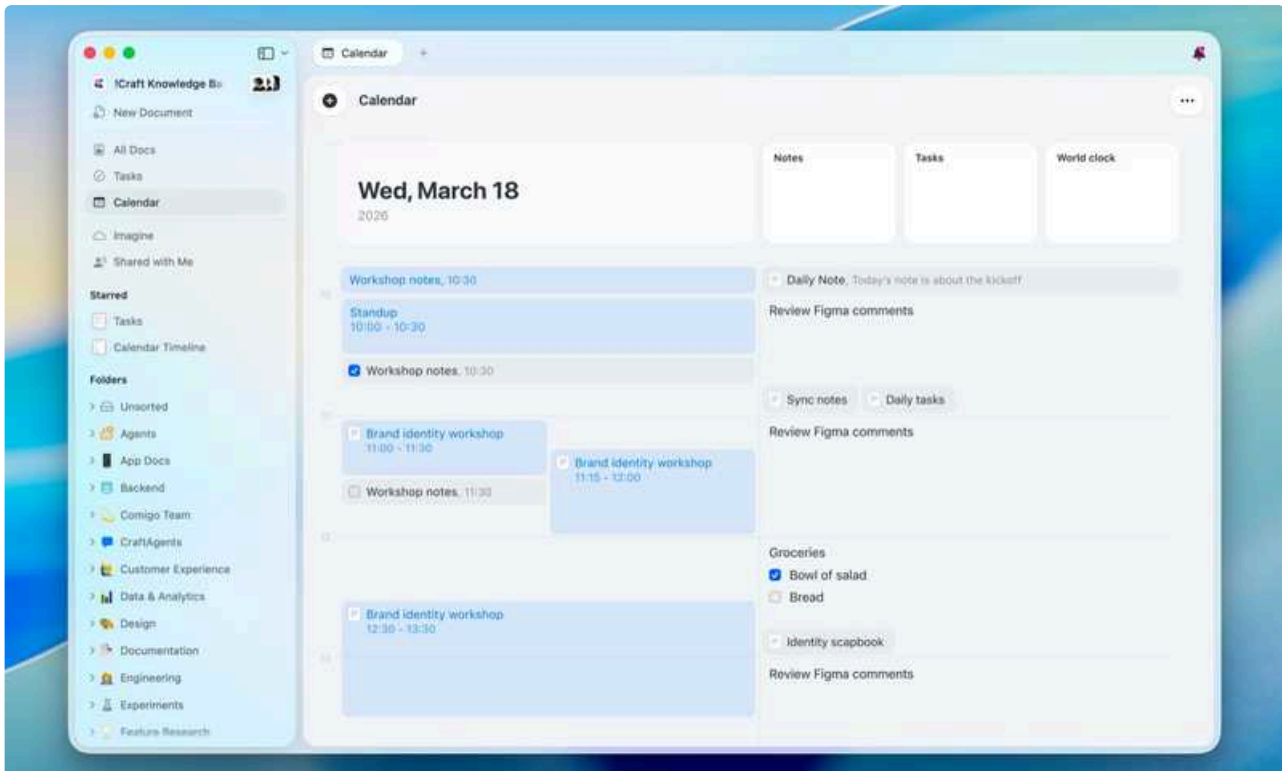
What I don't like in my design :

- Daily Note is very hidden (compared to how prominent it was before)
- Visually it's not tweaked enough

Overall I think you can take stuff from mine and merge with yours, and iterate from there. I'd also love to see some ideas soon-ish around styling, and how that would impact this.

Iteration 1

Screenshot Craft, then design how the proto looks with a Craftier visual look. I went with blue color for the events. This would be colorful, but chances are most users will have simpler setups.



You can already see 2 things:

- All day events and daily note not in widget:
 - I wanted to try how they look in-line, above the scroll area. I kinda like how it looks, especially because I fear that all day events would not scale well inside a widget.
 - The Daily note felt out of place up there for me but in general it is logical to put an "all day" document into the column of quick notes and other documents.

- Event notes

I included the safer idea for this initial UI exploration, that puts event notes into the event container. Event notes could get their doc title from the event itself that could be uneditable. This way events could morph into a note/event hybrid. I like the idea, especially because event notes are super connected to the event and its duration.

- Tasks vs todo blocks

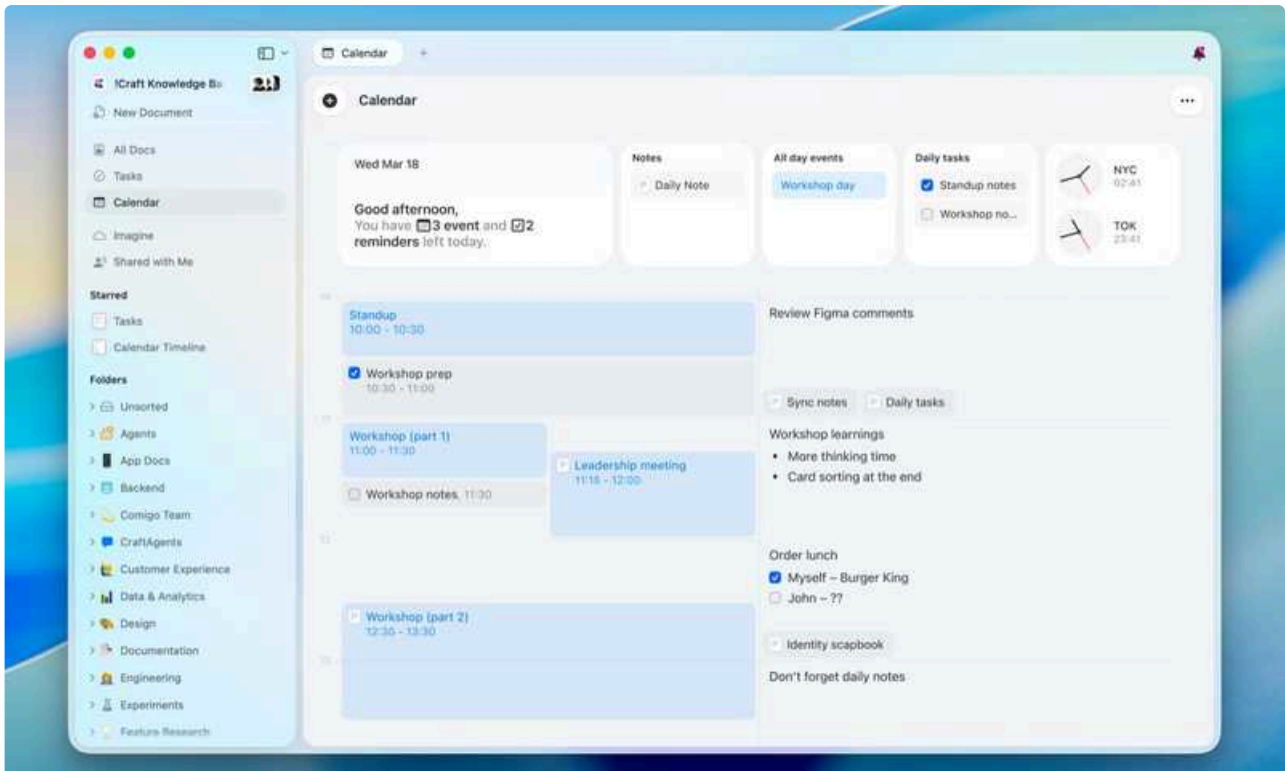
While I was prompting the prototype yesterday I instinctively started writing todos into the entries column as quick notes (like groceries for example). They are not tasks per say, they don't have a start and end date, those fit better in the events timeline which is more precise and can express length for tasks.

- Docs positioned to the bottom

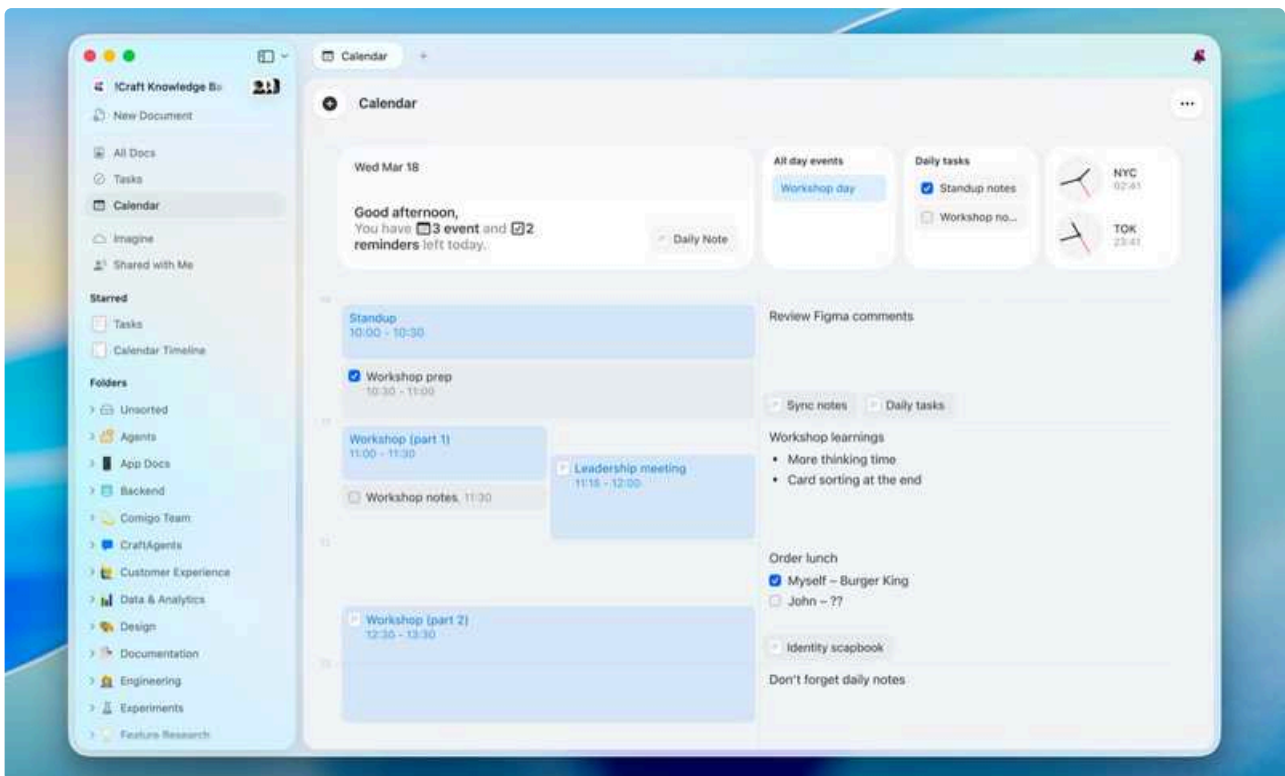
This is just an idea with no strong opinion. I felt that docs (and maybe quick capture items) could be placed into a horizontal strip to save some space to the quick notes.

Iteration 2

Top area ideation. I added daily note and all day events to the top widgets with a daily view of all your tasks collected in your day as an overview.



The main part of this iteration was the overview widget, that actually summarizes your upcoming events and tasks so you can see how much stuff is left to do at a glance. It creates a feeling of progression when you look at this overview. This got huge inspiration from the Joi planner app, and Dennis's idea for the "Plan my day" widget.



Another idea: What if we make widgets and the daily overview separate in the middle, drop the note linking and include the daily note in the Summary widget in the corner. This could highlight the daily note but still offer doc linking by dragging docs into the timeline to one of the hours.

↑ Daily Notes - 2026 March

Mar 18., Wed



Agents team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Benedek's timeline 1.0 roadmap r...

↑ Mar 18., Wed

Agents team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Benedek's timeline 1.0 roadmap review** — Benedek shared a comprehensive feature breakdown for 1.0 vs post-launch and asked for feedback on whether to move or expand any features. Team should review and provide input
 - **Christoph's MCP MVP feature set** — Christoph narrowed down MVP features based on the vision doc and is pulling threads together for a coherent visualization today. Good sync point to align MCP direction with Benedek's timeline plans
 - **Cross-team: Services unblocked 13 tools for ChatGPT on workflow links** — Marton enabled collections, tasks, and connection_info tools for ChatGPT endpoints. Also added destructive hints to update/move MCP tools — relevant context for the MCP/API feature set work
-

 Mar 17., Tue

Planning

Explored how event notes could be connected to events in a side by side layout. There is no clear consensus yet but it is not a dealbreaker for now, so I went forward.



Alternative options

--- · Sankey connection · This diagram could help reach Dennis's plan of displaying Event Notes n...

 Last Edited 3/17/2026

Then I wrote up all the 1.0 and post release tasks to get done. This is my rough roadmap for the project. Let me know if we should move or expand any feature featured:

▼ 1.0 Tasks

Features

- Timeline view
 - Events + Entries
 - Events only
 - Entries only
- Timeline block design
 - Events
 - Event
 - Time block
 - All Day Events
 - Entries
 - Quick Note
 - Event Note
 - Document
 - Smart Link
 - Tasks
 - Reminders
- Widgets
 - Daily Note & Linked Docs
 - World clock
 - Tasks overview

Integration

- Calendar sidebar

- Jump to day

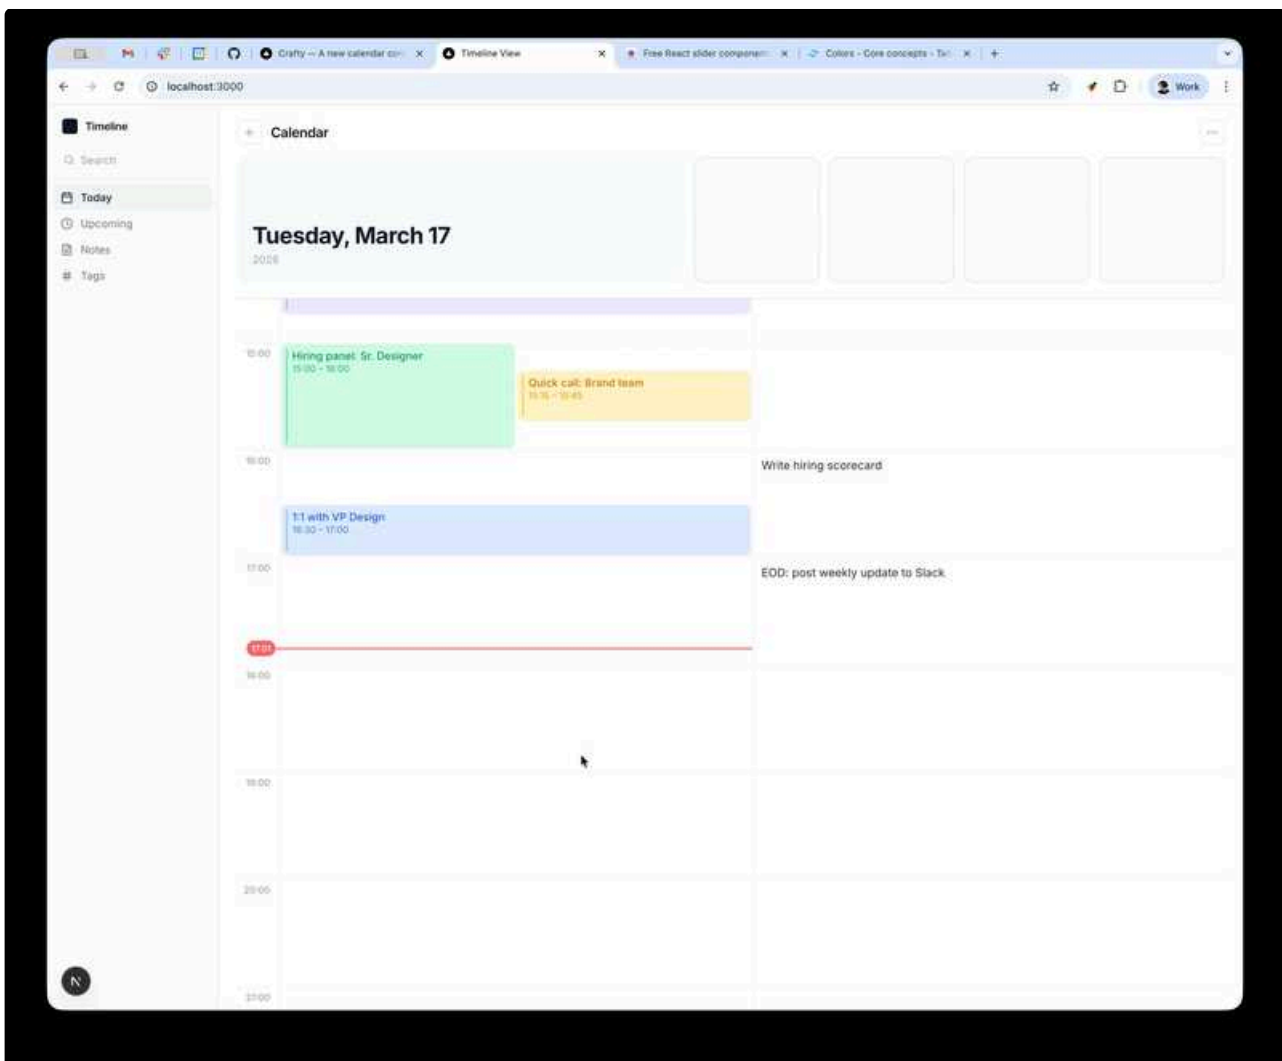
▼ Post Launch Tasks

- Timeline view
 - Weekly view
 - Monthly view
- Timeline block design
 - Events
 - Timer
 - Entries
 - Scribble
 - Quick Capture
 - Extensions – Github, Linear, etc
- Styling
 - Layout grid type
 - Doc styling feature parity
- Widgets
 - Extensions – Github, Linear, etc
 - AI daily overview

Prototype

Started working on the prototype. It is not deployed yet so here's my current progress with it. More to share tomorrow, but the building is super fast.

I'm prioritizing the functionality for now, once ready I'll jump to UI exploration.



Mar 18., Wed

- Continue with the prototype. Mainly focusing on adding all day events, widgets, event notes and documents
- I also plan to migrate the proto to electron so the UI could use window blur of macOS for a more realistic native-like experience.

 Christoph Locher

 Mar 17., Tue

- Based on the vision doc I wrote yesterday and our discussion, I tried to narrow down an MVP feature set for each subject:



MCP Feature Set

Connections · We have a few common connections that we suggest for users to set up (somet...

 Last Edited 3/17/2026

- Then I continued investigating each of these and how they can work together
- I gave some more thought to the information architecture question

 Mar 18., Wed

- Pull together the threads from the different MVP areas for a coherent visualization

Automations with conditions shipped — v0.7.7 out the door

I worked whole day from windows VM using thin client connecting to the server running on my MacOS. Also tested multi client connection to single server and it worked quite nice. Collab mode is not far.

Release & Features

- Shipped **v0.7.7** — 5-level thinking system (off/low/medium/high/max), automation conditions with logical composition, parallel `call_llm`, OAuth fixes, and 10+ bug fixes (PRs #259, #253)
- Migrated Anthropic backend from deprecated `maxThinkingTokens` to adaptive thinking + effort — Pi backend now maps to all 5 levels including Codex `xhigh`
- Built speculative prefetch for parallel `call_llm` on Pi backend (works around sequential SDK limitation)
- Added "Codex Done before 11 AM" conditional automation (status change + label + time gate)

Bug Fixes & Stability

- Fixed browser "Object has been destroyed" crashes with centralized guard method across 20+ call sites
- Fixed WebSocket client connection race — stale sockets causing 4x duplicate connections per viewer window
- Fixed Haiku thinking budgets (31999 → 6000/8000) and disabled thinking for title generation
- Fixed proxy credential leak in error messages (security), error classification ordering, and surfaced proxy-specific error messages to users
- Added Node.js proxy support in network interceptor (was silently ignoring proxy settings in Node subprocesses)
- Surfaced AI provider outages (500/502/503/529) as `provider\_error` instead of generic "unknown error"

Next steps

- Deliver remote workspace (POC was already working before RPC rework)
- Monitor v0.7.7 release

Assistant

- **Branch Integration**

- Merged PI framework + server side ID mapping + server side tools branches together
- Fixed all conflicts and setup the mac app (temporarily) to use server side tools and setup a full end to end local server setup (mac with server tools, local agent service, local AI tools)

- **Remove tool level ID mapping**

- Removed every trace of ID mapping, redis cache and related code pathes from the AI tools service
- This ended up as a very large change affecting 50 files, as the ID schemas in all tools had to be changed from integer to string, so there were a lot of replacements
- **Fun fact:**
 - 32% of the weekly codex quota is gone on day 1, and I even switched back to 5.3 so its cheaper. I will have to start using claude more and use codex only for complex changes as I will run out during this week

- **Testing**

- Fixed tier 3 tests that break during the integration
- Added new tier 1 ID mapping focused tests, so I can validate the mapping changes worked okay
- Previously we only tested the server tools with targeted tool tests, but now I fully rewired our full end to end tests to be able to run with server harness
 - This highlighted further feature gaps as only **23 / 58** succeeds

- **Server tools feature gap**

- After both the end to end tests and the manual end to end testing highlighted there is a large feature gap still between the mac and server tools
- I did a thorough analysis on the differences



- **Analysis**

Here's the full comparison. There are significant gaps between the Mac app and the server-side implement...

- The major differences are coming from the HTML sync save actions, not the regular tools
- **Conclusion**

- **In order for the agent to reliably cover all capabilities just having 100% test coverage on the function level is not enough, we need 100% line coverage on all possible combinations of block sync scenarios (and I will do this in the tier 1 tool unit tests as the full end to end tests are too complex to debug and maintain)**
- **That is a lot of tests, so today I will spin up multiple parallel sessions and try to create them, then update the server tools to feature parity using the tests as benchmark**

▼ Scenarios

Operation: `**create**`

- `parentMappedId`

- content
- htmlContent
- textStyle
- listStyle
- indent
- color
- fontStyle
- alignment
- todoCompleted
- codeLanguage
- (plus all style/property attributes from the default update path — same as block content update below)

Operation: `**delete**`

- mappedId (rejected with error)

Operation: `**reorder**`

- blockOrder

Operation: `**update**`, **type:** `**page**`

- pageTextColor / pageTextColorDark
- pageColor / pageColorDark
- fontStyle
- separatorStyle
- washiImageUrl
- washiPattern
- washiColor / washiColorDark
- themeId
- wideLayout
- backdropType
- backdropColor / backdropColorDark
- backdropFromColor / backdropFromColorDark
- backdropToColor / backdropToColorDark
- backdropDirection
- backdropStyle
- backdropImageUrl
- coverImageUrl
- content + runs

Operation: `**update**`, **type:** `**tableStructure**`

- tableRowCount
- tableColumnCount

Operation: `**update**`, **type:** `**object**`

- content
- propertyValues

Operation: ****update****, **type:** ****objectList****

- activeViewId
- content

Operation: ****update****, **type:** ****objectListViewUpdate****

- viewId
- viewData

Operation: ****update****, **type:** ****objectListViewCreate****

- viewData

Operation: ****update****, **type:** ****objectListViewDelete****

- viewId

Operation: ****update****, **type:** ****objectListSchemaUpdate****

- schemaAddProperty
- schemaRemoveProperty
- schemaUpdateProperty

Operation: ****update****, **type:** **default (block content)**

Style attributes: - textStyle - listStyle - indent - color / colorDark - fontStyle - alignment - decoration - imageFillStyle - imageSizeStyle - lineThemeStyleType - layoutStyle - lineStyle

Block properties: - codeLanguage - codeTheme - cardBgType - cardBgColor - cardBgFrom - cardBgTo - cardBgImageUrl - cardThemeAssetId - cardLayout - fileName - fileUrl - linkTitle - linkDescription - linkUrl - imageUrl - todoCompleted

Table styling: - hasRowsHeader - hasColumnsHeader - alternatingRowColor - tableColumnWidths

Task attributes: - taskSchedule - taskDeadline - repeatFrequency - repeatInterval - repeatType - repeatStart - repeatSkipWeekends - repeatWeekdays - repeatMonthdays - repeatMonths - repeatYeardays

Content: - content + runs

Total: ~80 distinct attributes across 4 operations and 9 update sub-types.

- Discussed and agreed with Levy and Tamas about the caching solution for efficient space structure load

📅 Mar 18., Wed

- Continue architecture changes

↑ Mar 18., Wed · Agents team

Analysis

Here's the full comparison. There are significant gaps between the Mac app and the server-side implementation.

Operations Comparison

Operation	Mac App	Server	Notes
create	✓	✓	Both supported
update	✓	✓	Server is much simpler (see below)
delete	✗ (rejected with error)	✓	Server supports it, Mac app doesn't
move	✗	✓	Server-only
reorder	✓	✓	Both supported

Update Types (`type` field within update operations)

The Mac app dispatches on the `type` field to specialized handlers. The server **ignores the `type` field for updates entirely** — it only uses it for block type assignment (text, code, image, etc.). All these specialized update types are **missing from the server**:

Update Type	Mac App	Server
page (page styling)	✓	✗
tableStructure (row/col count)	✓	✗
object (object block)	✓	✗
objectList (list config)	✓	✗
objectListViewUpdate	✓	✗
objectListViewCreate	✓	✗
objectListViewDelete	✓	✗
objectListSchemaUpdate	✓	✗
default (block content)	✓	✓

Style Attributes

Attribute	Mac App	Server
textStyle	✓	✓
listStyle	✓	✓
indent	✓	✓
color	✓	✓
colorDark	✓	✗
fontStyle	✓	✓
alignment	✓	✓
todoCompleted	✓	✓
codeLanguage	✓	✓
lineStyle	✓	✓
layoutStyle	✓	✓
imageSizeStyle	✓	✓
imageFillStyle	✓	✓
decoration (block, focus)	✓	✗
lineThemeStyleType (washi)	✓	✗

Block Properties (all missing from server)

Property	Mac App	Server
codeTheme	✓	✗
cardBgType / cardBgColor / cardBgFrom / cardBgTo	✓	✗
cardBgImageUrl / cardThemeAssetId / cardLayout	✓	✗
fileName	✓	✗
fileUrl	✓	✗
linkTitle / linkDescription / linkUrl	✓	✗
imageUrl	✓	✗

Page Properties (all missing from server)

The entire `processPageUpdate` handler has no server equivalent:

- `pageTextColor` / `pageTextColorDark`, `pageColor` / `pageColorDark`
 - `fontStyle` (page-level), `separatorStyle`, `washi*` properties
 - `themeld`, `wideLayout`
 - `backdrop*` properties (type, color, gradient, image)
 - `coverImageUrl`
-

Table Properties (all missing from server)

Property	Mac App	Server
<code>hasRowsHeader</code> / <code>hasColumnsHeader</code>	✓	✗
<code>alternatingRowColor</code>	✓	✗
<code>tableColumnWidths</code>	✓	✗
<code>tableRowCount</code> / <code>tableColumnCount</code> (tableStructure)	✓	✗

Task Properties (all missing from server)

Property	Mac App	Server
<code>taskSchedule</code> / <code>taskDeadline</code>	✓	✗
<code>repeatFrequency</code> / <code>repeatInterval</code> / <code>repeatType</code>	✓	✗
<code>repeatStart</code> / <code>repeatSkipWeekends</code>	✓	✗
<code>repeatWeekdays</code> / <code>repeatMonthdays</code> / <code>repeatMonths</code> / <code>repeatYeardays</code>	✓	✗

Content Encoding

- **Mac app:** `content` + `runs` (separate text content and formatting runs via `parseContentAndRuns`)
- **Server:** `content` + `htmlContent` (inline HTML for formatting)

Different encoding approaches — the server uses HTML which is more natural for the AI agent, while the Mac app uses the internal runs format.

Summary

The server covers the **basics well**: CRUD operations, core style attributes (text/list style, indent, color, font, alignment, todo, code language, line/layout/image styles). But it's missing:

1. **All specialized update types** (page, table structure, object/objectList, schema, views)
2. **Dark mode color variants** (colorDark, etc.)
3. **Block-level properties** (card styling, file/link/image URLs, code theme)
4. **Table styling** (headers, column widths, alternating colors)
5. **Task/repeat scheduling** properties
6. **Decorations** and **washi/line theme** styling

The question is — which of these gaps matter for the AI agent's actual use cases? Do you want to prioritize filling specific gaps, or is this just an inventory check?

↑ Daily Notes - 2026 March

Mar 17., Tue



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Page props fix needs this week's...

↑ Mar 17., Tue

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Page props fix needs this week's release** — Christoph discovered that page props (cover image, styling) are ignored when set on subpage creation. He has a client-side fix ready. Coordinate with App Team to merge Wednesday for this week's release.

- **Cross-team: call_llm sequential execution** — Marton (Services) found that call_llm runs sequentially even when calls appear parallel, which blocked his 500 API examples localization. Flagged to Gyula — discuss root cause and whether a fix is feasible.
 - **Figma MCP client whitelisting** — Gyula discovered Figma whitelists MCP clients by client name. This could be a blocker for Figma integration — discuss whether to reach out to Figma or find a workaround.
-

 Benedek Gobolos

 Mar 16., Mon

Calendar Timeline

Did an extensive audit on Dennis's [timeline prototype](#). Balint and Christoph added some comments that helped me align on some questions.

I also added some ideas regarding Quick Capture, that's gonna be a focus after the timeline 1.0 is out.

 CRAFT

Audit

In the afternoon I started covering missing parts of the prototype related to the layout. The goal was to solve possible blockers before prototyping would start:

- How the timeline looks on iOS
- How the timeline looks when the user hasn't added calendar access.

 CRAFT

Planning

 Mar 17., Tue

- Do another search for possible blockers that could bite us later.
- Start implementing the prototype on a web playground.

 Christoph Locher

 Mar 16., Mon

Assistant Vision Writeup

Based on my research from last week, I create a writeup of the components of the next iteration of the assistant, as I see them:



Assistant: Vision for the Next Version

Core Features · After reviewing the state-of-the art of AI agent apps, a clear feature-set emerges,...

Last Edited 3/17/2026

Small Fix of the Day

For my small improvement of the day, I looked into subpage creation for collection item, which is a subject we have gotten feedback on from a number of users, including this one: [From the CraftDocs community on Reddit](#)

Essentially for the agent it was very hard to understand that collection items can also have subpages, so it often misunderstood requests about this. Especiallt setting a cover image for each item in a collection seems to be a big user request, probably because of the new gallery feature.

I added information to the format reference for the agent to better understand collection items as blocks which can have subpages and added info on how to set their properties and add a cover image. I also discovered some validation issues where the validator gave errors when adding a collection with id's already set, but the errors where misleading, and fixed it. These were server-side fixes.

 Tighten collection HTML validation and document item subpages — Refine collection validation s...

On the server side, we also prevented the agent from writing a new subpage html, when it included the id of the block in the article tag. We do not allow writing new ids to a file, so the agent doesn't invent ids when creating new blocks. But in this case it was the known id of the page block itself. The error we gave was not clear to the agent. I added an exception that if the only newly added id is that of the current block, we allow it to account for this situation:

 Clarify allowed root id on new page write — Co-Authored-By: Craft Agent <agents-noreply@craft...

Most importantly, I discovered an issue: When page props (styling props like background and font, but also cover image) are set on the subpage on creation, they are ignored (only page content is added) . So when the agent create a new subpage html and already set a cover image right away (as it would when the only reason for the subpage to exist is the cover image), it was ignored. This is a client side issue that I fixed. We should merge this on Wednesday to be part of this week's release.

 Fix page create handling for page props — Co-Authored-By: Craft Agent <agents-noreply@craft....

It would be great if we could get this into this week's release.

 Mar 17., Tue

- Visualize IA ideas from the Vision doc

Release day meets Windows reality check

Release & CI testing improvements

- Shipped **v0.7.6** — MCP custom header support, custom endpoint persistence fix, OAuth refresh loop fix, and CI source validation improvements

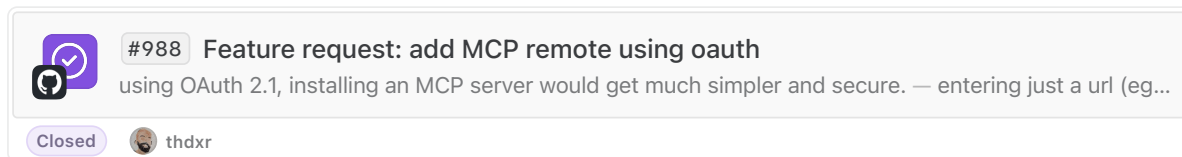
Windows Cross-Platform Fixes

- Massive debugging session on Windows thin-client ↔ Mac server connectivity — Windows paths sent to macOS server caused hangs via synchronous fs calls
- Planned & built **server-side directory browser** — new `fs:listDirectory` RPC channel, `ServerDirectoryBrowser` component, centralized `useDirectoryPicker` hook replacing native file dialogs in remote mode
- Added **cross-platform path validation** in `SessionManager` to reject wrong-platform paths with proper error events + toasts

Source Integrations

- Investigated Figma MCP issues, figured Figma whitelists some clients based on client name.

◦



- Found a nice workaround or dirty piggyback in OpenCode :)

Bugfixes

- Fixed automation subtitle showing wrong day-of-week — `computeNextRuns()` wasn't passing timezone to `Cron`

Planning & Architecture

- Explored remote server monetization concept — gaps in current architecture
- Roadmap discussion with Balint on remote setup — Phase 1 and Phase 2 planned out for the week

Other

- New MacBook setup & migration; used Windows laptop in the interim (slow but Craft Agents itself worked fine)

Next steps

- Continue Windows testing and fix remaining cross-platform issues
- Release **v0.7.7** with Windows fixes + automation conditions feature
- Start work on remote setup (Phase 1)

 Zsolt Varnai

 Mar 12., Thu

Assistant

- **Release**
 - Cherry picked collection improvements and push out a new release to prod
- **Server side ID mapping**
 - Added conditional logic to run legacy threads on new clients with client side mapping
 - New threads will be created with server side mapping
 - Remote threads downloaded to other clients will also continue with server side mapping
 - Wrote targeted tests to verify server side ID mapping
- **E2E Tests**
 - Further stabilization of flaky end to end tests
 - Now all tests are green
- **Server side tools**
 - Went through the new service code and server side tool setup with Levy
 - As Levy recommended a clean setup on the agent service side, just pulling code from his proto branch I moved over the relevant code from the new service and did a full clean integration of the server side tools into the agent service
 - Now the generate call has a feature flag, that controls the type of tool setup that we want (socket tool provider for client vs server tool provider for the API tools)
 - I changed the tier 1 tool tests, so each test runs twice, we fully execute all tests with the mac agent harness and with the API tool harness
 - Made some tweaks so we actually execute the tests via the tool provider to be closer to the real environment
 - Added bootstrap and teardown logic for the harness server, so we don't need to setup multiple services just to run the tests, the tier1 tests should be self contained
 - The integration with the tier 1 test setup runs fine, next step is merging my 2 branches: the server side ID mapping branch with the integrated server side tools, fix the relevant code and make everything work with server side mapped IDs

 Mar 17., Tue

- Continue architecture changes

↑ Daily Notes - 2026 March

Mar 16., Mon



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Levy's AI tools service — deplo...

↑ Mar 16., Mon

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Levy's AI tools service — deployed on sandbox, 58 tests passing, but likely still buggy.** Load testing against Craft KB still pending (planned for Saturday). Should confirm it happened and review results before moving to integration.
- **Craft 4 assistant scope alignment** — Christoph's writeup coming today covering pillars: Connections, Automations, Memory, Extended Capabilities, Rich Output, In-Situ Assistants. Goal is full alignment by end of week.
- **Automation conditions PR #253 needs UI design review** before release — design review with Balint scheduled for Tuesday. Mermaid flow diagram approach should be aligned on.
- **Zsolt's tool schema research conclusion:** tool results should be unified to JSON with LLM-specific transformations in the agent service. Cross-cutting architectural decision worth discussing as a team.

 Levy Melnikov

 CRAFT

AI Assistant - Server Tools | Integration Considerations

cc:  Zsolt Varnai

I finally was able to actual deploy my ai tools service and test it on sandbox E2E. It worked! Although I had to fix some bugs. Now in addition to Zsolt's phase 1 tests we also have:

- There's a separate test suite in `tools/agentik-tests/tests/*`
 - ~58 tests
 - Actual Unified API used
- Manual testing by prompting an agent to do different stuff seemed to work ok on sandbox
- Still probably contains a lot of bugs, after integration a more targeted testing should take place

I've created a writeup with nuances and concerns surfaced so far. I didn't have time to do the load testing (running prompts on Craft KB copy) - I see this as a last step.

 Mar 14., Sat

- Do load testing against Craft KB and suggested changes if anything critical pops up

After that I will get back to MCP CLI and Document History tracks

Christoph Locher

 Mar 13., Fri

- Researched tools and message formats
- Sketched out some information architectures to solve the global assistant management view, but also in-doc/folder assistant chat challenge – no clear solution yet, but moving forwards
- Quickly investigated an issue with collection item title problems

 Mar 16., Mon

Goal for the Week

At the end of the week I'd like for us to have reached full alignment on the scope, UI structure and features for the assistant for the Craft 4 release, so we have a roadmap on what to build and design. To get there I will create writeups, visualisations and prototypes.

Today

- Finish a writeup of my ideas for the assistant and the pillars of the next version (Connections / Automations / Memory / Extended Capabilities / Rich Output / In-Situ Assistants) for us to discuss

Goal(s) of the Week

- Align on how to introduce remote capabilities within Craft Agents
 - Will put together a feature list on Monday, that we can discuss and extend
 - Also a vision based on last weeks interviews with the team
- If the alignment is done, start the implementation for remote capabilities - have a POC by friday is the stretch goal
- Release Automations enrichment with conditions
- Prepare new youtube video about automations
 - Collect some use cases from Craft Agents Beta group and use the mas inspiration
- Thorough windows testing to catch Windows platform niggles

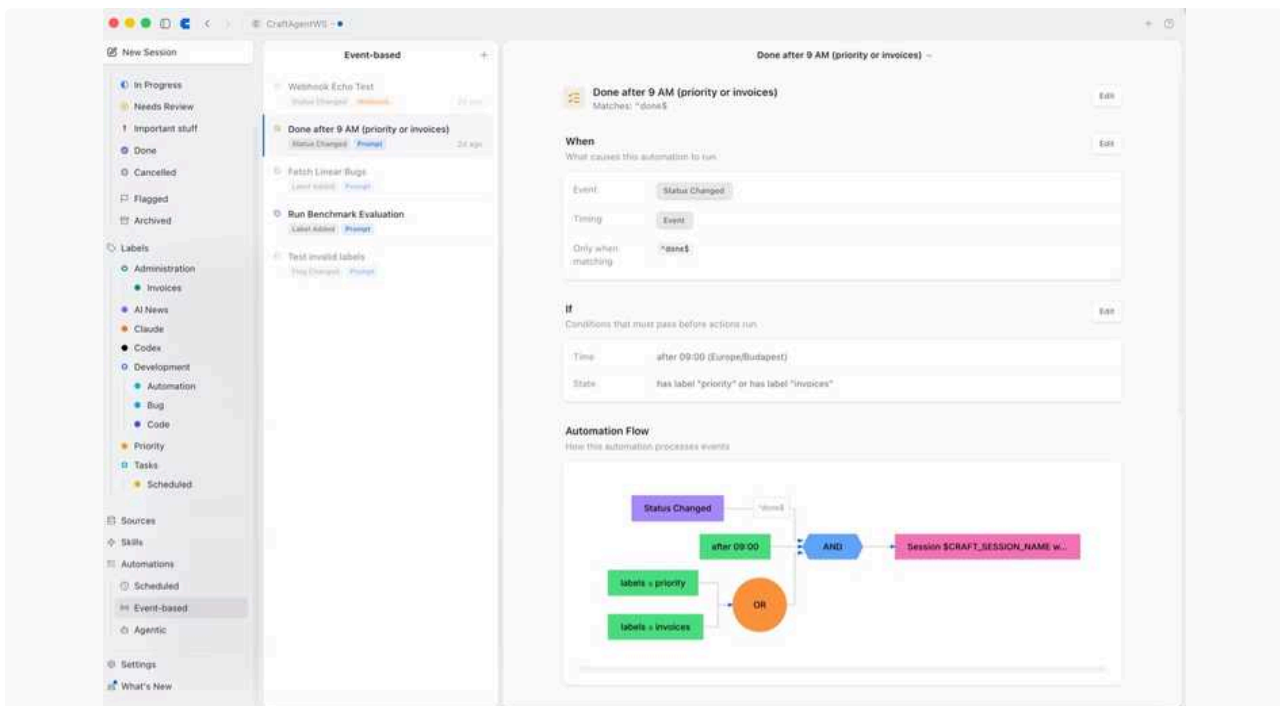
Friday the 13th: MCP custom headers prepared, automation conditions ready for review

Support MCP servers with custom headers

- [PR #254](#) — Custom headers for MCP HTTP sources (merged). Fixed #410: static `headers` in config, credential-store `headerNames` for secret headers, header precedence logic, and CI validation with `--validate-server`.

Feature Development

- [PR #253](#) — Automation conditions + flow diagram (open, needs UI design). Home Assistant-inspired conditions engine (`time`, `state`, `and/or/not`) gating automation actions. Added Mermaid-based flow visualization in the automation info panel to represent increasingly complex automation logic. 47 new tests, 203 total passing. **Needs UI design review** — the Mermaid diagram approach works well but should be aligned on before release cc  Balint Orosz



Source Integrations

- Set up Linear MCP and created BillyBob Craft source — thorough MCP integration testing

Demos

- Created a YouTube video demonstrating the Browser Use feature in Craft Agents

Next steps

- QUICK Hotfix release - 0.7.6
 - There was a persistence issue with Custom Endpoints, fixed it and also tested it with better-ccflare (Claude Proxy) on my machine, and will set it up in the CI this way we can test PiSDK as well as the Claude Agent SDK on all platforms !!!
 - Also with Marci's fix last time there was an issue where it assumed a token was refreshed for new MCP OAuths in some cases also fixed that
 - Also added support for MCP servers that use custom headers (like Google Stitch) and also added it to the e2e suite in Github
- Remote solution planning — test in Windows environment to identify gaps, put together a fix package for Windows support
- Remote design doc — cover UI aspects, feature scope, and architecture for the remote solution
- Design review with Bálint on Tuesday (March 17) — discuss remote design doc and automation conditions UI (PR #253 Mermaid flow diagram approach)

Assistant

- **PI integration**
 - Added OpenRouter support
 - Added support for passthrough mode (non agentic)
 - **Stability improvements**
 - Reviewed and fixed remaining issues with the implementation
 - Fixed issues with hook handling
 - Fixed timeout handling
- **Server Side ID Mapping**
 - Created optional server side ID mapping behind feature flag
 - Persist mapping state in thread context
 - Wire in the mapper throughout the generation flow when interacting with blocks or documents
 - Always interact with real IDs when calling client socket APIs
 - Added tool translation layer to replace IDs (input/output)
 - Mapping unit tests

- **Tool Schema Research**

- ▼ **Current Tools**

Category	Tool	Returns
Read	blocks_get	Block
	blocks_get-metadata	Block
	blocks_get-tasks	Block
	blocks_get-table-of-contents	Block
	space_list-tasks	Block
	space_list-calendar	Block
	document_find-blocks	Block
	document_get-comments	Comment
	list-documents	Document
	find-documents	Document
	space_view-folder-structure	Folder
Create	space_create_document	Document
	space_create_daily_note	Document

	space_duplicate_document	Document
	space_create_folder	Folder
Update	space_edit_document	Document
	space_edit_folder	Folder
Hierarchy	blocks_move	Block
	space_move_document	Document
	space_move_folders	Folder
Delete	blocks_delete	Block
	space_delete_document	Document
	space_delete_folder	Folder
Other tools	user_info_tool	User
Sync APIs	load_document	Block JSON
	save_document	Block JSON
	load_block	Block JSON
	load_space_structure	Document/folder JSON
Other APIs	resolve_urls	URLs

- Reviewed current tools and potential schema changes we should make for future proofness
- **Current Setup**
 - The majority of the current tools already implement raw CRUD operations
 - Create, update and delete are very straightforward, for property updates
 - Exceptions are block updates, which happen via HTML sync and hierarchy operations (move), which enforce strong consistency requirements
 - There are "shortcut" tools, which do something properly without exposing a huge complexity to the agent
 - Duplicate document recreates the document with all the smart capabilities to fix IDs, rewrite blocklinks, etc
 - Create daily note handles all the special metadata objects, that are behind daily notes
 - Read tools are giving specialized views for the entities and expose context dependent data
 - They often implement queries into underlying data, like regex search in the plaintext index, or query task metadata or daily note objectlink relations
- **Conclusions**
 - A lot of tools already implement very basic CRUD operations, which seem to be a good building block long term
 - For the "Read" category there is a lot more diversity in how the data can be queried and presented to the LLM

- eg we have some tools, that return plaintext, instead of JSON
- **We should unify the tool results to JSON and isolate the complexity of mapping them into LLM friendly content**
 - **(eg the folder tree should be created in the agent service, the tools should return JSON folder and document list)**
- For future tools the agent service should handle LLM specific transformations and potentially can reuse existing client side tools if available
 - Eg current server side **edit_block** is a good example, the tool allows single block editing by calling the underlying client socket API, which is enabling HTML sync, so the client didn't have to reimplement edit_block on its own
- Some tools could be moved to server side by introducing more dependencies on the server, that we previously tried to minimize
 - eg resolving URLs could happen on server side by calling the service that exports the S3 token

 Mar 16., Mon

- Continue PI integration

Weekly Goal

- Continue architecture changes to fully support multiple clients in a future proof way
 - Finish server side id mapping
 - Tool refactor
- **Investigations**
 - Semantic search
 - Custom compacting

 **Balint Orosz**

Integrated the task list into the app. Overall the direction is very promising, but lots of areas are touched I need to go through and we need a lot deeper integration in many areas.

As a start I refactored heavily navigation so it's more robust as we needed more advanced navigation constructs (sidebar items behave as tabs) in order to feel right.

I'm progressing with exposing stronger session state/task state and refining the UI all along.

↑ Daily Notes - 2026 March

Mar 13., Fri



Agents

Recommended for discussion · generated by Daily Team Discussion Points Agent · Viktor & Growth team request —...

↑ Mar 13., Fri

Agents

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Viktor & Growth team request — credit usage banners:** There's an open request (flagged by Christoph) from Viktor and the Growth team to show banners/messages to free users at 50% and 85% credit usage. Needs discussion and assignment.
- **Spaceld architecture question (Levy):** Levy is using a workaround (ignoring given spacelds, always using the default) to unblock Server Tools productization. Open question raised: does the agent actually need to know/receive spaceld given it operates on a single space at a time? Needs a decision.
- **9 unified APIs were mocks (Levy):** Levy discovered 9 unified APIs were only implemented as mocks. Worth discussing scope and prioritization of fixing these.

From  **Christoph Locher** : There was a request from Viktor & the growth team to display banners / special messages in the assistant for free users when they reach certain points of credit usage (50% of credits used, 85% of credits used)

Server Tools productization

I'm still in process of refining and polishing.

The most trouble is coming from id mapping. I thought that id mapping was only used in tools and that if we replicate that behavior there → everything will work. Turns out id mapping is used in tools, but also outside of tools in context generation: spaceld is mentioned in context. This means that no matter how I map ids on the server side there will be conflicting ids on the client.

A good breakdown of this is here:



Craft Agents Session Viewer

The issue is small-ish, we just have to come up with a different mechanism to set spaceld in context:

- Just propagate it from the backend somehow
- Use full id - there's only 1 space id in a session so the impact of not shortening it should be small
- Not specify it at all
 - Why does the agent need to know the id? Why is it even an argument in tools calls? Don't we operate on a singular space at a time?

To avoid this issue for now I made the code to ignore given spacelds and always use the default one. This allowed me to progress much more:

- Fixed more bugs - biggest one was that I found out that 9 unified apis were only implemented as mocks
- Refactored a lot of code, I now feel that agent starts to "get" the code again

Short Id Savings

I've found an easy self-contained way to measure token savings due to short ids. Based on 1 conversation it's **~10%**. But tomorrow I will try more different conversations:



Craft Agents Session Viewer

 Mar 14., Sat

- Finish polishing
- Deploy to sandbox, test prepare PR
- Prepare for integration
- Investigate more how much tokens short ids save

Christoph Locher

 Mar 12., Thu

- Ideation & Sketching
 - Today was a bit more messy than I would have liked because I had to do some fixes to get the folder icon PR into the release and ready for testing, and felt a bit under the weather in the afternoon. I have a complete focus day planned for Friday to catch up.
- Folder Icon PR
 - Peter gave me advice on the agents PR feedback multiple times (🙏) and since I was able to make the skill work with a separate icon library markdown file that the agent greps we merged it without the huge error message containing all icons (46kb)

Shipped 0.7.5 with webhooks

- Webhooks are finally out
- Community bug fixes and MCP improvements included - was a bit eager here, need to be less trigger happy

Investigations and testing

- We were looking into Numeric Citizens issue a but deeper,  Marton Wernigg did some improvements on Craft MCP tool description side, i have updated the craft source guides

Next steps

- Create the first video of the new series
- Add conditions Home Assistant style to automations that would improve the general usability
 - Trigger tells when to trigger
 - Condiitiin tells If it should trigger
- Start planning out the remote set up

 Mar 12., Thu

Assistant

- **PI integration**
 - **Tracing**
 - Added langsmith tracing support for request messages, tools, LLM results and usage stats
 - As relying on the SDK callbacks can lead to misleading traces (like it happened with Claude tracing) I tried implementing it lower level with the fetch interceptor, sending the traced request data to the main process for tracing calls. This seems to provide much more raw logs, allowing to catch issues easier.
 - **Tool permissions**
 - Handle explore/execute mode based tool permissions
 - Also saw a MultiEdit tool in the PI tool list, which seemed promising, so I tried to enable it for more efficient editing, that costs less
 - Unfortunately it looks like its not a real thing in PI. Claude didn't remember how it got there, but I found it in the craft agents repo's tool name list, probably not relevant for PI.
 - Integrated cancellation support via the subprocess
 - **Multi modal support**
 - Tested file handling support with various file types
 - Images work fine with the custom image resolution tool + PI's read tool
 - Other file types, like PDF were not supported
 - Although Anthropic backend API supports file message type to trigger the vision API, this would result in a non provider agnostic solution, that only works for Anthropic.
 - Also PI does not support the file message type either, so it would be problematic to implement
 - At the end I went with similar solution as craft agents and implemented custom text extraction, that we execute immediately in the file resolution tool, providing both the local file path and the extracted text content for certain types
 - PDF, XLS, DOC, PPT
 - ZIP extraction
 - I enabled ZIP extraction, which mostly worked, but result in long running extraction due to the content.
 - I added a timeout to make it viable
- **Release**
 - Release some new improvements to prod to support release testing
 - Also checked and merged Chris' gateway URL changes to access the gateway via the local network to avoid the timeouts. Deployed it to prod which seems to be stable.

 Mar 13., Fri

- Continue PI integration



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Server tools deployment — perfo...

↑ Mar 12., Thu

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Server tools deployment — performance issues expected once live:** Levy has server tools passing the same tests as client tools but they are not yet deployed. He expects deploying to the Craft KB copy will surface performance issues. Good to set expectations on timeline and who should be looped in.
- **Pi agent testing blocked by Minimax custom endpoint issues:** Gyula noted Pi's capabilities are hard to validate without Minimax set up, and there are still issues with custom endpoints. Needs a decision on a testing approach so Pi evaluation can proceed.
- **Webhooks still unfinished:** Listed by Gyula as the next priority but not yet started. Worth a brief status check given it's been flagged multiple days.

 Levy Melnikov

E2E Testing Loop

I managed to set everything up so that existing tier 1 tests from Zsolt are ran against the new server side tools.

Running actual tests surfaced a lot of big & small issues, I'm now in the process of solving them. Good news is that with a full agentic loop it was relatively easy to fix.

- Connected server tools solution to Zsolt's tests
- Fixed 72 tests
- Added the MemoryDB for storing long-short id mappings

 Mar 12., Thu

Current state: server tools pass the same tests client tools pass, but are not deployed

I will work on deploying and cleaning up. I expect that testing a deployed version on Craft KB copy will surface performance issues for fixing.

 **Christoph Locher**

 Mar 11., Wed

- Research & Writing
- Set up edit_folder_tool icon for PR for the Friday Release

 Mar 12., Thu

- Writing & Sketching

 **Gyula Halmos**

Branching wars, v0.7.4 shipped

Release & Deployment

- Shipped **v0.7.4** with extensive bugfixes
- Prepared release announcement

Bug Fixes & Hardening

- Spent way more time on branching fixes than expected — state recovery, lineage guards, Windows-specific issues
- Model selector now actually selects models
- Fixed several smaller, hard-to-reproduce bugs (automation triggers, CI, path expansion)

Product

- Worked on Craft Agents product messaging copy

CI/CD

- Successfully ran multi-platform e2e test for the server: [GitHub Actions run](#)

Next steps

- Finally finish webhooks
- Look into custom endpoints — still issues with Minimax
- Better understand Pi's capabilities, but hard to validate without Minimax set up — need to figure out a testing setup
- Expand e2e tests with more cases (branching, more providers beyond Anthropic) to cover the PiAgent path — would give much more confidence in releases

 Mar 11., Wed

Assistant

- **PI integration**

- Solved various issues with PI subprocess setup and communication
- Tool integration
 - Testing tool handling scenarios
 - Investigated and fixed schema conversion issues
 - Hide protocol differences with claude sdk
- Fixed pre and posttool hook handling
 - Integrated proper hook triggers for PI
- Now conversations work with PI, including tool calls, reading document via HTML and editing changes are also correctly applied back to Realm via HTML sync 🎉
 - I will continue checking/integrating more finegrained capabilities, like file and image handling
- Some LLM learnings
 - Compactions are crazy frequent now and there is a huge degradation in output quality when that happens
 - Checking files to get up to speed after compaction already consumes a huge percentage in the very beginning
 - Claude's compaction is terrible, because it even forgets stuff, that was the initial message of the thread to explain the entire goal
 - In my PI session in one of million compactions the agent forgot what and where craft agents repo is and even though I refreshed this callout every time, asking for it to "check why it works in craft agents and copy the solution" it completely ignored what I asked. "craft" and "agents" are very generic words in the other repos too, maybe because of that
 - I think until this becomes better I will copy paste and repost my initial prompt after every compaction and also have to very thoroughly check the tool call chains to see if it does what was asked or just checking files in the current repo

- **Support**

- Helped Levy setup the end to end environment and build the Mac agent harness for testing

 Mar 12., Thu

- Continue PI integration
 - Test and integrate additional capabilities
 -



Agents

Recommended for discussion · generated by Daily Team Discussion Points Agent · Server tools architecture — two o...

↑ Mar 11., Wed

Agents

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Server tools architecture — two open questions from Levy** — (1) Can ID shortening be dropped for server tools? (2) APIs that are easy client-side may be slower/harder server-side (e.g. returning all document titles with metadata) — how flexible can the tools/APIs be? These are explicitly raised and a decision is needed before Levy sculpts the tools.
- **Folder Icon skill blocked by oversized SKILL.md** — Christoph prototyped the skill but the full icon list makes the SKILL.md too large to be usable. How should large reference data be handled in skills? Needs a design decision.
- **Claude agent SDK bug #23440 impacting users** — Zsolt identified the root cause: interrupting during the "Thinking" phase writes invalid whitespace content, causing every subsequent API call to fail with a 400 error. Current workaround: start a new session. Worth discussing whether a better mitigation or fix is feasible.
- **Retry/fallback policy coordination with Services** — Chris (Services) is updating AI usage tracking; Zsolt noted the Agents SDK retry/fallback policy may need adjustment accordingly. Cross-team coordination with Christian needed.

 Levy Melnikov

Server Tools

I've got big shapes in places: vibe coded 14 tools and 7 rpc commands and the infrastructure around them. So far it's very rough, very sloppy - I have done only a simple manual test. I saw both tool error and success and it managed to answer a basic question of "What's in my space". So the foundation is in and it's time to refine.

Now that I have something e2e I see some nuances:

1. Solution uses id shortening? Can we ditch it for the server tools?
2. There are multiple APIs that are easy and fast on the client, but might be more complicated on the backend. For example returning all of the documents with their titles might be slower on the backend. Can we not return titles there? How flexible are we about the tools \ apis?

 Mar 11., Wed

Setup a e2e testing harness and start sculpting the tools to ensure that they're 1to1 exactly the same as the old thing.

Ship it: v0.7.3 release day with a PR review marathon

Release

- Released v0.7.3 from the 0.7.2-bugfixes branch

Features & Development

- Background task implementation — major feature work (PR #238, 579+ messages of deep implementation)
- Title generation rework — improvements, reviews, commit & push (PR #239, PR #240)

Code Review & Bug Fixes

- PR #241: Bugfixes — grumpy review + impact analysis on source creation & cross-platform compatibility
- PR #237: Corner case analysis deep-dive (500+ message investigation)
- PR #234: MCP OAuth token refresh — analyzed reconnect behavior on token refresh
- Investigated behaviour change in MCP connection handling between sessions from the Claude Agent SDK — connection lifecycle differs across session boundaries
- Plan validation review (500+ messages, bug fixes)
- Application startup ENOENT error investigation

Infrastructure & CI/CD

- Updated validate-server GitHub workflow to test on both Linux and macOS in parallel

Next steps

- Continue with webhook introduction
- Put together the foundation message

Christoph Locher

- Research & Ideation
- Investigated some user issues
- Fixed some validation issues
- Prototyped a Folder Icon skill (but the SKILL.md with the list of icons is too large)
- Synced with Bene about the coding theme

Assistant

• End To End Testing

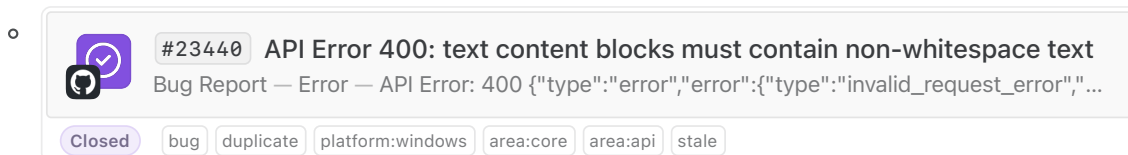
- ID handling
 - The block ID handling in the tests caused issues, as both the mapped and unmapped IDs could end up being equal ("1", "2", "3"), so replaced the test IDs in all test files to have a prefix
- Tier 4 tests
 - Worked on increasing the stability of the tier 4 tests
 - Added retry logic to reduce flakyness
 - Fixed race conditions in search indexing when the iOS agent harness starts

• Release Support

- Follow up on the 403 errors that a user reported
 - As they get the same error on iOS and web this issue seems unrelated to agent-v2 (as web is not using it), the 403 should be happeninng on the AI gateway
 - Most likely they are using some shared document, that we need the the information of to investigate
- Investigated the errors, that 2 users reported



- After some research found, that it's a claude agent SDK bug



- **Root cause:** interrupting during the "Thinking" phase saves "content": [{"type":"text","text":"\n\n"}] to the conversation JSONL. Every subsequent API call includes this invalid block and fails.
- Their workaround is "start a new session"
- Manually editing the session JSONL could fix the issue, but at this point it's most likely not worth for us to try fixing this

• AI Gateway

- Continued discussion with Chris about gateway cost differences
- Cancelled/aborted requests cannot be tracked properly, but Anthropic still counts them partially, so this is 1 source of the differences for sure
- Failed streaming requests in the agent SDK also fallback to HTTP, which usually end up causing more timeouts and disconnects
- Chris is making changes in the tracking and we might need to adjust the retry/fallback policy in the SDK
- **PI framework**
 - Framework abstraction
 - As a first step I created a general framework abstraction layer to isolate SDK specific code from the rest of the service logic
 - Introduced AgentRunner and ClaudeSDKRunner to abstract away the framework hooks
 - The change broke the tier2 and tier3 tests due to setting up the infra of the tests, so had to fix those.
 - Added PI
 - Going with the simplest setup that fits the existing architecture
 - We still need to channel everything through the gateway and might need the interceptor for adjusting the outgoing headers and other attributes
 - We will create custom subprocess architecture to isolate the sessions from the main service process
 - Replaced PI and started fixing the issues that happen during generation

 Mar 11., Wed

- Adjust retry mechanism of the production service
- Continue PI integration

↑ Daily Notes - 2026 March

Mar 10., Tue



Agents

Levy Melnikov · 1 Tool Example · I was able to implement the simplest server tool user_info_tool, integrate it into luki...

↑ Mar 10., Tue

Agents

 Levy Melnikov

1 Tool Example

I was able to implement the simplest server tool `user_info_tool`, integrate it into `luki-ai-v2` agent and test e2e locally. Throughout this process I had to solve a lot of devex issues → essentially none of it worked until I started all 4 service locally: native client, luki-ai-agent-v2, luki-ai-tools(new) and luki-ai-gateway. Now I have a nice(all local, all log files are available to the agent) setup. One thing that's missing is an ability to emulate ios client without running it, so that agent can have full feedback loop, I will consult with Zsolt about that.

With this I think I now have a strong grasp on the system and the task at hand.

Weekly Goal

Deliver server tools:

- Makes `luki-ai-agent-v2` independent of native tools
- Unblocks Web AI Assistant to use `luki-ai-agent-v2`

 Mar 10., Tue

Extend the 1 tool example to actually cover all of the client tools. Hopefully we will have "working, but flaky" server tools by EOD

Weekly goals:

- Improve perceived stability - fix as many niggles as possible
- Improve overall stability - improve core functionality edge cases - review architecture when and where needed
- Ideally create the first video of a new series conducted with Katya and based on her guidance, not ideally define the foundational message and story arch which can be used to chain the videos
- If getting there, start thinking about the pre-requisites for the remote exec, create plans

PR Marathon, 0.7.2 Release & Infrastructure Hardening

Release & Features

- **Released Craft Agents v0.7.2** — cut release from bugfix branch

Code Review & Bug Investigation

- Reviewed **5 PRs** across craft-agents and Luki.Serverless:
 - craft-agents PR #229 — Webhooks implementation (deep review, created worktree) WIP
 - craft-agents PR #230 — Reviewed
 - craft-agents PR #231 — Collection bugfix PR (architectural integrity focus)
 - craft-agents PR #232 — Reviewed in two sessions
 - Luki.Serverless PR #4411 — Reviewed
- OSS Issue Analysis — comprehensive severity/impact assessment

Architecture & Planning

- Explored **Tape** reference architecture for agent frameworks - can be an interesting approach to keep it in the back of our mind
- **Title generation rework** — created dedicated worktree, shared prompt architecture (WIP, targeting Wednesday)
 - Title generation should use the language set in the user preference to generate titles
 - Working on the logic on how to decide what the title should be, the last 3 messages is not accurate enough
- **Webhooks feature** — plan review, cross-validated with GPT-5.3 Codex (WIP, targeting Wednesday)
 - Considered feature parity with Claude Code, Zapier and Home Assistant
 - Took ideas where it made sense
- Identified search logic improvement opportunity (branched as idea)

Product & Strategy

- Discussion with Katya about product messaging framework — she shared a structured template covering user structure, problem narrative, belief shifts, category definition, positioning evolution, and messaging pillars

Research

- Claude Marketplace overview — explored the concept

Next steps

- Clean up PRs and OSS issues in GitHub after the 0.7.2 release
- Prepare another bugfix release
- Finish title generation rework and webhooks feature (by Wednesday)
- Fill in Katya's messaging framework template for Craft Agents and review

Christoph Locher

Goal for the Week

Gain a deep understanding of the state-of-the-art for AI Agent driven products by testing with and exploring the most common apps. I want to both get a feel for the industry standards, but also strengths and weaknesses of various tools, and map their challenges to ours. This should start triggering some ideas for how to bring advanced capabilities to the Craft agent.

Tools to Review

- OpenClaw
 - via Dashboard
 - via Telegram
 - via [Chowder](#)
- Notion AI
- Perplexity Computer
- Codex
- Claude Chat
- Claude Cowork
- Claude Code (UI)
- Cursor
- Dia
- Comet
- Craft Agents, of course

 Mar 9., Mon

I set up OpenClaw, but it was quite the struggle. I'm not sure why, but first I kept getting "API is Overloaded" errors all the time, then the configuration got messed up, then cron jobs got setup but never run. After resetting everything three times I managed to get a working setup and now have it all hooked up and running via telegram, experimenting with some different connections and input methods.

A friend of mine developed an OpenSource app that is an OpenClaw iOS client with some very nice LiveActivities, that I also want to try:



newmaterialco / chowder-iOS

github.com/newmaterialco/chowder-iOS



newmaterialco

Swift

Stars 532

In the evening, I added some more icon capabilities to the Edit Folder tool, along with icon validation and a useful error message:



Exten edit_folder tool to accept nucleo and streamline icons, validate icon name

Commit

b747f32f

Exten edit_folder tool to accept nucleo and streamline icons, validate icon name



lukilabs / Luki



christophlocher

3/9/2026

Assistant

- **End To End Testing**

- Worked on all 4 tiers of tests in parallel
- Implemented bunch of new tests, covered all tools and more edge cases
- Did several review rounds of what's missing from the coverage and extended wherever the test ideas had reasonable complexity to setup
- The tests that verify the tools ran too slow sequentially, so experimented with concurrent test running options
 - Even though bun has a concurrent option, it did not speed up the execution as expected, most likely because calling the iOS agent harness is a blocking call, so it won't help in a single process
 - At the end split up the tests into multiple files and run them in separate process, which made them much faster
 - Added test result aggregation and final summary at the end
 - For agent work having the ability to run the tests super fast will be very valuable
- iOS agent harness
 - Added several new capabilities to the harness as new type of tests required fully isolating the craft environment in harness mode (eg space level tests require workspace API and other things)
- Investigating and fixing issues with test flakiness
- Now tier 1, 2 and 3 are close to complete and all successful
 - I will do some final review on them, but otherwise should be good
 - We will be able to swap out the iOS agent harness to Levy's new APIs soon to test feature parity
- The tier 4 end to end tests are now significantly improved and went from the majority failing to the majority being successful, mainly due to fixing infrastructure issues with the test, isolate test environment and avoid shared state with the iOS harness
 - I will finish today by checking the remaining few failing tests

- **AI gateway cost tracking**

- Discussion and investigation with Chris about the cost tracking discrepancy
- Anthropic does not provide low level details, only the 1 day aggregate cost value, so it's hard to investigate
- We double checked API keys and what other components might use the account without the gateway
 - We found, that most of the stuff is already using a 2nd sandbox account, so not many things use the production account, but a slack thread has been started to find if anyone might be still using a production API key (we might also just disable any unknown key in the prod account just to make sure)

- Chris will continue checking the new logs and comparing gateway records and datadog logs to identify the root cause
- **Server side tools**
 - Had a knowledge sharing session with Levy and shared all the relevant details and needs for the server side tools

 Mar 10., Tue

- Continue assistant tech stack improvements
 - Some final test improvements
 - Check remaining few failing tier 4 tests
 - Final review
 - Start work on swapping clade agent SDK to PI framework
- **Goal for the week**
 - Have a working agent setup with PI framework locally
 - Including special flows, like asset handling etc
 - Identify bigger gaps and feature parity issues, where some functionality cannot be replicated in PI without bigger infra changes (Im writing this down just to be thorough, but based on Craft Agents learnings there might be no gaps at all and ideally we can fully replicate all capabilities)
 - **Note**
 - The main goal is to have PI working, but there are multiple levels to this where we can go and if the integration goes smooth we can most likely achieve more advanced levels this week, eg:
 - Agent SDK is replaced with PI
 - Integrating a fully custom agentic framework interface and have complete abstraction in front of any framework functionality
 - Potentially have Claude Agent SDK available as a swap in
 - The infra changes can even go further, but these are more unrelated from PI
 - eg isolating the socket infra from the agentic loop so we don't have to rely on a live socket connection for running the agent

↑ Daily Notes - 2026 March

Mar 9., Mon



Agents

--- · Christoph Locher · Testing latest updates in Codex, Perplexity Computer, Comet, Dia · Deep-dive into the SDK I...

↑ Mar 9., Mon

Agents

 Christoph Locher

- Testing latest updates in Codex, Perplexity Computer, Comet, Dia
- Deep-dive into the SDK layers of Craft Agents with Gyula
- Travel home

 Mar 9., Mon

- Deep dive into Notion's Assistant model & behavior



- There still were issues in branching - fixed it
- Annotations - this was a quite big work from a UI perspective, but I am quite glad how it turned out. 90% was creating a general annotation component which works across fullscreen etc, making it delightful etc. I use it now a lot, I think it makes asking questions a lot easier.
- Lots of other bugfixes (i.e. source creation was broken from EditPopover as deeplinking change with IPC, but also lots of other)
- All Changes now on Master

This week :

- Main Goal is getting Tasks inside the app. It's a bigger, multi step process so will be on separate branch for quite a while

v0.7.1 shipped, webhook automations in the pipe

Release & Bug Fixes (v0.7.1)

- Released v0.7.1 (10 commits, 27 files, +933/-165) — stability-focused patch release
- Fixed branch creation infinite hang — ERR_STREAM_WRITE_AFTER_END race in SDK's ProcessTransport, added 15s timeout with interrupt-based cancellation (#345)
- Fixed cross-machine session recovery — stale sdkSessionId never persisted as cleared (#342)
- Fixed base64 false positives — strict canonicalization pipeline replacing heuristic detection (#344)
- Fixed spawn ENOENT during auto-update — structured error fields + bounded single retry (#268)
- Fixed duplicate LLM connections on re-auth — stale React closure + updateOnly server-side guard
- Fixed Codex token expiry — routed through typed_error auth-retry pipeline, global refresh mutex
- Fixed Settings "Check Now" stuck at 0%
- Automation sessions now use automation name as session title

Feature Development

- Webhook actions for automations (PR #229, open) — new action type that sends HTTP requests from automation triggers. Full stack: types, handler, event bus, RPC test, UI. Supports all HTTP methods, custom headers, JSON/raw body, env var expansion.

Experiments & Browser Automation

- Created 15-question Office Work Environment Survey via Google Forms (browser automation + Apps Script)
- Built reusable google-forms skill for the two-phase approach
- Salesforce demo request flow, design tools pricing comparison (Figma vs Sketch vs Penpot)

Community

- Answered questions in the Slack craft-agents-beta channel
- Responded to discussions on GitHub (craft-agents-oss)

Next steps

- Prepare v0.7.2 bugfix release
- Research webhook implementations in Claude Code, Home Assistant, Zapier, n8n, etc. before finalizing our webhook action design (PR #229)
- Review & merge PR #229 after evaluation
- Meeting with Katya
 - Discuss goals
 - Unique selling point
 - Main message for the new series of videos - format discussions

 Mar 6., Fri

Assistant

- **Release Support**

- Investigated Nori's findings about background sessions and space switching

- **End To End Testing**

- I started implementing the tests for the different tiers
- The work here seemed to be very well parallelizable, so started the work on 8 separate sessions
 - 2 sessions per tier, 1 codex, 1 claude, I created the first test examples with both, to decide which one I will use for coding
 - Based on how they set up the mocking and test flow I decided to continue implementing the rest of the tests with claude
- **Agent Harness**
 - Also had a session only focusing on the iOS Agent Harness, that makes the cross platform tool calls work as a "CLI"
 - Added proper APIs to allow importing and exporting test data and the ability to call all available tools via CLI
 - Also simplified some boiler plate and created auto registerable tool setup here as well, to make it future proof
- **Tier 1**
 - Started implemented tests for all tools
- **Tier 2**
 - Created various tests for
 - Pretool use guardrails
 - Explore/execute mode handling
 - HTML sync
 - HTML validation
- **Tier 3**
 - Created stream event handling related tests
- **Tier 4**
 - Created end to end flow, where we have a special toolprovider, which runs the iOS agent harness instead of socket
 - Started porting the test from the iOS end to end list and making the end to end flow succeed

 Mar 9., Mon

- Continue assistant tech stack improvements
 - Finish end to end testing work



Agents

--- · Gyula Halmos · 0.7.0 ships while the codebase gets a structural reckoning · Release & CI · Managed the Craft A...

↑ Mar 6., Fri

Agents

0.7.0 ships while the codebase gets a structural reckoning

Release & CI

- Managed the Craft Agents 0.7.0 release — coordinated release branch prep and final checks
- Debugged GitHub Actions validate-server workflow failures on the release branch
- Verified repo state on latest main post-release

Architecture & Refactoring

- Deep architectural review of the IPC refactor branch — identified deduplication opportunities and created a streamlining plan
- Set up git worktree for CLI tool development off the IPC refactor branch, extensive CLI tool implementation work

Content

- Prepared intro YouTube video script (Browser Use) — reviewed by Katya

Next steps

- Monitor 0.7.0 release, fix any issues that come up
- Start working on bugfixes to improve stability
- Record 10-minute intro YouTube video based on Katya's feedback
- Follow up with Katya on Monday re: video
- Will need to jump out to doc at noon (driver's license renew)

Christoph Locher

- Release Alignment
- Synced about testing & found issues with Nori
- Design team sync & one on one with Balint
- Fixed three minor issues (styling prompt not always applied, linting, hide tailored style button when assistant is disabled)

Assistant

• Release Preparation

- Supported Nori with release testing and investigating potential issues

◦ Thread Handling

- Improved remote threads download timing
- Enforced space level filtering for displayed remote threads
- Manage and display unread badges per space

◦ Investigations

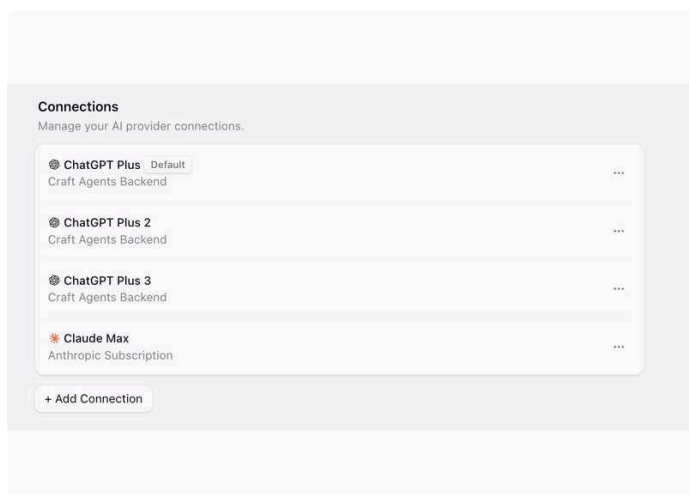
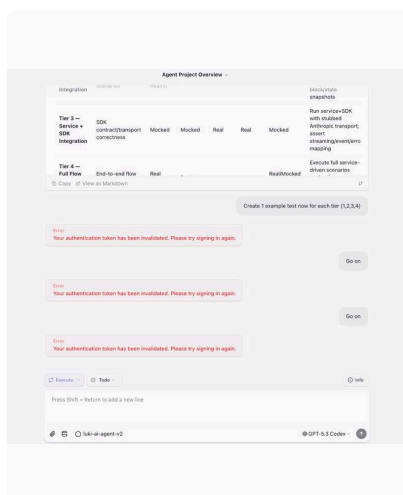
- Investigated credit handling edge cases with thread grouping
- Investigated space loading issue

◦ PR review and merge fixes for release build

- I tried this time using the PR skill with codex, which spawned up the subagents according to the skill correctly
- There were no false positives this time, which is great, although it was a smaller PR, so we have to keep testing
- The performance was an issue, the review took a very long time. Later I learned from Gyula that the parallel subagents are not yet actually parallel, so the slowness is due to the agents actually running sequentially. To switch to codex review we will probably need the parallel capability to avoid very slow reviews

• Codex feedbacks

- I switched to codex for the reviews and end to end test planning and overall have very positive experience with it
- One blocking technical issue is, that my token times out after 24 hours and my existing sessions become dead and cannot be restored
 - Reauthenticating creates another ChatGPT entry, so its not possible to switch the dead thread to the working one ( Gyula Halmos)



- **End to end testing**
 - Explored different end to end test setups and created a plan
 - **Swift Agent Test Harness**
 - The ability to run the swift tools from the service tests is an important capability, because it enables some of the testing tiers, that we couldn't do without it
 - I explored the options how can we invoke the swift files
 - Calling them with the swift interpreter technically is possible, but having to make all the imports works and abstract away any symbols and dependencies that are outside the file would be very complex
 - Alternatively the client could connect via websocket, but its problematic, because the goal would be to enable autonomous agentic testing, without manual intervention, making sure the app is running, etc
 - At the end I found an alternative, very convenient solution, a new **Agent Tool Harness** app target
 - Virtually it can be considered as an **Agent Tool CLI**, that the service tests can use to import test data, run any tool with provided JSON input, receive output and export realm data
 - The harness can execute any of the app's registered tools and uses an isolated temporary test realm based on the session ID parameter
 - It's technically not a real CLI, because in order to make it compile I had to keep mac catalyst target, but it runs without UI and terminates immediately after execution, so there is virtually no difference
 - **Test Tiers**
 - **Tier 1 — Harness Integration**
 - We can cover the tool behavior with tests, that we can execute both against the swift and server side tests to ensure feature parity
 - **Tier 2 — Service + Harness Integration**
 - We mock everything, that happens on the agent SDK side, but our side of the code remains with real integration to test end to end complex flows, like HTML sync for example.
 - For example: The test simulates an HTML write → Mock SDK triggers tool hooks → The agent code that we are testing does the HTML sync → We verify if the changes triggered the correct block updates
 - **Tier 3 — Service + SDK Integration**
 - Similar to tier 2, but the Agent SDK is real, we mock everthing on the anthropic API level. These tests can verify that tool hooks and other SDK behavior is correctly integrated when certain streamed events are triggered by the LLM.
 - **Tier 5**
 - Tier 5 are the full end to end tests from the native app, that do not mock everything, the entire stack is evaluated. These are the 80 tests, that we already have and can be manually run from the app
 - **Tier 4**

- Tier 4 is kind of a free win, that become possible due to the Agent Tool Harness. We just port the existing Tier 5 tests from the swift code and run them from server side, so the coding agent can easily run them anytime it wants to verify something
- Instead of using the real iOS app we just call the harness, the end to end tests can run on the temporary realms for the test sessions.

Tier	Purpose	Client with data	Client-side tools	Agent service	Claude Agent SDK	Anthropic API	Test Framework Responsibility
Tier 1 — Harness Integration	Swift tool/data correctness	Real (harness Realm)	Real (harness)	-	-	-	Call harness directly with fixture args; assert JSON + state
Tier 2 — Service + Harness Integration	Orchestration + HTML/sync scenarios	Real (harness Realm)	Real (harness)	Real	Mocked	Mocked	Run service tests against harness adapter; assert orchestration + final block/state snapshots
Tier 3 — Service + SDK Integration	SDK contract/transport correctness	Mocked	Mocked	Real	Real	Mocked	Run service+SDK with stubbed Anthropic transport; assert streaming/event/error mapping
Tier 4 — Full Flow E2E (Harness Client)	End-to-end flow without full app runtime	Real (harness Realm)	Real (harness)	Real	Real	Real/Mocked (suite-dependent)	Execute full service-driven scenarios against harness client; assert end-to-end output + data side effects
Tier 5 — True Full E2E (App Runtime)	Production-faithful end-to-end validation	Real (app runtime)	Real (app runtime)	Real	Real	Real	Run user-like end-to-end app+service scenarios; assert user-visible outcomes and data consistency

- Created example test setups for each tier

 Mar 6., Fri

- Continue assistant tech stack improvements

↑ Daily Notes - 2026 March

Mar 5., Thu



Agents

Recommended for discussion · generated by Daily Team Discussion Points Agent · Beta release readiness: Christop...

↑ Mar 5., Thu

151

Agents

created with 


Recommended for discussion

generated by Daily Team Discussion Points Agent

- **Beta release readiness:** Christoph was testing the release candidate and doing alignment conversations. Zsolt merged the PR to master for the beta build. The team should confirm: is the beta build ready to submit? Any last-minute blockers?
- **Zsolt's load test revealed CPU easily hits 100%** with only a few parallel sessions. He suggests investigating after the PI framework switch, but worth discussing whether this is an acceptable risk for the upcoming release.
- **Zsolt's baseline memory concern:** He noted the service baseline memory is "a bit too high" — not a blocker now, but flagging for awareness.
- **Zsolt found rows disappearing after collection row insert** — root cause was a race condition in UI stamp checks. He fixed it with batch processing (single update instead of 10-20), but the team should verify this fix holds in the beta.
- **AI code review quality concern:** Zsolt noted many false positives from AI review feedbacks that only surfaced after insisting multiple times. He suggests using better models (codex or main agent instead of haiku subagents) for PR reviews — worth discussing as a team practice.
- **Gyula's Phase 4 server extraction is underway** — 10,800 lines across 28 files being extracted into `@craft-agent/server-core`. Largest session was 1,796 messages (\$77). The team should be aware of the scope and check if any architectural decisions need broader input.
- **CX Team cross-reference:** Tamas (CX) found a multipanel focus-jump bug in Craft Agents and couldn't fix it. The focus management seems app-internal. Zsolt or the team may want to assist or at least be aware of the branch `fix/tab-focus-jump-inline-menus`.

 Christoph Locher

- Added a “Tailored Style” button that prompts the agent to use its skill and create a custom style based on the document content



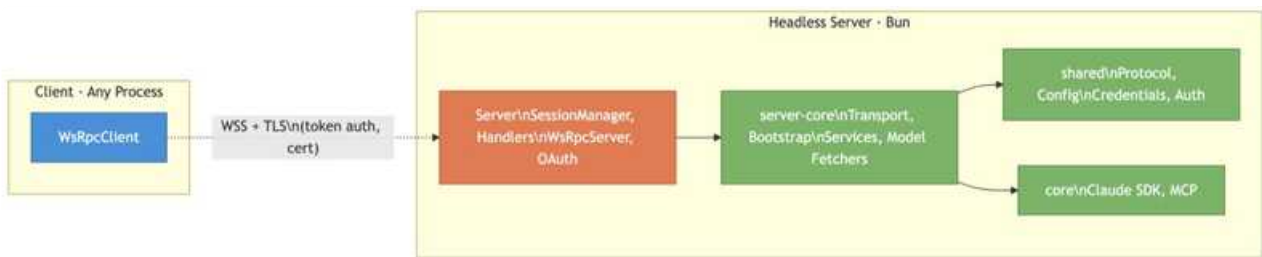
- Tweaked some of the suggested prompts
- Prepared visuals for the What's New slider
- One on one with Zita
- Lots of alignment conversations for the beta release and release submit
- Testing the release candidate

Phase 4 kicks off — 10,800 lines escaping Electron

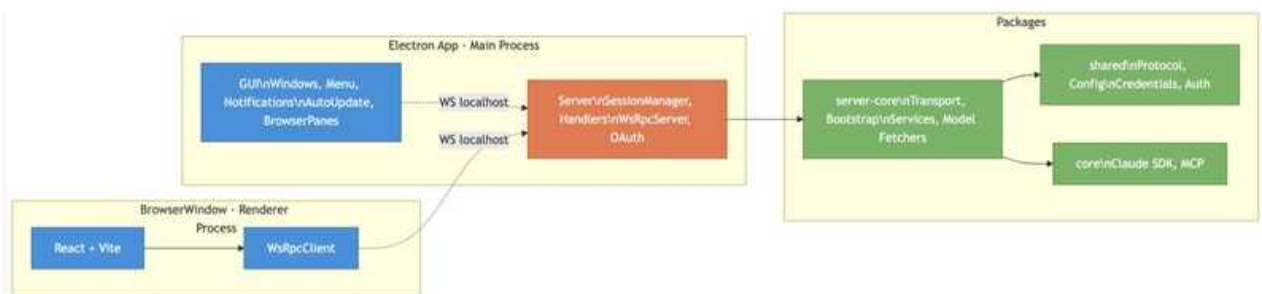
Architecture & Server Extraction

- Deep-dived WebSocket RPC communication flow on the IPC refactor branch — mapped how browser control sequences flow client→server end-to-end
- Reviewed Phase 3 standalone server plan against codebase — found it too high-level, produced comprehensive architecture report surfacing missing extraction prerequisites
- Authored Phase 4 plan: extract shared protocol, utilities, handlers, and SessionManager (~10,800 lines across 28 files) into `@craft-agent/server-core`, then scaffold `packages/server` as standalone CLI
- Began Phase 4 execution — largest session of the day (1,796 messages, \$77)
- Ran parallel serverless packaging analysis using GPT-5.3 Codex for cross-model validation
- Created architecture diagrams for full app component layout and headless server topology

Headless Architecture:



Normal standalone (Local) App architecture:



Package structure: before → after

Before — domain logic trapped in `apps/electron/`:

```

1 apps/electron/src/
2   └─ main/
3     └─ handlers/*           ← ALL handlers (core + GUI mixed)
4     └─ sessions.ts         ← SessionManager (5,812 lines)
5     └─ search.ts           ← platform services
6     └─ image-utils.ts
7   └─ shared/types.ts       ← IPC_CHANNELS + DTOs (1,789 lines)
8   └─ transport/types.ts    ← re-export barrel
9   └─ server/start.ts       ← imports ../main/* (6 of 8 imports)

```

After — clean extraction into packages:

```

1 packages/
2   └─ server/               ← NEW standalone CLI (bunx @craft-agent/server)
3     └─ src/index.ts
4   └─ server-core/
5     └─ transport/
6     └─ runtime/
7     └─ bootstrap/
8     └─ handlers/          ← EXTRACTED core handlers (14 domains)
9     └─ session/           ← EXTRACTED SessionManager
10    └─ services/          ← EXTRACTED search, image-utils, broker
11    └─ domain/            ← EXTRACTED pure utilities
12  └─ shared/
13    └─ protocol/          ← NEW: RPC_CHANNELS, BroadcastEventMap, DTOs
14  └─ core/
15
16 apps/electron/src/
17   └─ main/handlers/       ← GUI-only remainders (browser, system-gui)
18   └─ server/start.ts      ← thin wrapper, zero ../main/ imports

```

Next steps

- Continue testing and wrap up the PR
- Do some cross-model architecture validation
- Test TLS with remote server set-up
- Explore multi-tenancy (not planned to do now, just evaluate the possibilities)
- Explore combination of remote and local (if there are any architectural blockers)

Assistant

• Stability

- Checked Fargate available disk storage space for live sessions (20 GB by default, can be increased later)
- Added disk usage metrics and warnings for storage space running out
- Created load test script to see if the service is unstable under load
 - There were no stability issues
 - I feel the baseline memory usage of the service is a bit too high, but it did not increase drastically during the test, so it should be fine for now
 - The load test was easily able to put the CPU to 100% load even with only a few parallel sessions, which I don't really like
 - We can investigate this further, although maybe its worth putting the effort to this after the PI framework switch, to avoid profiling something, that we are going to replace either way
- Checked the existing end to end tests (79/80 succeeded)
 - The last one was flaky due to the prompt not being specific enough, so improved the test prompt
- Talked to Chris and retested gateway OpenRouter usage tracking via the Anthropic skin. It worked for me again, so we are not sure where did the invalid cost tracking IDs came from earlier, might have been a transient OpenRouter issue on their side.

• Compacting

- Disabled compacting and tested the context limit error
- Improved error handling for context limit error
- Added staging support for credit site

• Schema Handling

- Did a research about potential backward compatibility issues with the schema
- Most of it already supported graceful degradation, but added a few more improvements to a few places to make it even more flexible
- Did some test changes in the schema and verified that the client did not crash or cause visible issues

• Models API

- Added model list cache and default model list to avoid the Max and Fast models disappearing when there is no connection to the server

• Improvements

- Review PR and improved code quality based on the AI review feedbacks (there were a lot of false positives, that only came out when I insisted multiple times, so it might be worth to try reviewing with better models, eg codex or if using claude, then use the main agent, not the haiku subagents)
- Investigated rows disappearing after collection row insert
 - Replaced API usage to dedicated collection entity APIs

- The rows disappearing is most likely caused by some race condition in the UI stamp checks and state updates
- Added batch processing, so the collection is updated only once, instead of 10-20 times, which solved the issue
- Background session improvements to avoid false positive unread badges
- Deployed and tested prod service, fixed some AWS specific issues with skills.
- Merged the PR to master for the beta build

 Mar 5., Thu

- Assistant tech stack improvements

↑ Daily Notes - 2026 March

Mar 4., Wed



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · IPC→WebSocket rework: scope...

↑ Mar 4., Wed

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **IPC→WebSocket rework: scope was 3-4x the initial estimate.** Gyula's "just swap the transport" turned into touching 224 IPC channels across a 3,400-line ipc.ts. The rework is functionally complete and running on a remote Linux server, but Gyula notes WS may be noticeably slower than IPC. He plans to build Windows/Mac versions and manually test all key features today. Discuss whether there are specific scenarios to prioritize in testing.
- **New message protocol issues identified.** Christoph did extensive testing with the new message protocol, identified and reproduced issues, and discussed them with Zsolt. The team should align on the priority and fix timeline for these issues.
- **Christoph ran a session with CX and Growth teams.** He demoed the agent, approach, and possibilities. Worth discussing any feedback or follow-up actions from that session.
- **OAuth redesign: tokens no longer egress to client wire.** Gyula redesigned OAuth so the server owns all tokens — a security improvement worth noting. Verify this doesn't break any existing flows during testing.

Ripping out IPC — the surgery was bigger than the scan suggested

Architecture & Planning

- Planned and executed a fundamental shift from Electron IPC to WebSocket as the universal transport layer — what started as "just swap the transport" turned into touching 224 IPC channels across a monolithic 3,400-line `ipc.ts`
- Iterated through 6 versions of the unified WS architecture plan, arriving at a symmetric WS core + command-based JSONL protocol with clear capability boundaries between client (shell) and server (runtime)
- Redesigned OAuth so the server owns all tokens — they never egress to the client wire anymore

Infrastructure & Deployment

- Spun up headless WS server on Kamatera Ubuntu VM (EU, 2CPU/4GB, port 9000)
- Worked through bootstrap issues: config creation in headless mode, logger mode detection, xvfb setup for Electron-dependent code paths
- Binary payload serialization (Uint8Array) over WS transport required dedicated fixes

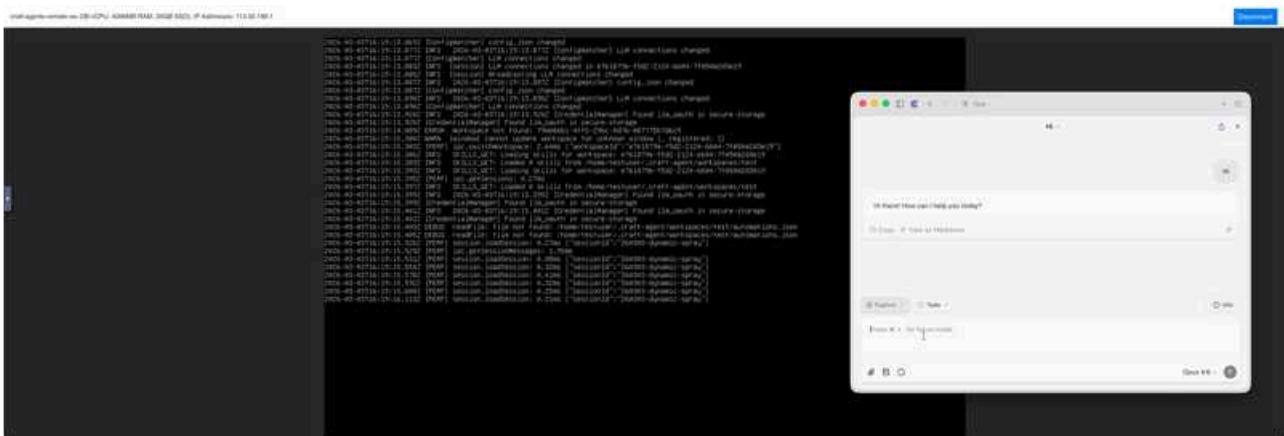


refactor(ipc): namespace channels + typed broadcast dispatch

#223 · Opened by rjulius23

Merged

Craft Agents running on a remote Linux Server using WebSocket transport layer. In itself it is not a big thing, but i put extra effort to make the architecture clean. Removed IPC and decoupled the UI/Client from the Server (main process) logic as much as possible. Claude wanted to always cut corners, but Codex was strict.



Research & Analysis

- Deep-dived Electron renderer↔main process communication paths to map all IPC touchpoints before cutting
- Analyzed Bartek's `remote-v2` branch for mobile app architecture patterns

- Reviewed OAuth flow rework to fit the thin-client/server-authority model
- Explored OpenClaw's credential management: static secrets use a plain-text `secrets.json` with a provider option for external fetchers (e.g. 1Password CLI). OAuth tokens are stored separately as plain text, continuously refreshed — no encryption on either.

Lessons Learned

- **IPC was load-bearing everywhere** — what looked like a transport swap revealed IPC was deeply entangled with domain logic, OAuth flows, credential storage, and even binary data handling. The "blast radius" of removing it was 3-4x the initial estimate.
- **Plan versioning pays off** — iterating from v1→v6 on the unified WS plan caught architectural gaps (OAuth locality, binary payloads, bidirectional RPC) that would have been painful bugs in implementation.
- **Remote deployment exposes hidden assumptions** — things like "open a browser for OAuth" or "read a local config file" break silently when the main process runs on a headless Linux box. Every `shell.openExternal()` call is a locality violation in remote mode.

Next steps

- Built my branch for win and MacOS so today i will use it for my daily chores (prepare a smaller bug fix package)
- I want to cover all aspects manually as well, added unit tests, but the key features want to test and also WS can be a bit slower than IPC so want to see the feel if noticable

Christoph Locher

- Implemented a single-line check in the **automatic order-fixing** pre-tool hook, did some more testing and committed it to the main branch

 Automatically fix order of classes and attributes when writing the doc

Commit

cbeccd651

Automatically fix order of classes and attributes when writing the doc

 lukilabs / Luki.Serverless  christophlocher 3/2/2026

 Only automatically fix order if it's a single line

- Had a long **session with customer success and growth** to show them the agent, the approach and the possibilities and answer all of their questions
- Tweaked the **agent / model / mode badge** – it still needed lots of tweaking to get the animations right for the state transitions, because there are so many different possible combinations of going from one state to the other with various icons appearing and changing position and the labels updating.

 Consolidate AI chat badge and improve transitions — Summary — Merge model, agent, and...

Commit

34d79f97

Consolidate AI chat badge and improve transitions — Summary — Merge model, agen...

 lukilabs / Luki  christophlocher 3/3/2026

- While tweaking the assistant badge on iOS I also fixed an issue where the right mode was not getting selected when loading a thread, which only happened right after app start. It took a bunch of digging with Claude to identify the root cause, but then the fix was straightforward

 Fix mode badge not showing on fresh app launch — When restarting the app and opening t...

Commit

0a313213

Fix mode badge not showing on fresh app launch — When restarting the app and open...

 lukilabs / Luki  christophlocher 3/3/2026

- A lot of testing with the new **message protocol**, identified and reproduced issues and discussed them with Zsolt, fixed some UI glitches

Assistant Editing

- **Message Schema V2**

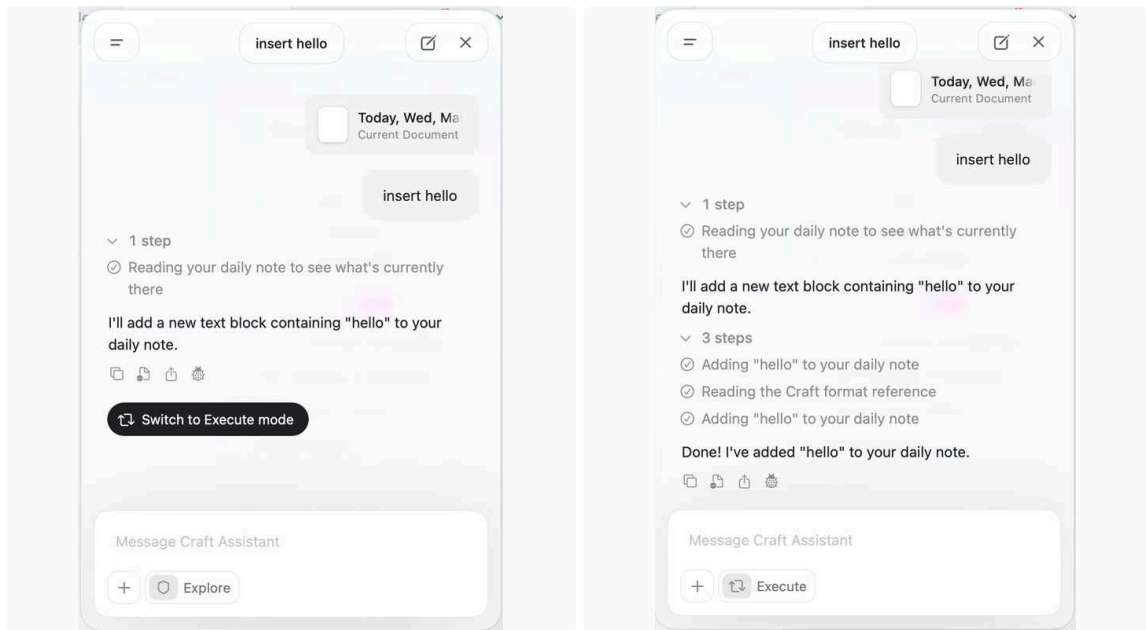
- Updated remote thread fetching API to V2 schema
- Fully removed V1 schema from the service code
- Cleaned up all feature flags including for V1/V2 schema support and other capabilities that we added in the last few weeks
 - As discussed we are abandoning the beta version to reduce technical debt
 - The only remaining feature flag is for space/document HTML sync mode, because not all agent types support editing
- The server → clients messages were using V2 types, but the generate call's message input was still using legacy models
 - Updated the generate API to also use the same new message types as we have in the results

- **Context handling**

- With the new messages I added JSON metadata to the context messages coming from the client
- The `<context>...</context>` string goes to the LLM, the JSON metadata object is used for reconstructing the context item in the history (eg when loading a remote thread)

- **Grouping**

- Updated grouping based on the discussion with Christoph
- Fixed issues with 2 groups being merged due to execute mode switch
- Fixed issue with grouping staying in pending state due to execute mode switch
- Added special logic to handle the execute mode switch tool call
 - The tool call itself is intentionally hidden
 - Added a **messageToUserAboutIntendedChanges** parameter, that the LLM fills
 - We create a special text message from the content of this parameter and insert it as a message outside the group with the switch button
 - The generation also gets force canceled when this tool call happens to make sure there is no summary response displayed by the LLM



- **Tooling**
 - Created new tool for daily note creation
- **Session Handling**
 - Improved background session activity tracking to avoid unintended badges appearing when background thread connection state changes
 - Fixed timing issues with persisting session storage

📅 Mar 4., Wed

- Assistant testing and improvements

↑ Daily Notes - 2026 March

Mar 3., Tue



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · IPC → WebSocket migration plan...

↑ Mar 3., Tue

Agents Team

Recommended for discussion

generated by Daily Team Discussion Points Agent

- **IPC → WebSocket migration plan submitted.** Gyula completed a deep architecture analysis (12 sessions, multiple models) and submitted a formal implementation plan to replace Electron IPC with WebSocket. The goal is local and remote sessions executing the same way. Discuss the plan, timeline, and any concerns before implementation begins.
- **Levy (Services Team) asking about mod system traction.** Levy explicitly asks Gyula: "is there any traction on the mod system?" He is actively using the mod repo and branch that supports it, manually updating to the latest versions. This needs a direct answer.
- **Assistant editing performance experiment — promising early results.** Christoph experimented with auto-fixing attribute/class ordering in the PreTool hook instead of making the agent fix it manually. Early results suggest it speeds up editing without negative consequences. Discuss whether to merge this approach and expand the experiment.
- **New assistant message protocol alignment.** Christoph aligned with Zsolt on the new message protocol for the assistant. Worth a brief team update on what was decided and implications for the codebase.
- **Release and beta planning alignment.** Christoph met with Peter, Viktor, Lisa, and Nori regarding release and beta planning. Share key outcomes with the rest of the Agents team.
- **OSS bug fixes shipped (PR #222).** Gyula fixed LiteLLM/OpenRouter config saving, cross-workspace source leaking, and window drag region issues. Quick awareness for the team.

IPC tear-down day — rethinking the entire transport layer

Architecture & Planning

- Deep-dived into the full IPC channel architecture of craft-agents across 12 sessions, using both Claude Opus and GPT-5.3 Codex for varied perspectives
- Designed a protocol abstraction layer to consolidate many IPC channels into a unified messaging protocol
- Planned clean separation of concerns: transport logic (auth, connection management, heartbeat, failure handling) vs. business/protocol logic
- Iterated through multiple plan reviews culminating in a formally submitted implementation plan for WebSocket migration
- Direction: throw out Electron IPC entirely, switch to WebSocket, simplify the architecture

OSS Bug Fixes (PR [#222](#))

- [#326](#) — LiteLLM/OpenRouter users can't save connection config (model name rejection)
- [#323](#) — Sources from other workspaces leaking into current workspace UI
- Window drag region blocked by TopBar no-drag containers

Next steps

- Iron out architectural quirks — ensure the design is as simple as possible
- Goal: local and remote sessions execute the same way via WebSocket
- Start testing and exploring corner cases

- Fixed a small issue with coding languages in the validator
- Aligned with Zsolt on the new message protocol for the assistant
- Aligned with Peter, Viktor, Lisa and Nori regarding release and beta planning

- **Background Thread Handling on iOS**

More testing and fixing with background thread completion and closed foreground thread completion, mostly on iOS. Now displaying and indicator in the closed state on iOS too when there are unread messages

- **Experiment: Increasing Performance by Auto-Fixing Attribute Order in the PreTool Hook**

I was thinking about ways to improve the tool use of the agent, and after reading through the logs to check the most common validation errors the agent makes when editing, started wondering if we really have to force the agent to manually fix the attribute and class order.

In order to guarantee stability and be able to compare the before and after effectively, we discovered a few weeks ago that we need to enforce a consistent order of attributes and classes (both are now alphabetically sorted). We added this to the prompt and added validation for it and since haven't had any problems with it. But the agent still orders the classes wrong on the first try so so many times, and then needs to call the edit tool again with the fixed order, which slows everything down.

In this experiment (on a separate branch), I **extended the PreTool hook** where the validation happens, to **auto-fix the order of attributes and classes** before writing the files, and also to add a `systemMessage` to the agent to remind them of this rule for the future.



Github

Directory not found

<https://github.com/lukilabs/Luki.Serverless/tree/zsolt/ai-agent-v2-autofix-order>

The prompt I used for testing is this: `Please turn the paragraphs under the sub headlines into a list and use serif font for them`. The agent almost always uses the order `"list-bullet font-serif"` as that is the order of the requests in the prompt, but it should sort them alphabetically.

On this branch, the edit tool fixes them before writing.

I was worried that this would slow it down even more because changing the file would force the agent to reread it after every edit (like it has to when blocks are added and their ids are generated), but so far I did not observe this. Instead so far it seems like this speeds up the editing without negative consequences. I will keep investigating this, but it could be a good way to improve performance of edits.

 Zsolt Varnai

 Mar 2., Mon

Assistant Editing

◦ Message Schema V2

- Brainstorming and planning for the goal that we would like to achieve with the new message schema
- Created core schema for the most important messages
 - Messages have a **role** and **type**
 - **Roles:** user, system, assistant
 - **assistant types:** text, reasoning, tool-call
 - **user types:** prompt (later could be permission approval, and other actions)
 - **system types:** prompt, context, error, switchMode
- Added JSON metadata for context items (eg current page, all docs, etc)
- Put the new schema behind feature flag to support old clients
 - Added a conversion layer that maps to V1/V2 schemas
- Created JSONL based append only session storage for the messages
- Grouping
 - Added server side group ID and grouping logic
 - Removed client side grouping
- The JSONL goes through a consolidation before the remote thread is fetched
 - Tool call start events are removed when tool result is available
 - Hidden system messages are filtered out
 - Unfinished tool calls are canceled

◦ Improvements

- Fixed unnecessary update loop for remote thread fetching

 Mar 3., Tue

- Continue with V2 schema support

↑ Daily Notes - 2026 March

Mar 2., Mon



Agents Team

Recommended for discussion · generated by Daily Team Discussion Points Agent · Assistant Editing beta released —...

↑ Mar 2., Mon

Agents Team

166

Recommended for discussion

generated by Daily Team Discussion Points Agent

1. **Assistant Editing beta released** — Zsolt added feature flags for backward compatibility, reviewed PR changes, and deployed the service. Discuss early feedback and any issues observed.
2. **Planning mode — legacy vs new client capability** — Zsolt reviewed Christoph's prototype PR and added backward compatibility for legacy plan handling. Ensure alignment on which clients support the new planning mode and when legacy support can be dropped.
3. **Custom skill handling approach** — Zsolt implemented a simpler skill loading approach (returning content in tool result) vs. the built-in method (injecting as user prompt). Decision made to go with simpler approach unless issues arise. Monitor for session compaction problems.
4. **Stale session handling refactored** — Zsolt refactored stale session checks to be more conservative — only on assistant panel open, no stamp comparisons during active changes. Christoph also worked on thread restoration (5-min and unread items) for both macOS and iOS. Verify these changes work well together.
5. **New message API/format** — Christoph plans to design a new improved message API/format based on Turn display today. Discuss requirements and scope.

 Feb 27., Fri

Assistant Editing

- **Beta release**

- Added feature flags for backward compatibility
- Reviewed PR changes
- Merged to master and deployed service

- **Planning mode changes**

- Reviewed Christophs prototype PR
- Added back legacy plan handling with plan.md special permissions purely for backward compatibility
- Handle both scenarios and tools based on client capability

- **Custom Skill Handling**

- Created custom skill handling solution and removed built in skill usage
- Skills are copied to the user's claude session
- The custom skill logic automatically discovers the list of available skills and injects it to the model's context (into the skill tool's description)
- The agent can load any of the skills on demand, just like with the built in solution
- The skill tool actually just reads the skill's markdown file and returns the content to the agent in tool result
- I did an exploration session with claude to understand how the built in skill loading works and it turned out that the built in skill tool just returns a "loading skill" status in the tool result, the skill content is not returned. The skill is actually injected as a user prompt immediately after.
- There are some speculations why its done this way, which is definitely more complex than just returning in the tool result. It might react better to session compacting for example, although this is unconfirmed.
- Until we actually see any issues with it I decided to go with the simpler solution (and we don't want session compacting anyway)

- **Session Handling**

- Updated remote session download flow so session configuration is also applied immediately, when the remote session becomes available on the device and opened
- Refactored how stale session handling works to reduce the number of edge cases that can happen when we are running multiple background sessions and then switch back to them, triggering stale session checks on sessions, that are currently being mutated both server and client side
 - Now stale session handling is more conservative, its done on assistant panel open to determine staleness flag, but no later stamp comparisons are done, when the threads are actively changing

 Mar 2., Mon

- Continue with editing support

 Christoph Locher

 Feb 27., Fri

- **Improved & Tested Assistant Session/Thread Handling and Background Operation**
 - macOS
 - Brought back 5 minute thread restoration
 - Introduced thread restoration when thread has unread items
 - iOS
 - Introduced 5 minute thread restoration
 - Introduced thread restoration when thread has unread items
 - Fixed thread unread indicator
 - Fixed sending a thread into background mode by starting a new thread
 - Fixed setting of model when bringing a thread back from background mode

 Mar 2., Mon

- Plan new improved message API/format based on Turn display

Bugfixes and remote spike, e2e with automations

- Prepared a set of bugfixes in this PR

**fix: stability improvements (oversized images, metadata leak, OSS build)**
#221 · Opened by rjulius23 · Merged by balintorosz

Merged

- Started exploring remote setup:
 - It requires a new interface for the IPC calls and a transport manager for the websocket connections
 - I managed to spin it up on a remote Linux server and connect to it
 - Credentials needs to be remote as well so for now i set up that infra in CloudFlare
 - The POC worked fine with API keys for now
- I have created 2 skills one to benchmark and another one to eval the sessions
- The e2e suite uses automations to spin up different providers and models and run the benchmark skill, they can be triggered by label or time schedule
- The eval skill checks the sessions with the relevant tag and scores them

Next steps

- Discuss the remote setup corner stones
 - I would make it possible in the OSS to run in server mode, and as a next step figure out how to provide hosting via our infra
 - So the first refactor would just make it possible to spin up Workspaces that are remote
- Finish e2e eval